

```
In [ ]: import yfinance as yf
import pandas as pd
import math
from sklearn.metrics import mean_squared_error
import random
from datetime import datetime as dt
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
import numpy as np
import matplotlib.pyplot as plt
from statsmodels.tsa.stattools import adfuller
from statsmodels.tsa.arima.model import ARIMA
from sklearn.preprocessing import MinMaxScaler
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout, LSTM
import warnings

warnings.filterwarnings('ignore')

# Set display options for better output
pd.set_option('display.max_columns', None)
pd.set_option('display.width', 1000)

random.seed(42)
np.random.seed(42)
tf.random.set_seed(42)
```

Prepare Dataset

```
In [ ]: # Data Download Configuration
# S&P 500: January 1, 2002 to December 31, 2023
# Bitcoin: January 1, 2015 to December 31, 2023

# Define date ranges
sp500_start = "2002-01-01"
sp500_end = "2023-12-31"
bitcoin_start = "2015-01-01"
bitcoin_end = "2023-12-31"

print("Downloading S&P 500 data...")
sp500_data = yf.download("^GSPC", start=sp500_start, end=sp500_end, progress=False)
```

```
print("Downloading Bitcoin data...")
bitcoin_data = yf.download("BTC-USD", start=bitcoin_start, end=bitcoin_end, progress=False)

# Display basic information about downloaded data
print(f"\nS&P 500 Data Shape: {sp500_data.shape}")
print(f"S&P 500 Date Range: {sp500_data.index.min()} to {sp500_data.index.max()}")
print(f"Total S&P 500 observations: {len(sp500_data)}")

print(f"\nBitcoin Data Shape: {bitcoin_data.shape}")
print(f"Bitcoin Date Range: {bitcoin_data.index.min()} to {bitcoin_data.index.max()}")
print(f"Total Bitcoin observations: {len(bitcoin_data)}")

sp500_data.columns = sp500_data.columns.set_levels(['Adj Close' if x == 'Close' else x for x in sp500_data.columns.levels[0]], 1)
bitcoin_data.columns = bitcoin_data.columns.set_levels(['Adj Close' if x == 'Close' else x for x in bitcoin_data.columns.levels[0]], 1)
```

```
In [ ]: sp500_data['Log_Returns'] = np.log(sp500_data['Adj Close'] / sp500_data['Adj Close'].shift(1))

bitcoin_data['Log_Returns'] = np.log(bitcoin_data['Adj Close'] / bitcoin_data['Adj Close'].shift(1))

sp500_clean = sp500_data.dropna()
bitcoin_clean = bitcoin_data.dropna()
```

```
In [ ]: statistics_data = {
    'Observations': [len(sp500_clean), len(bitcoin_clean)],
    'Mean_Daily_Return': [sp500_clean['Log_Returns'].mean(), bitcoin_clean['Log_Returns'].mean()],
    'Standard_Deviation': [sp500_clean['Log_Returns'].std(), bitcoin_clean['Log_Returns'].std()],
    'Minimum_Return': [sp500_clean['Log_Returns'].min(), bitcoin_clean['Log_Returns'].min()],
    'Maximum_Return': [sp500_clean['Log_Returns'].max(), bitcoin_clean['Log_Returns'].max()],
    'Skewness': [sp500_clean['Log_Returns'].skew(), bitcoin_clean['Log_Returns'].skew()],
    'Kurtosis': [sp500_clean['Log_Returns'].kurtosis(), bitcoin_clean['Log_Returns'].kurtosis()],
    'Infinite_Values': [np.isinf(sp500_clean['Log_Returns']).sum(), np.isinf(bitcoin_clean['Log_Returns']).sum()],
    'NaN_Values': [sp500_clean['Log_Returns'].isnull().sum(), bitcoin_clean['Log_Returns'].isnull().sum()]
}

returns_statistics = pd.DataFrame(statistics_data, index=['S&P_500', 'Bitcoin'])

returns_statistics.T
```

```
In [ ]: # Create visualizations
fig, ((ax1, ax2), (ax3, ax4)) = plt.subplots(2, 2, figsize=(15, 10))

# S&P 500 Price Series
ax1.plot(sp500_clean.index, sp500_clean['Adj Close'])
```

```

ax1.set_title('S&P 500 Adjusted Close Price (2002-2023)')
ax1.set_ylabel('Price ($)')
ax1.grid(True, alpha=0.3)

# Bitcoin Price Series
ax2.plot(bitcoin_clean.index, bitcoin_clean['Adj Close'])
ax2.set_title('Bitcoin Price (2015-2023)')
ax2.set_ylabel('Price ($)')
ax2.grid(True, alpha=0.3)

# S&P 500 Log Returns
ax3.plot(sp500_clean.index, sp500_clean['Log Returns'])
ax3.set_title('S&P 500 Logarithmic Returns')
ax3.set_ylabel('Log Returns')
ax3.set_xlabel('Date')
ax3.axhline(y=0, color='red', linestyle='--', alpha=0.5)
ax3.grid(True, alpha=0.3)

# Bitcoin Log Returns
ax4.plot(bitcoin_clean.index, bitcoin_clean['Log Returns'])
ax4.set_title('Bitcoin Logarithmic Returns')
ax4.set_ylabel('Log Returns')
ax4.set_xlabel('Date')
ax4.axhline(y=0, color='red', linestyle='--', alpha=0.5)
ax4.grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

# Display sample data
print("S&P 500 Data (First 5 rows):")
print(sp500_clean[['Open', 'High', 'Low', 'Adj Close', 'Volume', 'Log Returns']].head())

print("\nBitcoin Data (First 5 rows):")
print(bitcoin_clean[['Open', 'High', 'Low', 'Adj Close', 'Volume', 'Log Returns']].head())

```

Stationary Test

```

In [ ]: from statsmodels.tsa.stattools import adfuller
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf

sp500_adf_test = adfuller(sp500_clean['Log Returns'])
print('ADF Statistic: %f' % sp500_adf_test[0])

```

```
print('p-value: %f' % sp500_adf_test[1])

plot_acf(sp500_clean['Log_RetURNS'], lags=40)
plot_pacf(sp500_clean['Log_RetURNS'], lags=40)
plt.show()
```

```
In [ ]: bitcoin_adf_test = adfuller(bitcoin_clean['Log_RetURNS'])
print('ADF Statistic: %f' % bitcoin_adf_test[0])
print('p-value: %f' % bitcoin_adf_test[1])

plot_acf(bitcoin_clean['Log_RetURNS'], lags=40)
plot_pacf(bitcoin_clean['Log_RetURNS'], lags=40)
plt.show()
```

Cross Validation

```
In [ ]: import matplotlib.dates as mdates
from datetime import datetime, timedelta
from dateutil.relativedelta import relativedelta

def create_sp500_cv_splits(data, start_date=None):

    if start_date is None:
        start_date = data.index.min()

    cv_splits = []
    window_start = start_date

    while True:
        # Define window boundaries
        train_start = window_start
        train_end = train_start + relativedelta(years=3) - timedelta(days=1)

        # Validation periods (8, 16, 24 months)
        val_start = train_end + timedelta(days=1)
        val1_end = val_start + relativedelta(months=8) - timedelta(days=1) # 8 months
        val2_end = val_start + relativedelta(months=16) - timedelta(days=1) # 16 months
        val3_end = val_start + relativedelta(months=24) - timedelta(days=1) # 24 months (2 years)

        # Test period (1 year)
        test_start = val3_end + timedelta(days=1)
        test_end = test_start + relativedelta(years=1) - timedelta(days=1)
```

```
# Check if we have enough data
if test_end.year > 2024:
    break

# Create splits for this window
train_data = data[(data.index >= train_start) & (data.index <= train_end)]

# Three validation folds
val1_data = data[(data.index >= val_start) & (data.index <= val1_end)]
val2_data = data[(data.index >= val_start) & (data.index <= val2_end)]
val3_data = data[(data.index >= val_start) & (data.index <= val3_end)]

test_data = data[(data.index >= test_start) & (data.index <= test_end)]

cv_splits.append({
    'window_id': len(cv_splits) + 1,
    'train': {
        'data': train_data,
        'start': train_start,
        'end': train_end,
        'size': len(train_data)
    },
    'validation': [
        {
            'fold': 1,
            'data': val1_data,
            'start': val_start,
            'end': val1_end,
            'size': len(val1_data),
            'months': 8
        },
        {
            'fold': 2,
            'data': val2_data,
            'start': val_start,
            'end': val2_end,
            'size': len(val2_data),
            'months': 16
        },
        {
            'fold': 3,
            'data': val3_data,
            'start': val_start,
            'end': val3_end,
            'size': len(val3_data)
        }
    ]
})
```

```

        'size': len(val3_data),
        'months': 24
    }
],
'test': {
    'data': test_data,
    'start': test_start,
    'end': test_end,
    'size': len(test_data)
}
})

# Move window forward by 1 year
window_start += relativedelta(years=1)

return cv_splits

def create_bitcoin_cv_splits(data, start_date=None):

    if start_date is None:
        # Start from a date that allows for proper window construction
        start_date = datetime(2015, 1, 1)

    cv_splits = []
    window_start = start_date

    # Define the testing period constraint
    test_period_start = datetime(2018, 1, 1)
    test_period_end = datetime(2023, 12, 31)

    while True:
        # Define window boundaries
        train_start = window_start
        train_end = train_start + relativedelta(years=2) - timedelta(days=1)

        # Validation periods (4, 8, 12 months)
        val_start = train_end + timedelta(days=1)
        val1_end = val_start + relativedelta(months=4) - timedelta(days=1) # 4 months
        val2_end = val_start + relativedelta(months=8) - timedelta(days=1) # 8 months
        val3_end = val_start + relativedelta(months=12) - timedelta(days=1) # 12 months

        # Test period (6 months)
        test_start = val3_end + timedelta(days=1)
        test_end = test_start + relativedelta(months=6) - timedelta(days=1)

```

```
# Check constraints
if test_end.year > 2023:
    break

# Only include windows where test period is within 2018-2023
if test_start < test_period_start:
    window_start += relativedelta(months=6)
    continue

# Create splits for this window
train_data = data[(data.index >= train_start) & (data.index <= train_end)]

# Three validation folds
val1_data = data[(data.index >= val_start) & (data.index <= val1_end)]
val2_data = data[(data.index >= val_start) & (data.index <= val2_end)]
val3_data = data[(data.index >= val_start) & (data.index <= val3_end)]

test_data = data[(data.index >= test_start) & (data.index <= test_end)]

cv_splits.append({
    'window_id': len(cv_splits) + 1,
    'train': {
        'data': train_data,
        'start': train_start,
        'end': train_end,
        'size': len(train_data)
    },
    'validation': [
        {
            'fold': 1,
            'data': val1_data,
            'start': val_start,
            'end': val1_end,
            'size': len(val1_data),
            'months': 4
        },
        {
            'fold': 2,
            'data': val2_data,
            'start': val_start,
            'end': val2_end,
            'size': len(val2_data),
            'months': 8
        },
    ],
})
```

```

        {
            'fold': 3,
            'data': val3_data,
            'start': val_start,
            'end': val3_end,
            'size': len(val3_data),
            'months': 12
        }
    ],
    'test': {
        'data': test_data,
        'start': test_start,
        'end': test_end,
        'size': len(test_data)
    }
}

# Move window forward by 6 months
window_start += relativedelta(months=6)

return cv_splits

```

```

In [ ]: # Apply Cross-Validation Schemes to Data

# Generate S&P 500 cross-validation splits
sp500_cv_splits = create_sp500_cv_splits(sp500_clean)

# Generate Bitcoin cross-validation splits
bitcoin_cv_splits = create_bitcoin_cv_splits(bitcoin_clean)

# Create summary DataFrames
def create_cv_summary(cv_splits, asset_name):
    summary_data = []

    for split in cv_splits:
        # Add training data info
        summary_data.append({
            'Asset': asset_name,
            'Window_ID': split['window_id'],
            'Split_Type': 'Train',
            'Fold': 'N/A',
            'Start_Date': split['train']['start'].strftime('%Y-%m-%d'),
            'End_Date': split['test']['end'].strftime('%Y-%m-%d')
        })

```

```

        'End_Date': split['train']['end'].strftime('%Y-%m-%d'),
        'Size': split['train']['size'],
        'Duration_Months': 'N/A'
    })

# Add validation data info
for val_fold in split['validation']:
    summary_data.append({
        'Asset': asset_name,
        'Window_ID': split['window_id'],
        'Split_Type': 'Validation',
        'Fold': val_fold['fold'],
        'Start_Date': val_fold['start'].strftime('%Y-%m-%d'),
        'End_Date': val_fold['end'].strftime('%Y-%m-%d'),
        'Size': val_fold['size'],
        'Duration_Months': val_fold['months']
    })

# Add test data info
summary_data.append({
    'Asset': asset_name,
    'Window_ID': split['window_id'],
    'Split_Type': 'Test',
    'Fold': 'N/A',
    'Start_Date': split['test']['start'].strftime('%Y-%m-%d'),
    'End_Date': split['test']['end'].strftime('%Y-%m-%d'),
    'Size': split['test']['size'],
    'Duration_Months': 'N/A'
})

return pd.DataFrame(summary_data)

# Create summary DataFrames
sp500_cv_summary = create_cv_summary(sp500_cv_splits, 'S&P_500')
bitcoin_cv_summary = create_cv_summary(bitcoin_cv_splits, 'Bitcoin')

# Combined summary
cv_summary_combined = pd.concat([sp500_cv_summary, bitcoin_cv_summary], ignore_index=True)

# Display first few windows for each asset
print("\nS&P 500 CV WINDOWS (First set)")
for i, split in enumerate(sp500_cv_splits[:1]):
    print(f"\nWindow {split['window_id']}:")
    print(f"  Train: {split['train']['start'].strftime('%Y-%m-%d')} to {split['train']['end'].strftime('%Y-%m-%d')} ({split['tra")
    print(f"  Validation Folds:")

```

```

for val_fold in split['validation']:
    print(f"    Fold {val_fold['fold']} ({val_fold['months']}mo): {val_fold['start'].strftime('%Y-%m-%d')} to {val_fold['end'].strftime('%Y-%m-%d')}")
print(f"    Test: {split['test']['start'].strftime('%Y-%m-%d')} to {split['test']['end'].strftime('%Y-%m-%d')} ({split['test'].shape[0]})")

print("\nBITCOIN CV WINDOWS (First set)")
for i, split in enumerate(bitcoin_cv_splits[:1]):
    print(f"\nWindow {split['window_id']}:")
    print(f"    Train: {split['train']['start'].strftime('%Y-%m-%d')} to {split['train']['end'].strftime('%Y-%m-%d')} ({split['train'].shape[0]})")
    print(f"    Validation Folds:")
    for val_fold in split['validation']:
        print(f"        Fold {val_fold['fold']} ({val_fold['months']}mo): {val_fold['start'].strftime('%Y-%m-%d')} to {val_fold['end'].strftime('%Y-%m-%d')}")
    print(f"    Test: {split['test']['start'].strftime('%Y-%m-%d')} to {split['test']['end'].strftime('%Y-%m-%d')} ({split['test'].shape[0]})")

# Display summary statistics
print("\nCV SPLIT STATISTICS")
split_stats = cv_summary_combined.groupby(['Asset', 'Split_Type']).agg({
    'Size': ['mean', 'std', 'min', 'max'],
    'Window_ID': 'count'
}).round(0)
print(split_stats)

```

In []: # Visualize Cross-Validation Scheme

```

def plot_cv_timeline(cv_splits, asset_name, max_windows=8):

    fig, ax = plt.subplots(figsize=(16, max(6, len(cv_splits[:max_windows]) * 1.5)))

    # Colors for different split types
    colors = {
        'train': '#2E8B57',
        'val_fold1': '#4169E1',
        'val_fold2': '#1E90FF',
        'val_fold3': '#87CEEB',
        'test': '#DC143C'
    }

    y_positions = []

    for i, split in enumerate(cv_splits[:max_windows]):
        y_pos = len(cv_splits[:max_windows]) - i - 1
        y_positions.append(y_pos)

        # Plot training period
        ax.bars(y_pos, (split['train']['end'] - split['train']['start']).days,

```

```

        left=split['train']['start'], height=0.6,
        color=colors['train'], alpha=0.8, label='Train' if i == 0 else ""))
    # Plot validation periods
    val_colors = ['val_fold1', 'val_fold2', 'val_fold3']
    for j, val_fold in enumerate(split['validation']):
        ax.barh(y_pos + 0.1 + j*0.15, (val_fold['end'] - val_fold['start']).days,
                left=val_fold['start'], height=0.12,
                color=colors[val_colors[j]], alpha=0.8,
                label=f'Val Fold {j+1} ({val_fold["months"]}mo)' if i == 0 else "")
    # Plot test period
    ax.barh(y_pos, (split['test']['end'] - split['test']['start']).days,
            left=split['test']['start'], height=0.6,
            color=colors['test'], alpha=0.8, label='Test' if i == 0 else "")
    # Add window labels
    ax.text(split['train']['start'], y_pos, f'W{split["window_id"]}',
            verticalalignment='center', fontsize=9, fontweight='bold')
    # Formatting
    ax.set_xlim(-0.5, len(cv_splits[:max_windows]) - 0.5)
    ax.set_ylabel('CV Windows (Newest to Oldest)', fontsize=12)
    ax.set_xlabel('Time Period', fontsize=12)
    ax.set_title(f'{asset_name} Cross-Validation Timeline\n{len(cv_splits)} Total Windows, Showing First {min(max_windows, len(cv_splits))}', fontsize=14, fontweight='bold')
    # Format x-axis
    ax.xaxis.set_major_formatter(mdates.DateFormatter('%Y-%m'))
    ax.xaxis.set_major_locator(mdates.YearLocator())
    plt.xticks(rotation=45)
    # Add legend
    ax.legend(loc='upper right', bbox_to_anchor=(1, 1), frameon=True, fancybox=True, shadow=True)
    # Add grid
    ax.grid(True, alpha=0.3, axis='x')
    plt.tight_layout()
    return fig, ax
def create_cv_comparison_chart():
    fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(20, 8))

```

```

# S&P 500 scheme visualization
y_pos = 1

# S&P 500 scheme
ax1.barh(y_pos, 3*365, left=0, height=0.6, color="#2E8B57", alpha=0.8, label='Train (3yr)')
ax1.barh(y_pos+0.1, 8*30, left=3*365, height=0.15, color="#4169E1", alpha=0.8, label='Val Fold 1 (8mo)')
ax1.barh(y_pos+0.25, 16*30, left=3*365, height=0.15, color="#1E90FF", alpha=0.8, label='Val Fold 2 (16mo)')
ax1.barh(y_pos+0.4, 24*30, left=3*365, height=0.15, color="#87CEEB", alpha=0.8, label='Val Fold 3 (24mo)')
ax1.barh(y_pos, 1*365, left=5*365, height=0.6, color="#DC143C", alpha=0.8, label='Test (1yr)')

ax1.set_xlim(0, 6*365)
ax1.set_ylim(0.5, 1.8)
ax1.set_xlabel('Days')
ax1.set_title('S&P 500 CV Scheme\n(6-year windows)', fontweight='bold')
ax1.legend()
ax1.grid(True, alpha=0.3)

# Bitcoin scheme
ax2.barh(y_pos, 2*365, left=0, height=0.6, color="#2E8B57", alpha=0.8, label='Train (2yr)')
ax2.barh(y_pos+0.1, 4*30, left=2*365, height=0.15, color="#4169E1", alpha=0.8, label='Val Fold 1 (4mo)')
ax2.barh(y_pos+0.25, 8*30, left=2*365, height=0.15, color="#1E90FF", alpha=0.8, label='Val Fold 2 (8mo)')
ax2.barh(y_pos+0.4, 12*30, left=2*365, height=0.15, color="#87CEEB", alpha=0.8, label='Val Fold 3 (12mo)')
ax2.barh(y_pos, 6*30, left=3*365, height=0.6, color="#DC143C", alpha=0.8, label='Test (6mo)')

ax2.set_xlim(0, 3.5*365)
ax2.set_ylim(0.5, 1.8)
ax2.set_xlabel('Days')
ax2.set_title('Bitcoin CV Scheme\n(~3.5-year windows)', fontweight='bold')
ax2.legend()
ax2.grid(True, alpha=0.3)

plt.suptitle('Time Series Cross-Validation Scheme Comparison', fontsize=16, fontweight='bold')
plt.tight_layout()
return fig

# fig3 = create_cv_comparison_chart()
# plt.show()

# Create summary table
cv_scheme_summary = pd.DataFrame({
    'Asset': ['S&P 500', 'Bitcoin'],
    'Total_Windows': [len(sp500_cv_splits), len(bitcoin_cv_splits)],
    'Window_Length': ['6 years', '~3.5 years'],
    'Train_Period': ['3 years', '2 years'],
}

```

```
'Validation_Folds': ['8/16/24 months', '4/8/12 months'],
'Test_Period': ['1 year', '6 months'],
'Window_Shift': ['1 year', '6 months'],
'Test_Coverage': [
    f'{sp500_cv_splits[0]['test']]['start'].strftime('%Y-%m-%d')} to {sp500_cv_splits[-1]['test']]['end'].strftime('%Y-%m-%d')
    f'{bitcoin_cv_splits[0]['test']]['start'].strftime('%Y-%m-%d')} to {bitcoin_cv_splits[-1]['test']]['end'].strftime('%Y-%m-%d')
]
})
```

In []: # Utility Functions for Cross-Validation Data Access

```
def get_cv_data(cv_splits, window_id, fold=None, return_type='data'):

    split = next((s for s in cv_splits if s['window_id'] == window_id), None)
    if split is None:
        raise ValueError(f"Window ID {window_id} not found")

    if return_type == 'train':
        return split['train']['data']
    elif return_type == 'test':
        return split['test']['data']
    elif return_type == 'validation':
        if fold is None:
            raise ValueError("Fold number must be specified for validation data")
        if fold not in [1, 2, 3]:
            raise ValueError("Fold must be 1, 2, or 3")
        return split['validation'][fold-1]['data']
    else:
        return split

# Example usage functions
def demonstrate_cv_usage():

    # Example 1: Get training data from first S&P 500 window
    train_data_sp500 = get_cv_data(sp500_cv_splits, window_id=1, return_type='train')
    print(f"S&P 500 Window 1 - Training data shape: {train_data_sp500.shape}")
    print(f"Training period: {train_data_sp500.index.min()} to {train_data_sp500.index.max()}\n\n"

    # Example 2: Get validation fold 2 data from first S&P 500 window
    val_data_sp500 = get_cv_data(sp500_cv_splits, window_id=1, fold=2, return_type='validation')
    print(f"\nS&P 500 Window 1 - Validation Fold 2 shape: {val_data_sp500.shape}")
    print(f"Validation period: {val_data_sp500.index.min()} to {val_data_sp500.index.max()}\n")
```

```

# Example 3: Get test data from first Bitcoin window
test_data_bitcoin = get_cv_data(bitcoin_cv_splits, window_id=1, return_type='test')
print(f"\nBitcoin Window 1 - Test data shape: {test_data_bitcoin.shape}")
print(f"Test period: {test_data_bitcoin.index.min()} to {test_data_bitcoin.index.max()}")


return train_data_sp500, val_data_sp500, test_data_bitcoin

# Run demonstration
sample_train, sample_val, sample_test = demonstrate_cv_usage()

# Save CV splits for later use (optional)
cv_implementation_summary = {
    'sp500_cv_splits': sp500_cv_splits,
    'bitcoin_cv_splits': bitcoin_cv_splits,
    'sp500_summary': sp500_cv_summary,
    'bitcoin_summary': bitcoin_cv_summary,
    'combined_summary': cv_summary_combined,
    'scheme_comparison': cv_scheme_summary
}

```

ARIMA + LSTM HYBRID

ARIMA

```

In [ ]: import itertools
from statsmodels.tsa.arima.model import ARIMA
from statsmodels.tsa.stattools import adfuller, kpss
from statsmodels.stats.diagnostic import acorr_ljungbox
import warnings
from sklearn.metrics import mean_squared_error, mean_absolute_error
from scipy import stats
import time

# Suppress convergence warnings for cleaner output
warnings.filterwarnings('ignore', category=UserWarning)
warnings.filterwarnings('ignore', category=RuntimeWarning)

def find_optimal_arima_order(data, max_p=5, max_d=2, max_q=5, seasonal=False,
                             information_criterion='aic', verbose=False):

```

```
best_ic = np.inf
best_params = None
best_model = None
results_log = []

# Create parameter grid
if seasonal:
    # For seasonal ARIMA (not implemented in this study)
    param_grid = itertools.product(range(max_p+1), range(max_d+1), range(max_q+1),
                                    range(2), range(2), range(2), [12])
else:
    # Standard ARIMA grid search
    param_grid = itertools.product(range(max_p+1), range(max_d+1), range(max_q+1))

total_combinations = (max_p+1) * (max_d+1) * (max_q+1)

if verbose:
    print(f"Testing {total_combinations} ARIMA parameter combinations...")
    start_time = time.time()

for i, params in enumerate(param_grid):
    p, d, q = params[:3]

    # Skip if model is too simple (all parameters zero)
    if p == 0 and d == 0 and q == 0:
        continue

    try:
        # Fit ARIMA model
        model = ARIMA(data, order=(p, d, q))
        fitted_model = model.fit()

        # Get information criterion value
        if information_criterion.lower() == 'aic':
            ic_value = fitted_model.aic
        elif information_criterion.lower() == 'bic':
            ic_value = fitted_model.bic
        elif information_criterion.lower() == 'hqic':
            ic_value = fitted_model.hqic
        else:
            ic_value = fitted_model.aic

        # Store results
        results_log.append({
            'order': (p, d, q),
```

```

        'aic': fitted_model.aic,
        'bic': fitted_model.bic,
        'hqic': fitted_model.hqic,
        'llf': fitted_model.llf,
        'converged': fitted_model.mle_retvals['converged'] if hasattr(fitted_model, 'mle_retvals') else True
    })

    # Update best model if current is better
    if ic_value < best_ic:
        best_ic = ic_value
        best_params = (p, d, q)
        best_model = fitted_model

except Exception as e:
    # Log failed fits
    results_log.append({
        'order': (p, d, q),
        'aic': np.nan,
        'bic': np.nan,
        'hqic': np.nan,
        'llf': np.nan,
        'converged': False,
        'error': str(e)
    })

    if verbose and i % 10 == 0:
        print(f"Failed to fit ARIMA{params}: {str(e)[:50]}...")

    if verbose and (i + 1) % 20 == 0:
        elapsed = time.time() - start_time
        progress = (i + 1) / total_combinations * 100
        print(f"Progress: {progress:.1f}% ({i+1}/{total_combinations}) | "
              f"Best so far: ARIMA{best_params} ({information_criterion.upper()}={best_ic:.4f})")

if verbose:
    total_time = time.time() - start_time
    print(f"\nGrid search completed in {total_time:.2f} seconds")
    print(f"Best model: ARIMA{best_params} with {information_criterion.upper()}={best_ic:.4f}")

# Create results summary
results_df = pd.DataFrame(results_log)
successful_fits = results_df[results_df['converged'] == True]

return {
    'best_order': best_params,

```

```

        'best_model': best_model,
        'best_ic_value': best_ic,
        'information_criterion': information_criterion,
        'results_df': results_df,
        'successful_fits': len(successful_fits),
        'total_attempts': len(results_log),
        'success_rate': len(successful_fits) / len(results_log) * 100
    }

def evaluate_arima_model(model, train_data, test_data, model_order):

    # Generate forecasts
    n_forecast = len(test_data)
    forecast_result = model.get_forecast(steps=n_forecast)
    forecasts = forecast_result.predicted_mean
    forecast_ci = forecast_result.conf_int()

    # Calculate performance metrics
    mse = mean_squared_error(test_data, forecasts)
    rmse = np.sqrt(mse)
    mae = mean_absolute_error(test_data, forecasts)
    mape = np.mean(np.abs((test_data - forecasts) / test_data)) * 100

    # Calculate R2 score for fair comparison with LSTM and SVM
    ss_res = np.sum((test_data - forecasts) ** 2)
    ss_tot = np.sum((test_data - np.mean(test_data)) ** 2)
    r2 = 1 - (ss_res / ss_tot) if ss_tot > 0 else 0

    # Direction accuracy (for returns)
    direction_actual = np.sign(test_data.values[1:])
    direction_forecast = np.sign(forecasts.values[1:])
    direction_accuracy = np.mean(direction_actual == direction_forecast) * 100

    # Residual diagnostics
    residuals = model.resid

    # Ljung-Box test for serial correlation in residuals
    lb_test = acorr_ljungbox(residuals, lags=10, return_df=False)

    # Normality test (Jarque-Bera)
    jb_stat, jb_pvalue = stats.jarque_bera(residuals)

    # Heteroskedasticity test (simple approach)
    residuals_squared = residuals ** 2
    arch_stat, arch_pvalue = acorr_ljungbox(residuals_squared, lags=5, return_df=False)

```

```

    return {
        'model_order': model_order,
        'forecasts': forecasts,
        'forecast_ci': forecast_ci,
        'performance_metrics': {
            'mse': mse,
            'rmse': rmse,
            'mae': mae,
            'mape': mape,
            'r2': r2,
            'direction_accuracy': direction_accuracy
        },
        'diagnostic_tests': {
            'ljung_box_stat': lb_test['lb_stat'].iloc[-1],
            'ljung_box_pvalue': lb_test['lb_pvalue'].iloc[-1],
            'jarque_bera_stat': jb_stat,
            'jarque_bera_pvalue': jb_pvalue,
            'arch_stat': arch_stat[-1] if isinstance(arch_stat, np.ndarray) else arch_stat,
            'arch_pvalue': arch_pvalue[-1] if isinstance(arch_pvalue, np.ndarray) else arch_pvalue
        },
        'residuals': residuals
    }

```

HYBRID ARIMA-LSTM MODEL

```

In [ ]: def run_arima_cross_validation(cv_splits, data_clean, asset_name, max_p=3, max_d=2, max_q=3,
                                     information_criterion='aic', verbose=True):

    print(f"\n{asset_name.upper()} ARIMA CROSS-VALIDATION ===")
    print(f"Running AIC-based model selection across {len(cv_splits)} windows...")
    print(f"Parameter search space: p∈[0,{max_p}], d∈[0,{max_d}], q∈[0,{max_q}]")
    print("-" * 80)

    all_results = []
    model_selection_summary = []

    for window_idx, split in enumerate(cv_splits):
        window_id = split['window_id']

        if verbose:
            print(f"\nProcessing Window {window_id}/{len(cv_splits)}...")
            print(f"    Train: {split['train']['start'].strftime('%Y-%m-%d')} to {split['train']['end'].strftime('%Y-%m-%d')}\t({s

```

```
print(f"  Test: {split['test']['start'].strftime('%Y-%m-%d')} to {split['test']['end'].strftime('%Y-%m-%d')}) ({spl

# Extract data
train_data = split['train']['data']['Log_Returns']
test_data = split['test']['data']['Log_Returns']

# STEP 1: Model Selection using Training Data
if verbose:
    print(f"  Model selection using {information_criterion.upper()} criterion...")

selection_result = find_optimal_arima_order(
    train_data,
    max_p=max_p,
    max_d=max_d,
    max_q=max_q,
    information_criterion=information_criterion,
    verbose=False # Keep individual window selection quiet
)

if selection_result['best_model'] is None:
    print(f"  Failed to find suitable model for Window {window_id}")
    continue

best_order = selection_result['best_order']

# STEP 2: Hyperparameter Validation using Validation Folds
if verbose:
    print(f"  📈 Validating ARIMA{best_order} across 3 validation folds...")

validation_scores = []

for val_fold in split['validation']:
    fold_num = val_fold['fold']
    val_data = val_fold['data']['Log_Returns']

    try:
        # Fit model on training data and evaluate on validation fold
        val_model = ARIMA(train_data, order=best_order).fit()
        val_forecasts = val_model.get_forecast(steps=len(val_data)).predicted_mean
        val_rmse = np.sqrt(mean_squared_error(val_data, val_forecasts))
        validation_scores.append(val_rmse)

    except Exception as e:
        if verbose:
            print(f"      ⚠️ Validation fold {fold_num} failed: {str(e)[:50]}...")
```

```

        validation_scores.append(np.inf)

    avg_validation_rmse = np.mean(validation_scores)

    # STEP 3: Final Model Training and Out-of-Sample Evaluation
    if verbose:
        print(f"  ⚡ Final evaluation on test data...")

    try:
        # Re-fit the model on training data
        final_model = ARIMA(train_data, order=best_order).fit()

        # Evaluate on test data
        evaluation = evaluate_arima_model(final_model, train_data, test_data, best_order)

        # Store comprehensive results
        window_result = {
            'window_id': window_id,
            'asset': asset_name,
            'train_period': f'{split["train"]["start"].strftime("%Y-%m-%d")} to {split["train"]["end"].strftime("%Y-%m-%d")}',
            'test_period': f'{split["test"]["start"].strftime("%Y-%m-%d")} to {split["test"]["end"].strftime("%Y-%m-%d")}',
            'train_size': split['train']['size'],
            'test_size': split['test']['size'],
            'best_order': best_order,
            'model_selection': selection_result,
            'validation_scores': validation_scores,
            'avg_validation_rmse': avg_validation_rmse,
            'evaluation': evaluation,
            'final_model': final_model
        }

        all_results.append(window_result)

        # Summary for quick reference
        model_selection_summary.append({
            'Window': window_id,
            'Best_Order': f"ARIMA{best_order}",
            'AIC': selection_result['best_ic_value'],
            'Validation_RMSE': avg_validation_rmse,
            'Test_RMSE': evaluation['performance_metrics']['rmse'],
            'Test_MAE': evaluation['performance_metrics']['mae'],
            'Test_R2': evaluation['performance_metrics']['r2'],
            'Direction_Accuracy': evaluation['performance_metrics']['direction_accuracy'],
            'Ljung_Box_p': evaluation['diagnostic_tests']['ljung_box_pvalue']
        })
    
```

```

    if verbose:
        print(f"  ARIMA{best_order}: Test RMSE={evaluation['performance_metrics']['rmse']:.6f}, "
              f"Direction Acc={evaluation['performance_metrics']['direction_accuracy']:.1f}%")

    except Exception as e:
        print(f"    Final evaluation failed for Window {window_id}: {str(e)}")
        continue

    # Create summary DataFrame
    summary_df = pd.DataFrame(model_selection_summary)

    # Calculate overall performance statistics
    if len(summary_df) > 0:
        performance_summary = {
            'total_windows': len(cv_splits),
            'successful_windows': len(summary_df),
            'success_rate': len(summary_df) / len(cv_splits) * 100,
            'avg_test_rmse': summary_df['Test_RMSE'].mean(),
            'std_test_rmse': summary_df['Test_RMSE'].std(),
            'avg_test_mae': summary_df['Test_MAE'].mean(),
            'avg_r2': summary_df['Test_R2'].mean(),
            'avg_direction_accuracy': summary_df['Direction_Accuracy'].mean(),
            'avg_validation_rmse': summary_df['Validation_RMSE'].mean(),
            'most_common_order': summary_df['Best_Order'].mode().iloc[0] if len(summary_df) > 0 else None
        }
    else:
        performance_summary = None

    print(f"\n{'='*80}")
    print(f"{asset_name.upper()} ARIMA CROSS-VALIDATION COMPLETE")
    print(f"{'='*80}")

    if performance_summary:
        print(f" Successfully processed {performance_summary['successful_windows']}/{performance_summary['total_windows']} windows")
        print(f" Average Test RMSE: {performance_summary['avg_test_rmse']:.6f} ± {performance_summary['std_test_rmse']:.6f}")
        print(f" Average Direction Accuracy: {performance_summary['avg_direction_accuracy']:.2f}%")
        print(f" Most Common Model: {performance_summary['most_common_order']}")
    else:
        print(" No successful model fits achieved")

    return {
        'asset_name': asset_name,
        'all_results': all_results,
        'summary_df': summary_df,
    }

```

```
'performance_summary': performance_summary,
'methodology': {
    'approach': 'AIC-based automated selection',
    'information_criterion': information_criterion,
    'parameter_space': f'p∈[0,{max_p}], d∈[0,{max_d}], q∈[0,{max_q}]',
    'cross_validation': '3-fold temporal validation',
    'evaluation_metric': 'Out-of-sample RMSE and direction accuracy'
}
}
```

```
In [ ]: import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense, Dropout
from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import mean_squared_error, mean_absolute_error

def create_sequences_from_residuals(residuals, lookback=60):

    X, y = [], []
    for i in range(lookback, len(residuals)):
        X.append(residuals[i-lookback:i])
        y.append(residuals[i])
    return np.array(X), np.array(y)

def build_residual_lstm_model(lookback=60, units=50, dropout=0.2, learning_rate=0.001):

    model = Sequential([
        LSTM(units=units, return_sequences=True, input_shape=(lookback, 1)),
        Dropout(dropout),
        LSTM(units=units, return_sequences=False),
        # Dense(16, activation='relu'),
        Dropout(dropout),
        Dense(units=1)
    ])

    optimizer = tf.keras.optimizers.Adam(learning_rate=learning_rate)
    model.compile(optimizer=optimizer, loss='mean_squared_error', metrics=['mae'])

    return model

def run_hybrid_arima_lstm_cv(cv_splits, data_clean, asset_name,
```

```

max_p=3, max_d=1, max_q=3,
lookback=60, lstm_units=50, dropout=0.2,
epochs=50, batch_size=32, verbose=True):

print(f"\n{'='*100}")
print(f" HYBRID ARIMA-LSTM CROSS-VALIDATION: {asset_name.upper()}")
print(f"{'='*100}")
print(f"Architecture: ARIMA (linear) + LSTM on residuals (non-linear)")
print(f"ARIMA space: p∈[0,{max_p}], d∈[0,{max_d}], q∈[0,{max_q}]")
print(f"LSTM config: lookback={lookback}, units={lstm_units}, dropout={dropout}")
print(f"Total windows: {len(cv_splits)}")
print("-" * 100)

all_results = []
scaler = MinMaxScaler(feature_range=(0, 1))

for window_idx, split in enumerate(cv_splits):
    window_id = split['window_id']

    if verbose:
        print(f"\n Window {window_id}/{len(cv_splits)}")
        print(f"  Train: {split['train']['start'].strftime('%Y-%m-%d')} to {split['train']['end'].strftime('%Y-%m-%d')}")
        print(f"  Test:  {split['test']['start'].strftime('%Y-%m-%d')} to {split['test']['end'].strftime('%Y-%m-%d')}")
        print(f"{'='*100}")

    try:
        # Extract data
        train_data = split['train']['data']['Log_Returns'].values
        test_data = split['test']['data']['Log_Returns'].values

        if verbose:
            print(f"  [1/4] ↗ Fitting ARIMA model...")

        # Find optimal ARIMA order
        from statsmodels.tsa.arima.model import ARIMA
        best_aic = np.inf
        best_order = None
        best_model = None

        for p in range(max_p + 1):
            for d in range(max_d + 1):
                for q in range(max_q + 1):
                    try:
                        model = ARIMA(train_data, order=(p, d, q))
                        fitted_model = model.fit()
                    except:
                        continue
                    aic = fitted_model.aic
                    if aic < best_aic:
                        best_aic = aic
                        best_order = (p, d, q)
                        best_model = model
```
```

```

 if fitted_model.aic < best_aic:
 best_aic = fitted_model.aic
 best_order = (p, d, q)
 best_model = fitted_model
 except:
 continue

 if best_model is None:
 print(f" Failed to fit ARIMA for window {window_id}")
 continue

 if verbose:
 print(f" Best ARIMA{best_order}, AIC={best_aic:.2f}")

 if verbose:
 print(f" [2/4] Calculating ARIMA residuals...")

Get in-sample predictions and residuals
arima_train_pred = best_model.fittedvalues
train_residuals = train_data - arima_train_pred

Align arrays (ARIMA might drop initial values)
if len(train_residuals) < len(train_data):
 n_dropped = len(train_data) - len(train_residuals)
 train_data_aligned = train_data[n_dropped:]
 train_residuals_aligned = train_residuals
else:
 train_data_aligned = train_data
 train_residuals_aligned = train_residuals

if verbose:
 print(f" ✓ Residuals: mean={np.mean(train_residuals_aligned):.6f}, std={np.std(train_residuals_aligned):.6f}")

if len(train_residuals_aligned) <= lookback:
 print(f" ⚠️ Insufficient data for LSTM (need > {lookback} points)")
 continue

if verbose:
 print(f" [3/4] 🧠 Training LSTM on residuals...")

Scale residuals
residuals_reshaped = train_residuals_aligned.reshape(-1, 1)
scaled_residuals = scaler.fit_transform(residuals_reshaped)

Create sequences from residuals

```

```

X_train, y_train = create_sequences_from_residuals(scaled_residuals, lookback)

if len(X_train) == 0:
 print(f" No sequences created from residuals")
 continue

X_train = X_train.reshape((X_train.shape[0], X_train.shape[1], 1))

Build and train LSTM
lstm_model = build_residual_lstm_model(lookback=lookback, units=lstm_units, dropout=dropout)

early_stop = EarlyStopping(monitor='loss', patience=5, restore_best_weights=True, verbose=0)
reduce_lr = ReduceLROnPlateau(monitor='loss', factor=0.5, patience=3, min_lr=1e-6, verbose=0)

history = lstm_model.fit(
 X_train, y_train,
 epochs=epochs,
 batch_size=batch_size,
 callbacks=[early_stop, reduce_lr],
 verbose=0,
 validation_split=0.1
)

if verbose:
 print(f" ✓ LSTM trained ({len(history.history['loss'])} epochs)")

if verbose:
 print(f" [4/4] ⚡ Generating hybrid forecasts...")

Make ARIMA forecasts for test period
arima_test_pred = best_model.forecast(steps=len(test_data))

For LSTM residual prediction, we need historical residuals
Use last 'lookback' residuals from training + generate iteratively for test

Get last lookback residuals from training
last_residuals = scaled_residuals[-lookback:].flatten()
lstm_residual_predictions = []

Predict residuals for each test point
for t in range(len(test_data)):
 # Prepare input sequence
 X_input = last_residuals[-lookback:].reshape(1, lookback, 1)

 # Predict next residual

```

```

residual_pred_scaled = lstm_model.predict(X_input, verbose=0)
residual_pred = scaler.inverse_transform(residual_pred_scaled)[0, 0]
lstm_residual_predictions.append(residual_pred)

Update sequence with actual residual for next iteration
if t < len(test_data) - 1:
 actual_residual = test_data[t] - arima_test_pred[t]
 actual_residual_scaled = scaler.transform([[actual_residual]])[0, 0]
 last_residuals = np.append(last_residuals[1:], actual_residual_scaled)

lstm_residual_predictions = np.array(lstm_residual_predictions)

HYBRID PREDICTION = ARIMA + LSTM_residuals
hybrid_predictions = arima_test_pred + lstm_residual_predictions

Metrics for ARIMA only
arima_rmse = np.sqrt(mean_squared_error(test_data, arima_test_pred))
arima_mae = mean_absolute_error(test_data, arima_test_pred)

Metrics for HYBRID
hybrid_rmse = np.sqrt(mean_squared_error(test_data, hybrid_predictions))
hybrid_mae = mean_absolute_error(test_data, hybrid_predictions)

Direction accuracy
arima_direction = np.mean(np.sign(test_data[1:]) == np.sign(arima_test_pred[1:])) * 100
hybrid_direction = np.mean(np.sign(test_data[1:]) == np.sign(hybrid_predictions[1:])) * 100

improvement = ((arima_rmse - hybrid_rmse) / arima_rmse) * 100

if verbose:
 print(f" RESULTS:")
 print(f" ARIMA only: RMSE={arima_rmse:.6f}, MAE={arima_mae:.6f}, Direction={arima_direction:.1f}%")
 print(f" HYBRID: RMSE={hybrid_rmse:.6f}, MAE={hybrid_mae:.6f}, Direction={hybrid_direction:.1f}%")
 print(f" Improvement: {improvement:+.2f}% RMSE")

Store results
window_results = {
 'window_id': window_id,
 'train_start': split['train']['start'],
 'train_end': split['train']['end'],
 'test_start': split['test']['start'],
 'test_end': split['test']['end'],
 'arima_order': best_order,
 'arima_aic': best_aic,
 # ARIMA metrics
}

```

```

 'arima_rmse': arima_rmse,
 'arima_mae': arima_mae,
 'arima_direction_accuracy': arima_direction,
 'arima_predictions': arima_test_pred,
 # LSTM residual metrics
 'lstm_residual_predictions': lstm_residual_predictions,
 # HYBRID metrics
 'hybrid_rmse': hybrid_rmse,
 'hybrid_mae': hybrid_mae,
 'hybrid_direction_accuracy': hybrid_direction,
 'hybrid_predictions': hybrid_predictions,
 # Comparison
 'rmse_improvement_pct': improvement,
 'actuals': test_data
 }

 all_results.append(window_results)

except Exception as e:
 print(f" Error in window {window_id}: {str(e)}")
 import traceback
 traceback.print_exc()
 continue

if len(all_results) > 0:
 avg_arima_rmse = np.mean([r['arima_rmse'] for r in all_results])
 avg_hybrid_rmse = np.mean([r['hybrid_rmse'] for r in all_results])
 avg_improvement = np.mean([r['rmse_improvement_pct'] for r in all_results])

 avg_arima_mae = np.mean([r['arima_mae'] for r in all_results])
 avg_hybrid_mae = np.mean([r['hybrid_mae'] for r in all_results])

 avg_arima_direction = np.mean([r['arima_direction_accuracy'] for r in all_results])
 avg_hybrid_direction = np.mean([r['hybrid_direction_accuracy'] for r in all_results])

 print("\n'*100")
 print(f'{asset_name.upper()} HYBRID MODEL COMPLETE')
 print('*100")
 print(f"Windows processed: {len(all_results)}/{len(cv_splits)}")
 print("\n AVERAGE PERFORMANCE:")
 print(f" ARIMA only: RMSE={avg_arima_rmse:.6f}, MAE={avg_arima_mae:.6f}, Direction={avg_arima_direction:.2f}%")
 print(f" HYBRID: RMSE={avg_hybrid_rmse:.6f}, MAE={avg_hybrid_mae:.6f}, Direction={avg_hybrid_direction:.2f}%")
 print(f" Improvement: {avg_improvement:+.2f}% RMSE")
 print('*100")

```

```

results_dict = {
 'asset_name': asset_name,
 'model_type': 'HYBRID_ARIMA_LSTM',
 'windows_processed': len(all_results),
 'total_windows': len(cv_splits),
 # ARIMA metrics
 'avg_arima_rmse': avg_arima_rmse,
 'avg_arima_direction': avg_arima_direction,
 'avg_arima_mae': avg_arima_mae,
 # HYBRID metrics
 'avg_hybrid_rmse': avg_hybrid_rmse,
 'avg_hybrid_direction': avg_hybrid_direction,
 'avg_hybrid_mae': avg_hybrid_mae,
 # Improvement
 'avg_improvement_pct': avg_improvement,
 'window_results': all_results,
 'hyperparameters': {
 'arima': f'p∈[0,{max_p}], d∈[0,{max_d}], q∈[0,{max_q}]',
 'lstm_lookback': lookback,
 'lstm_units': lstm_units,
 'dropout': dropout,
 'epochs': epochs,
 'batch_size': batch_size
 }
}

return results_dict
else:
 print(f"\n No windows successfully processed for {asset_name}")
 return None

```

```

In []: print("\n PHASE 1: S&P 500 HYBRID ARIMA-LSTM")

sp500_hybrid_results = run_hybrid_arima_lstm_cv(
 cv_splits=sp500_cv_splits,
 data_clean=sp500_clean,
 asset_name='S&P 500',
 max_p=3,
 max_d=1,
 max_q=3,
 lookback=60,
 lstm_units=50,
 dropout=0.2,
 epochs=50,

```

```

 batch_size=32,
 verbose=True
)

print("\n\n PHASE 2: BITCOIN HYBRID ARIMA-LSTM")

bitcoin_hybrid_results = run_hybrid_arima_lstm_cv(
 cv_splits=bitcoin_cv_splits,
 data_clean=bitcoin_clean,
 asset_name='Bitcoin',
 max_p=3,
 max_d=1,
 max_q=3,
 lookback=60,
 lstm_units=50,
 dropout=0.2,
 epochs=50,
 batch_size=32,
 verbose=True
)

print("\n" + "=" * 100)
print(" FINAL SUMMARY: HYBRID MODEL PERFORMANCE ACROSS BOTH ASSETS")
print("=" * 100)

summary_data = []

if sp500_hybrid_results:
 summary_data.append({
 'Asset': 'S&P 500',
 'ARIMA_RMSE': sp500_hybrid_results['avg_arima_rmse'],
 'Hybrid_RMSE': sp500_hybrid_results['avg_hybrid_rmse'],
 'Improvement_%': sp500_hybrid_results['avg_improvement_pct'],
 'ARIMA_MAE': sp500_hybrid_results['avg_arima_mae'],
 'Hybrid_MAE': sp500_hybrid_results['avg_hybrid_mae'],
 'ARIMA_Direction_%': sp500_hybrid_results['avg_arima_direction'],
 'Hybrid_Direction_%': sp500_hybrid_results['avg_hybrid_direction'],
 'Windows': f"{sp500_hybrid_results['windows_processed']}/{sp500_hybrid_results['total_windows']}"
 })

if bitcoin_hybrid_results:
 summary_data.append({
 'Asset': 'Bitcoin',
 'ARIMA_RMSE': bitcoin_hybrid_results['avg_arima_rmse'],

```

```

'Hybrid_RMSE': bitcoin_hybrid_results['avg_hybrid_rmse'],
'Improvement_%': bitcoin_hybrid_results['avg_improvement_pct'],
'ARIMA_MAE': bitcoin_hybrid_results['avg_arima_mae'],
'Hybrid_MAE': bitcoin_hybrid_results['avg_hybrid_mae'],
'ARIMA_Direction_%': bitcoin_hybrid_results['avg_arima_direction'],
'Hybrid_Direction_%': bitcoin_hybrid_results['avg_hybrid_direction'],
'Windows': f"{bitcoin_hybrid_results['windows_processed']}/{bitcoin_hybrid_results['total_windows']}"
})

if summary_data:
 summary_df = pd.DataFrame(summary_data)
 print("\n", summary_df.to_string(index=False))

 print("\n" + "=" * 100)
 print("✓ KEY INSIGHTS:")
 print("=" * 100)

 for data in summary_data:
 asset = data['Asset']
 improvement = data['Improvement_%']

 if improvement > 0:
 print(f"✓ {asset}: Hybrid model OUTPERFORMS ARIMA by {improvement:.2f}% RMSE")
 elif improvement < 0:
 print(f"✗ {asset}: Hybrid model underperforms ARIMA by {abs(improvement):.2f}% RMSE")
 else:
 print(f"= {asset}: Hybrid model equals ARIMA performance")

```

```

In []: def plot_hybrid_decomposition(results, window_idx=0):
 if not results or len(results['window_results']) == 0:
 print("No results to plot")
 return

 window = results['window_results'][window_idx]
 asset_name = results['asset_name']

 # Extract data
 actuals = window['actuals']
 arima_pred = window['arima_predictions']
 lstm_residual = window['lstm_residual_predictions']
 hybrid_pred = window['hybrid_predictions']

 # Create figure with subplots
 fig, axes = plt.subplots(3, 2, figsize=(16, 12))

```

```

fig.suptitle(f'{asset_name} - Hybrid ARIMA-LSTM Decomposition (Window {window["window_id"]})',
 fontsize=16, fontweight='bold')

Plot 1: Actual vs ARIMA
ax1 = axes[0, 0]
ax1.plot(actuals, label='Actual', color='black', linewidth=2, alpha=0.7)
ax1.plot(arima_pred, label='ARIMA', color='blue', linewidth=1.5, linestyle='--', alpha=0.7)
ax1.set_title('ARIMA Component (Linear)', fontsize=12, fontweight='bold')
ax1.set_ylabel('Log Returns')
ax1.legend()
ax1.grid(True, alpha=0.3)
ax1.text(0.02, 0.98, f'RMSE: {window["arima_rmse"]:.6f}',
 transform=ax1.transAxes, verticalalignment='top',
 bbox=dict(boxstyle='round', facecolor='wheat', alpha=0.5))

Plot 2: ARIMA Residuals
ax2 = axes[0, 1]
arima_residuals = actuals - arima_pred
ax2.plot(arima_residuals, color='red', linewidth=1, alpha=0.7)
ax2.axhline(y=0, color='black', linestyle='-', linewidth=0.5)
ax2.fill_between(range(len(arima_residuals)), arima_residuals, 0, alpha=0.3, color='red')
ax2.set_title('ARIMA Residuals (Non-linear to Model)', fontsize=12, fontweight='bold')
ax2.set_ylabel('Residuals')
ax2.grid(True, alpha=0.3)
ax2.text(0.02, 0.98, f'Std: {np.std(arima_residuals):.6f}',
 transform=ax2.transAxes, verticalalignment='top',
 bbox=dict(boxstyle='round', facecolor='wheat', alpha=0.5))

Plot 3: LSTM Residual Predictions
ax3 = axes[1, 0]
ax3.plot(arima_residuals, label='Actual Residuals', color='red', linewidth=1.5, alpha=0.7)
ax3.plot(lstm_residual, label='LSTM Predicted', color='green', linewidth=1.5, linestyle='--', alpha=0.7)
ax3.axhline(y=0, color='black', linestyle='-', linewidth=0.5)
ax3.set_title('LSTM Residual Predictions', fontsize=12, fontweight='bold')
ax3.set_ylabel('Residuals')
ax3.legend()
ax3.grid(True, alpha=0.3)

Plot 4: Actual vs Hybrid
ax4 = axes[1, 1]
ax4.plot(actuals, label='Actual', color='black', linewidth=2, alpha=0.7)
ax4.plot(hybrid_pred, label='Hybrid (ARIMA+LSTM)', color='purple', linewidth=1.5, linestyle='--', alpha=0.7)
ax4.set_title('Final Hybrid Forecast', fontsize=12, fontweight='bold')
ax4.set_ylabel('Log Returns')
ax4.legend()

```

```

ax4.grid(True, alpha=0.3)
ax4.text(0.02, 0.98, f'RMSE: {window["hybrid_rmse"]:.6f}',
 transform=ax4.transAxes, verticalalignment='top',
 bbox=dict(boxstyle='round', facecolor='wheat', alpha=0.5))

Plot 5: Component Comparison
ax5 = axes[2, 0]
sample_size = min(50, len(actuals))
x = range(sample_size)
ax5.plot(x, actuals[:sample_size], label='Actual', color='black', linewidth=2, alpha=0.7)
ax5.plot(x, arima_pred[:sample_size], label='ARIMA', color='blue', linewidth=1.5, linestyle='--', alpha=0.7)
ax5.plot(x, hybrid_pred[:sample_size], label='Hybrid', color='purple', linewidth=1.5, linestyle='--', alpha=0.7)
ax5.set_title(f'First {sample_size} Predictions Comparison', fontsize=12, fontweight='bold')
ax5.set_xlabel('Time Step')
ax5.set_ylabel('Log Returns')
ax5.legend()
ax5.grid(True, alpha=0.3)

Plot 6: Performance Metrics
ax6 = axes[2, 1]
metrics = ['RMSE', 'MAE', 'Direction Acc (%)']
arima_metrics = [window['arima_rmse'], window['arima_mae'], window['arima_direction_accuracy']]
hybrid_metrics = [window['hybrid_rmse'], window['hybrid_mae'], window['hybrid_direction_accuracy']]

x_pos = np.arange(len(metrics))
width = 0.35

bars1 = ax6.bar(x_pos - width/2, arima_metrics, width, label='ARIMA', alpha=0.8)
bars2 = ax6.bar(x_pos + width/2, hybrid_metrics, width, label='Hybrid', alpha=0.8)

ax6.set_ylabel('Value')
ax6.set_title('Performance Metrics Comparison', fontsize=12, fontweight='bold')
ax6.set_xticks(x_pos)
ax6.set_xticklabels(metrics)
ax6.legend()
ax6.grid(True, alpha=0.3, axis='y')

Add value labels
for bars in [bars1, bars2]:
 for bar in bars:
 height = bar.get_height()
 ax6.text(bar.get_x() + bar.get_width()/2., height,
 f'{height:.4f}', ha='center', va='bottom', fontsize=8)

plt.tight_layout()

```

```

plt.show()

Print detailed stats
print(f"\n{'='*80}")
print(f"DETAILED STATISTICS - Window {window['window_id']}")
print(f"{'='*80}")
print(f"Period: {window['test_start'].strftime('%Y-%m-%d')} to {window['test_end'].strftime('%Y-%m-%d')}")
print(f"ARIMA Order: {window['arima_order']}, AIC: {window['arima_aic']:.2f}")
print(f"\nARIMA Performance:")
print(f" • RMSE: {window['arima_rmse']:.6f}")
print(f" • MAE: {window['arima_mae']:.6f}")
print(f" • Direction Accuracy: {window['arima_direction_accuracy']:.2f}%")
print(f"\nHybrid Performance:")
print(f" • RMSE: {window['hybrid_rmse']:.6f}")
print(f" • MAE: {window['hybrid_mae']:.6f}")
print(f" • Direction Accuracy: {window['hybrid_direction_accuracy']:.2f}%")
print(f"\nImprovement: {window['rmse_improvement_pct']+2} % RMSE")
print(f"{'='*80}")

def plot_all_windows_comparison(results):
 if not results or len(results['window_results']) == 0:
 print("No results to plot")
 return

 asset_name = results['asset_name']
 windows = [r['window_id'] for r in results['window_results']]
 arima_rmse = [r['arima_rmse'] for r in results['window_results']]
 hybrid_rmse = [r['hybrid_rmse'] for r in results['window_results']]
 improvement = [r['rmse_improvement_pct'] for r in results['window_results']]

 fig, axes = plt.subplots(1, 2, figsize=(14, 5))
 fig.suptitle(f'{asset_name} - Hybrid Model Performance Across All Windows',
 fontsize=14, fontweight='bold')

 # Plot 1: RMSE Comparison
 ax1 = axes[0]
 ax1.plot(windows, arima_rmse, marker='o', label='ARIMA', linewidth=2, markersize=6)
 ax1.plot(windows, hybrid_rmse, marker='s', label='Hybrid', linewidth=2, markersize=6)
 ax1.set_xlabel('Window ID')
 ax1.set_ylabel('RMSE')
 ax1.set_title('RMSE by Window', fontsize=12, fontweight='bold')
 ax1.legend()
 ax1.grid(True, alpha=0.3)

```

```

Plot 2: Improvement
ax2 = axes[1]
colors = ['green' if x > 0 else 'red' for x in improvement]
ax2.bar(windows, improvement, color=colors, alpha=0.7)
ax2.axhline(y=0, color='black', linestyle='-', linewidth=1)
ax2.set_xlabel('Window ID')
ax2.set_ylabel('RMSE Improvement (%)')
ax2.set_title('Hybrid Improvement over ARIMA', fontsize=12, fontweight='bold')
ax2.grid(True, alpha=0.3, axis='y')

plt.tight_layout()
plt.show()

```

In [ ]: `plot_all_windows_comparison(sp500_hybrid_results)`  
`plot_all_windows_comparison(bitcoin_hybrid_results)`

## EURUSD Dataset Analysis

### ARIMA + LSTM Hybrid

Dataset: EUR/USD Exchange Rate (2009-08-11 to 2019-08-11)

```

In []: # Download EURUSD data
print("Downloading EURUSD data...")
eurusd_data = yf.download("EURUSD=X", start="2009-08-11", end="2019-08-11", progress=False)

print(f"\nEURUSD Data Shape: {eurusd_data.shape}")
print(f"EURUSD Date Range: {eurusd_data.index.min()} to {eurusd_data.index.max()}")
print(f"Total EURUSD observations: {len(eurusd_data)}")

Calculate log returns
eurusd_data['Log_Returns'] = np.log(eurusd_data['Close'] / eurusd_data['Close'].shift(1))

Clean data
eurusd_clean = eurusd_data.dropna()

print(f"\nAfter cleaning: {len(eurusd_clean)} observations")
print(f"Mean daily return: {eurusd_clean['Log_Returns'].mean():.6f}")
print(f"Standard deviation: {eurusd_clean['Log_Returns'].std():.6f}")

```

```
In []: from dateutil.relativedelta import relativedelta
from datetime import timedelta

def create_eurusd_cv_splits(data, start_date=None):
 if start_date is None:
 start_date = data.index.min()

 cv_splits = []
 window_start = start_date

 while True:
 # Define window boundaries
 train_start = window_start
 train_end = train_start + relativedelta(years=2) - timedelta(days=1)

 # Validation periods (8, 16, 24 months)
 val_start = train_end + timedelta(days=1)
 val1_end = val_start + relativedelta(months=8) - timedelta(days=1)
 val2_end = val_start + relativedelta(months=16) - timedelta(days=1)
 val3_end = val_start + relativedelta(months=24) - timedelta(days=1)

 # Test period (6 months)
 test_start = val3_end + timedelta(days=1)
 test_end = test_start + relativedelta(months=6) - timedelta(days=1)

 # Check if we have enough data
 if test_end > data.index.max():
 break

 # Create splits for this window
 train_data = data[(data.index >= train_start) & (data.index <= train_end)]

 # Three validation folds
 val1_data = data[(data.index >= val_start) & (data.index <= val1_end)]
 val2_data = data[(data.index >= val_start) & (data.index <= val2_end)]
 val3_data = data[(data.index >= val_start) & (data.index <= val3_end)]

 test_data = data[(data.index >= test_start) & (data.index <= test_end)]

 cv_splits.append({
 'window_id': len(cv_splits) + 1,
 'train': {
 'data': train_data,
 'start': train_start,
```

```
 'end': train_end,
 'size': len(train_data)
 },
 'validation': [
 {
 'fold': 1,
 'data': val1_data,
 'start': val_start,
 'end': val1_end,
 'size': len(val1_data),
 'months': 8
 },
 {
 'fold': 2,
 'data': val2_data,
 'start': val_start,
 'end': val2_end,
 'size': len(val2_data),
 'months': 16
 },
 {
 'fold': 3,
 'data': val3_data,
 'start': val_start,
 'end': val3_end,
 'size': len(val3_data),
 'months': 24
 }
],
 'test': {
 'data': test_data,
 'start': test_start,
 'end': test_end,
 'size': len(test_data)
 }
})
window_start += relativedelta(years=1)

return cv_splits

eurusd_cv_splits = create_eurusd_cv_splits(eurusd_clean)

print(f"\n==== EURUSD Cross-Validation Setup ===")
print(f"Total CV windows: {len(eurusd_cv_splits)}")
```

```
print(f"\nFirst window details:")
print(f" Train: {eurusd_cv_splits[0]['train']['start'].date()} to {eurusd_cv_splits[0]['train']['end'].date()} ({eurusd_cv_splits[0]['train'].date().strftime('%Y-%m-%d')})
print(f" Val 1: {eurusd_cv_splits[0]['validation'][0]['start'].date()} to {eurusd_cv_splits[0]['validation'][0]['end'].date()} ({eurusd_cv_splits[0]['validation'][0].date().strftime('%Y-%m-%d')})
print(f" Test: {eurusd_cv_splits[0]['test']['start'].date()} to {eurusd_cv_splits[0]['test']['end'].date()} ({eurusd_cv_splits[0]['test'].date().strftime('%Y-%m-%d'))}
```

```
In []: eurusd_hybrid_results = run_hybrid_arima_lstm_cv(
 cv_splits=eurusd_cv_splits,
 data_clean=eurusd_clean,
 asset_name='EURUSD',
 max_p=3,
 max_d=1,
 max_q=3,
 lookback=60,
 lstm_units=50,
 dropout=0.2,
 epochs=50,
 batch_size=32,
 verbose=True
)
```

```
In []: plot_all_windows_comparison(eurusd_hybrid_results)
```

## Trading Performance Metrics

```
In []: def volatility_predictions_to_returns_new(predictions, true_values, actual_returns, transaction_costs=0.0):

 min_len = min(len(predictions), len(true_values), len(actual_returns))
 predictions = predictions[:min_len]
 true_values = true_values[:min_len]
 actual_returns = (actual_returns.iloc[:min_len]
 if isinstance(actual_returns, pd.Series)
 else actual_returns[:min_len])

 actual_returns_array = (actual_returns.values
 if isinstance(actual_returns, pd.Series)
 else actual_returns)

 signal = np.where(predictions > transaction_costs, 1, -1)

 if transaction_costs > 0.0:
 signals = np.where(np.abs(actual_returns_array) > transaction_costs, signal, 0)
```

```
strategy_returns = signals * actual_returns_array

return pd.Series(strategy_returns)
```

```
In []: def get_all_predictions(model_results, data_clean, model_type="S&P", window_indices=None):

 if model_type == "S&P":
 cost = 0.005
 elif model_type == "Bitcoin":
 cost = 0.01
 else:
 cost = 0.001
 all_strategy_returns = []

 windows_to_use = model_results['window_results']
 if window_indices is not None:
 windows_to_use = [w for w in windows_to_use if w['window_id'] in window_indices]

 for window_result in windows_to_use:
 try:
 test_start = window_result['test_start']
 test_end = window_result['test_end']

 test_data = data_clean[test_start:test_end]

 predictions = window_result['hybrid_predictions']
 true_values = test_data['Log_Returns'].values[-len(predictions):]

 actual_returns = test_data['Log_Returns'].iloc[-len(predictions):]

 window_returns = volatility_predictions_to_returns_new(
 predictions, true_values, actual_returns.values, transaction_costs=cost
)

 all_strategy_returns.append(window_returns)

 except Exception as e:
 print(f"Warning: Failed to process window {window_result.get('window_id', '?')}: {str(e)}")
 continue

 if all_strategy_returns:
```

```
 return pd.concat(all_strategy_returns, ignore_index=True)
 else:
 return pd.Series([])
```

```
In []: def annualized_return(daily_returns):
 cumulative = (1 + daily_returns).prod()
 n = daily_returns.shape[0]
 return cumulative ** (TRADING_DAYS / n) - 1

def annualized_std(daily_returns):
 return daily_returns.std() * np.sqrt(TRADING_DAYS)

def max_drawdown(daily_returns):
 equity = (1 + daily_returns).cumprod()
 peak = equity.cummax()
 drawdown = (equity - peak) / peak
 return np.abs(drawdown.min()) # Paper uses absolute value

def information_ratio(strategy_returns, benchmark_returns):
 arc = annualized_return(strategy_returns)
 asd = annualized_std(strategy_returns)

 if asd == 0:
 return np.nan
 return arc / asd

def modified_information_ratio(strategy_returns, benchmark_returns):
 arc = annualized_return(strategy_returns)
 asd = annualized_std(strategy_returns)
 md = max_drawdown(strategy_returns)

 if asd == 0 or md == 0:
 return np.nan

 return (arc * np.sign(arc) * arc) / (asd * md)

def sortino_ratio(daily_returns, risk_free_rate=0):
```

```

negative_returns = daily_returns[daily_returns < 0]

if len(negative_returns) == 0:
 return np.nan

downside_std = np.std(negative_returns, ddof=1)
asd_downside = downside_std * np.sqrt(TRADING_DAYS)

arc = annualized_return(daily_returns)

if asd_downside == 0:
 return np.nan

return arc / asd_downside

def compute_performance_indicators(strategy_returns, benchmark_returns):
 return {
 "ARC": annualized_return(strategy_returns),
 "ASD": annualized_std(strategy_returns),
 "MD": abs(max_drawdown(strategy_returns)),
 "IR": information_ratio(strategy_returns, benchmark_returns),
 "IR*": modified_information_ratio(strategy_returns, benchmark_returns),
 "SR": sortino_ratio(strategy_returns)
 }

```

## S&P 500

In [ ]: TRADING\_DAYS = 232

```

Get benchmark returns (Buy-and-Hold)
sp500_bnh_returns = sp500_clean['Log_Returns'].loc["2007-01-01":"2023-12-29"].values

sp500_hybrid_predictions = get_all_predictions(sp500_hybrid_results, sp500_clean)

sp500_hybrid_strategy_returns = sp500_hybrid_predictions
sp500_bnh_aligned = sp500_bnh_returns[-len(sp500_hybrid_strategy_returns):]

results_sp500 = []

HYBRID

```

```

hybrid_metrics = compute_performance_indicators(
 pd.Series(sp500_hybrid_strategy_returns),
 pd.Series(sp500_bnh_aligned)
)
hybrid_metrics['Model'] = 'HYBRID'
hybrid_metrics['Num_Trades'] = int(np.sum(np.abs(np.diff(sp500_hybrid_strategy_returns > 0)) > 0))
results_sp500.append(hybrid_metrics)

table2_sp500 = pd.DataFrame(results_sp500)

print("TABLE: S&P 500 Long-Short Strategy Results")
print(table2_sp500[['Model', 'ARC', 'ASD', 'MD', 'IR', 'IR*', 'SR']].to_string(index=False))

table2_sp500.to_csv('table2_sp500.csv', index=False)
print("\n Results saved to 'table2_sp500.csv'"')

```

## Bitcoin

```

In []: TRADING_DAYS = 345

Get benchmark returns (Buy-and-Hold)
bitcoin_bnh_returns = bitcoin_clean['Log_Returns'].values

bitcoin_hybrid_predictions = get_all_predictions(bitcoin_hybrid_results, bitcoin_clean)

bitcoin_hybrid_strategy_returns = bitcoin_hybrid_predictions
bitcoin_bnh_aligned = bitcoin_bnh_returns[-len(bitcoin_hybrid_strategy_returns):]

results_bitcoin = []

HYBRID
hybrid_metrics = compute_performance_indicators(
 pd.Series(bitcoin_hybrid_strategy_returns),
 pd.Series(bitcoin_bnh_aligned)
)
hybrid_metrics['Model'] = 'HYBRID'
hybrid_metrics['Num_Trades'] = int(np.sum(np.abs(np.diff(bitcoin_hybrid_strategy_returns > 0)) > 0))
results_bitcoin.append(hybrid_metrics)

table2_bitcoin = pd.DataFrame(results_bitcoin)

```

```

print("TABLE: Bitcoin Long-Short Strategy Results")
print(table2_bitcoin[['Model', 'ARC', 'ASD', 'MD', 'IR', 'IR*', 'SR']].to_string(index=False))

table2_bitcoin.to_csv('table2_bitcoin.csv', index=False)
print("\n Results saved to 'table2_bitcoin.csv'")

```

## EURUSD

```

In []: TRADING_DAYS = 252

Get benchmark returns (Buy-and-Hold)
eurusd_bnh_returns = eurusd_clean['Log_Returns'].values

eurusd_hybrid_predictions = get_all_predictions(eurusd_hybrid_results, eurusd_clean)

eurusd_hybrid_strategy_returns = eurusd_hybrid_predictions
eurusd_bnh_aligned = eurusd_bnh_returns[-len(eurusd_hybrid_strategy_returns):]

results_eurusd = []

HYBRID
hybrid_metrics = compute_performance_indicators(
 pd.Series(eurusd_hybrid_strategy_returns),
 pd.Series(eurusd_bnh_aligned)
)
hybrid_metrics['Model'] = 'HYBRID'
hybrid_metrics['Num_Trades'] = int(np.sum(np.abs(np.diff(eurusd_hybrid_strategy_returns > 0)) > 0))
results_eurusd.append(hybrid_metrics)

table2_eurusd = pd.DataFrame(results_eurusd)

print("TABLE: EURUSD Long-Short Strategy Results")
print(table2_eurusd[['Model', 'ARC', 'ASD', 'MD', 'IR', 'IR*', 'SR']].to_string(index=False))

table2_eurusd.to_csv('table2_eurusd.csv', index=False)
print("\n Results saved to 'table2_eurusd.csv'")

best_idx = table2_eurusd['IR'].idxmax()

```