

```
In [ ]: import yfinance as yf
import pandas as pd
import math
from sklearn.metrics import mean_squared_error
import random
from datetime import datetime as dt
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
import numpy as np
import matplotlib.pyplot as plt
from statsmodels.tsa.stattools import adfuller
from statsmodels.tsa.arima.model import ARIMA
from sklearn.preprocessing import MinMaxScaler, StandardScaler
from sklearn.svm import SVR
from sklearn.model_selection import GridSearchCV
import warnings
import tensorflow as tf
import random

warnings.filterwarnings('ignore')

# Set display options for better output
pd.set_option('display.max_columns', None)
pd.set_option('display.width', 1000)

random.seed(42)
np.random.seed(42)
tf.random.set_seed(42)
```

Prepare Dataset

```
In [ ]: # Data Download Configuration
# S&P 500: January 1, 2002 to December 31, 2023
# Bitcoin: January 1, 2015 to December 31, 2023

# Define date ranges
sp500_start = "2002-01-01"
sp500_end = "2023-12-31"
bitcoin_start = "2015-01-01"
bitcoin_end = "2023-12-31"

print("Downloading S&P 500 data...")
```

```
sp500_data = yf.download("^GSPC", start=sp500_start, end=sp500_end, progress=False)

print("Downloading Bitcoin data...")
bitcoin_data = yf.download("BTC-USD", start=bitcoin_start, end=bitcoin_end, progress=False)

# Display basic information about downloaded data
print(f"\nS&P 500 Data Shape: {sp500_data.shape}")
print(f"S&P 500 Date Range: {sp500_data.index.min()} to {sp500_data.index.max()}")
print(f"Total S&P 500 observations: {len(sp500_data)}")

print(f"\nBitcoin Data Shape: {bitcoin_data.shape}")
print(f"Bitcoin Date Range: {bitcoin_data.index.min()} to {bitcoin_data.index.max()}")
print(f"Total Bitcoin observations: {len(bitcoin_data)}")

sp500_data.columns = sp500_data.columns.set_levels(['Adj Close' if x == 'Close' else x for x in sp500_data.columns.levels[0]], level=0)
bitcoin_data.columns = bitcoin_data.columns.set_levels(['Adj Close' if x == 'Close' else x for x in bitcoin_data.columns.levels[0]], level=0)

print("\nVerification - S&P 500 has Adj Close:", ('Adj Close', '^GSPC') in sp500_data.columns)
print("Verification - Bitcoin has Adj Close:", ('Adj Close', 'BTC-USD') in bitcoin_data.columns)
```

```
In [ ]: sp500_data['Log_Returns'] = np.log(sp500_data['Adj Close']) / sp500_data['Adj Close'].shift(1)

bitcoin_data['Log_Returns'] = np.log(bitcoin_data['Adj Close']) / bitcoin_data['Adj Close'].shift(1)

sp500_clean = sp500_data.dropna()
bitcoin_clean = bitcoin_data.dropna()
```

```
In [ ]: statistics_data = {
    'Observations': [len(sp500_clean), len(bitcoin_clean)],
    'Mean_Daily_Return': [sp500_clean['Log_Returns'].mean(), bitcoin_clean['Log_Returns'].mean()],
    'Standard_Deviation': [sp500_clean['Log_Returns'].std(), bitcoin_clean['Log_Returns'].std()],
    'Minimum_Return': [sp500_clean['Log_Returns'].min(), bitcoin_clean['Log_Returns'].min()],
    'Maximum_Return': [sp500_clean['Log_Returns'].max(), bitcoin_clean['Log_Returns'].max()],
    'Skewness': [sp500_clean['Log_Returns'].skew(), bitcoin_clean['Log_Returns'].skew()],
    'Kurtosis': [sp500_clean['Log_Returns'].kurtosis(), bitcoin_clean['Log_Returns'].kurtosis()],
    'Infinite_Values': [np.isinf(sp500_clean['Log_Returns']).sum(), np.isinf(bitcoin_clean['Log_Returns']).sum()],
    'NaN_Values': [sp500_clean['Log_Returns'].isnull().sum(), bitcoin_clean['Log_Returns'].isnull().sum()]
}
```

```
# Create DataFrame with asset names as index
returns_statistics = pd.DataFrame(statistics_data, index=['S&P_500', 'Bitcoin'])
```

```
returns_statistics.T
```

```
In [ ]: # Data Visualization and Final Verification

# Create visualizations
fig, ((ax1, ax2), (ax3, ax4)) = plt.subplots(2, 2, figsize=(15, 10))

# S&P 500 Price Series
ax1.plot(sp500_clean.index, sp500_clean['Adj Close'])
ax1.set_title('S&P 500 Adjusted Close Price (2002-2023)')
ax1.set_ylabel('Price ($)')
ax1.grid(True, alpha=0.3)

# Bitcoin Price Series
ax2.plot(bitcoin_clean.index, bitcoin_clean['Adj Close'])
ax2.set_title('Bitcoin Price (2015-2023)')
ax2.set_ylabel('Price ($)')
ax2.grid(True, alpha=0.3)

# S&P 500 Log Returns
ax3.plot(sp500_clean.index, sp500_clean['Log_Returns'])
ax3.set_title('S&P 500 Logarithmic Returns')
ax3.set_ylabel('Log Returns')
ax3.set_xlabel('Date')
ax3.axhline(y=0, color='red', linestyle='--', alpha=0.5)
ax3.grid(True, alpha=0.3)

# Bitcoin Log Returns
ax4.plot(bitcoin_clean.index, bitcoin_clean['Log_Returns'])
ax4.set_title('Bitcoin Logarithmic Returns')
ax4.set_ylabel('Log Returns')
ax4.set_xlabel('Date')
ax4.axhline(y=0, color='red', linestyle='--', alpha=0.5)
ax4.grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

# Display sample data
print("\n==== SAMPLE DATA PREVIEW ====\n")
print("S&P 500 Data (First 5 rows):")
print(sp500_clean[['Open', 'High', 'Low', 'Adj Close', 'Volume', 'Log_Returns']].head())
```

```
print("\nBitcoin Data (First 5 rows):")
print(bitcoin_clean[['Open', 'High', 'Low', 'Adj Close', 'Volume', 'Log_Returns']].head())
```

Stationary Test

```
In [ ]: from statsmodels.tsa.stattools import adfuller
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf

sp500_adf_test = adfuller(sp500_clean['Log_Returns'])
# Output the results
print('ADF Statistic: %f' % sp500_adf_test[0])
print('p-value: %f' % sp500_adf_test[1])

plot_acf(sp500_clean['Log_Returns'], lags=40)
plot_pacf(sp500_clean['Log_Returns'], lags=40)
plt.show()
```

```
In [ ]: bitcoin_adf_test = adfuller(bitcoin_clean['Log_Returns'])
# Output the results
print('ADF Statistic: %f' % bitcoin_adf_test[0])
print('p-value: %f' % bitcoin_adf_test[1])

plot_acf(bitcoin_clean['Log_Returns'], lags=40)
plot_pacf(bitcoin_clean['Log_Returns'], lags=40)
plt.show()
```

Cross Validation

```
In [ ]: # Time Series Cross-Validation Implementation
# Novel 3-fold cross-validation scheme with rolling windows

import matplotlib.dates as mdates
from datetime import datetime, timedelta
from dateutil.relativedelta import relativedelta

# print("==== TIME SERIES CROSS-VALIDATION IMPLEMENTATION ====\n")

def create_sp500_cv_splits(data, start_date=None):
    if start_date is None:
        start_date = data.index.min()
```

```

cv_splits = []
window_start = start_date

while True:
    # Define window boundaries
    train_start = window_start
    train_end = train_start + relativedelta(years=3) - timedelta(days=1)

    # Validation periods (8, 16, 24 months)
    val_start = train_end + timedelta(days=1)
    val1_end = val_start + relativedelta(months=8) - timedelta(days=1) # 8 months
    val2_end = val_start + relativedelta(months=16) - timedelta(days=1) # 16 months
    val3_end = val_start + relativedelta(months=24) - timedelta(days=1) # 24 months (2 years)

    # Test period (1 year)
    test_start = val3_end + timedelta(days=1)
    test_end = test_start + relativedelta(years=1) - timedelta(days=1)

    # Check if we have enough data
    # if test_end > data.index.max():
    if test_end.year > 2024:
        break

    # Create splits for this window
    train_data = data[(data.index >= train_start) & (data.index <= train_end)]

    # Three validation folds
    val1_data = data[(data.index >= val_start) & (data.index <= val1_end)]
    val2_data = data[(data.index >= val_start) & (data.index <= val2_end)]
    val3_data = data[(data.index >= val_start) & (data.index <= val3_end)]

    test_data = data[(data.index >= test_start) & (data.index <= test_end)]

    cv_splits.append({
        'window_id': len(cv_splits) + 1,
        'train': {
            'data': train_data,
            'start': train_start,
            'end': train_end,
            'size': len(train_data)
        },
        'validation': [
            {
                'fold': 1,

```

```

        'data': val1_data,
        'start': val_start,
        'end': val1_end,
        'size': len(val1_data),
        'months': 8
    },
    {
        'fold': 2,
        'data': val2_data,
        'start': val_start,
        'end': val2_end,
        'size': len(val2_data),
        'months': 16
    },
    {
        'fold': 3,
        'data': val3_data,
        'start': val_start,
        'end': val3_end,
        'size': len(val3_data),
        'months': 24
    }
],
'test': {
    'data': test_data,
    'start': test_start,
    'end': test_end,
    'size': len(test_data)
}
})
)

# Move window forward by 1 year
window_start += relativedelta(years=1)

return cv_splits
}

def create_bitcoin_cv_splits(data, start_date=None):

    if start_date is None:
        # Start from a date that allows for proper window construction
        start_date = datetime(2015, 1, 1)

    cv_splits = []
    window_start = start_date

```

```

# Define the testing period constraint
test_period_start = datetime(2018, 1, 1)
test_period_end = datetime(2023, 12, 31)

while True:
    # Define window boundaries
    train_start = window_start
    train_end = train_start + relativedelta(years=2) - timedelta(days=1)

    # Validation periods (4, 8, 12 months)
    val_start = train_end + timedelta(days=1)
    val1_end = val_start + relativedelta(months=4) - timedelta(days=1) # 4 months
    val2_end = val_start + relativedelta(months=8) - timedelta(days=1) # 8 months
    val3_end = val_start + relativedelta(months=12) - timedelta(days=1) # 12 months

    # Test period (6 months)
    test_start = val3_end + timedelta(days=1)
    test_end = test_start + relativedelta(months=6) - timedelta(days=1)

    # Check constraints
    # if test_end > data.index.max() or test_end > test_period_end:
    #     break
    if test_end.year > 2023:
        break

    # Only include windows where test period is within 2018–2023
    if test_start < test_period_start:
        window_start += relativedelta(months=6)
        continue

    # Create splits for this window
    train_data = data[(data.index >= train_start) & (data.index <= train_end)]

    # Three validation folds
    val1_data = data[(data.index >= val_start) & (data.index <= val1_end)]
    val2_data = data[(data.index >= val_start) & (data.index <= val2_end)]
    val3_data = data[(data.index >= val_start) & (data.index <= val3_end)]

    test_data = data[(data.index >= test_start) & (data.index <= test_end)]

    cv_splits.append({
        'window_id': len(cv_splits) + 1,
        'train': {
            'data': train_data,
            'start': train_start,

```

```
        'end': train_end,
        'size': len(train_data)
    },
    'validation': [
        {
            'fold': 1,
            'data': val1_data,
            'start': val_start,
            'end': val1_end,
            'size': len(val1_data),
            'months': 4
        },
        {
            'fold': 2,
            'data': val2_data,
            'start': val_start,
            'end': val2_end,
            'size': len(val2_data),
            'months': 8
        },
        {
            'fold': 3,
            'data': val3_data,
            'start': val_start,
            'end': val3_end,
            'size': len(val3_data),
            'months': 12
        }
    ],
    'test': {
        'data': test_data,
        'start': test_start,
        'end': test_end,
        'size': len(test_data)
    }
}

# Move window forward by 6 months
window_start += relativedelta(months=6)

return cv_splits
```

In []: # Apply Cross-Validation Schemes to Data

```
# Generate S&P 500 cross-validation splits
sp500_cv_splits = create_sp500_cv_splits(sp500_clean)

# Generate Bitcoin cross-validation splits
bitcoin_cv_splits = create_bitcoin_cv_splits(bitcoin_clean)

# Create summary DataFrames
def create_cv_summary(cv_splits, asset_name):
    summary_data = []

    for split in cv_splits:
        # Add training data info
        summary_data.append({
            'Asset': asset_name,
            'Window_ID': split['window_id'],
            'Split_Type': 'Train',
            'Fold': 'N/A',
            'Start_Date': split['train']['start'].strftime('%Y-%m-%d'),
            'End_Date': split['train']['end'].strftime('%Y-%m-%d'),
            'Size': split['train']['size'],
            'Duration_Months': 'N/A'
        })

        # Add validation data info
        for val_fold in split['validation']:
            summary_data.append({
                'Asset': asset_name,
                'Window_ID': split['window_id'],
                'Split_Type': 'Validation',
                'Fold': val_fold['fold'],
                'Start_Date': val_fold['start'].strftime('%Y-%m-%d'),
                'End_Date': val_fold['end'].strftime('%Y-%m-%d'),
                'Size': val_fold['size'],
                'Duration_Months': val_fold['months']
            })

        # Add test data info
        summary_data.append({
            'Asset': asset_name,
            'Window_ID': split['window_id'],
            'Split_Type': 'Test',
            'Fold': 'N/A',
            'Start_Date': split['test']['start'].strftime('%Y-%m-%d'),
```

```

        'End_Date': split['test']['end'].strftime('%Y-%m-%d'),
        'Size': split['test']['size'],
        'Duration_Months': 'N/A'
    })

    return pd.DataFrame(summary_data)

# Create summary DataFrames
sp500_cv_summary = create_cv_summary(sp500_cv_splits, 'S&P_500')
bitcoin_cv_summary = create_cv_summary(bitcoin_cv_splits, 'Bitcoin')

# Combined summary
cv_summary_combined = pd.concat([sp500_cv_summary, bitcoin_cv_summary], ignore_index=True)

# Display first few windows for each asset
print("\n==== S&P 500 CV WINDOWS (First set) ===")
for i, split in enumerate(sp500_cv_splits[:1]):
    print(f"\nWindow {split['window_id']}:")
    print(f"  Train: {split['train']['start'].strftime('%Y-%m-%d')} to {split['train']['end'].strftime('%Y-%m-%d')} ({split['train'].shape[0]})")
    print(f"  Validation Folds:")
    for val_fold in split['validation']:
        print(f"    Fold {val_fold['fold']} ({val_fold['months']}mo): {val_fold['start'].strftime('%Y-%m-%d')} to {val_fold['end'].strftime('%Y-%m-%d')}")
    print(f"  Test: {split['test']['start'].strftime('%Y-%m-%d')} to {split['test']['end'].strftime('%Y-%m-%d')} ({split['test'].shape[0]})")

print("\n==== BITCOIN CV WINDOWS (First set) ===")
for i, split in enumerate(bitcoin_cv_splits[:1]):
    print(f"\nWindow {split['window_id']}:")
    print(f"  Train: {split['train']['start'].strftime('%Y-%m-%d')} to {split['train']['end'].strftime('%Y-%m-%d')} ({split['train'].shape[0]})")
    print(f"  Validation Folds:")
    for val_fold in split['validation']:
        print(f"    Fold {val_fold['fold']} ({val_fold['months']}mo): {val_fold['start'].strftime('%Y-%m-%d')} to {val_fold['end'].strftime('%Y-%m-%d')}")
    print(f"  Test: {split['test']['start'].strftime('%Y-%m-%d')} to {split['test']['end'].strftime('%Y-%m-%d')} ({split['test'].shape[0]})")

# Display summary statistics
print("\n==== CV SPLIT STATISTICS ===")
split_stats = cv_summary_combined.groupby(['Asset', 'Split_Type']).agg({
    'Size': ['mean', 'std', 'min', 'max'],
    'Window_ID': 'count'
}).round(0)
print(split_stats)

```

In []: # Visualize Cross-Validation Scheme

```

def plot_cv_timeline(cv_splits, asset_name, max_windows=8):

    fig, ax = plt.subplots(figsize=(16, max(6, len(cv_splits[:max_windows]) * 1.5)))

    # Colors for different split types
    colors = {
        'train': '#2E8B57',
        'val_fold1': '#4169E1',
        'val_fold2': '#1E90FF',
        'val_fold3': '#87CEEB',
        'test': '#DC143C'
    }

    y_positions = []

    for i, split in enumerate(cv_splits[:max_windows]):
        y_pos = len(cv_splits[:max_windows]) - i - 1
        y_positions.append(y_pos)

        # Plot training period
        ax.barh(y_pos, (split['train']['end'] - split['train']['start']).days,
                 left=split['train']['start'], height=0.6,
                 color=colors['train'], alpha=0.8, label='Train' if i == 0 else "")

        # Plot validation periods
        val_colors = ['val_fold1', 'val_fold2', 'val_fold3']
        for j, val_fold in enumerate(split['validation']):
            ax.barh(y_pos + 0.1 + j*0.15, (val_fold['end'] - val_fold['start']).days,
                     left=val_fold['start'], height=0.12,
                     color=colors[val_colors[j]], alpha=0.8,
                     label=f'Val Fold {j+1} ({val_fold["months"]} mo)' if i == 0 else "")

        # Plot test period
        ax.barh(y_pos, (split['test']['end'] - split['test']['start']).days,
                 left=split['test']['start'], height=0.6,
                 color=colors['test'], alpha=0.8, label='Test' if i == 0 else "")

        # Add window labels
        ax.text(split['train']['start'], y_pos, f'W{split["window_id"]}',
                verticalalignment='center', fontsize=9, fontweight='bold')

    # Formatting
    ax.set_xlim(-0.5, len(cv_splits[:max_windows]) - 0.5)
    ax.set_ylabel('CV Windows (Newest to Oldest)', fontsize=12)
    ax.set_xlabel('Time Period', fontsize=12)

```

```

    ax.set_title(f'{asset_name} Cross-Validation Timeline\n{len(cv_splits)} Total Windows, Showing First {min(max_windows, len(cv_splits))} Windows, From {cv_splits[0].start} to {cv_splits[-1].end}\n', fontweight='bold')

    # Format x-axis
    ax.xaxis.set_major_formatter(mdates.DateFormatter('%Y-%m'))
    ax.xaxis.set_major_locator(mdates.YearLocator())
    plt.xticks(rotation=45)

    # Add legend
    ax.legend(loc='upper right', bbox_to_anchor=(1, 1), frameon=True, fancybox=True, shadow=True)

    # Add grid
    ax.grid(True, alpha=0.3, axis='x')

    plt.tight_layout()
    return fig, ax

# Create a comprehensive comparison chart
def create_cv_comparison_chart():
    """Create a side-by-side comparison of CV schemes"""

    fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(20, 8))

    # S&P 500 scheme visualization
    y_pos = 1

    # S&P 500 scheme
    ax1.barh(y_pos, 3*365, left=0, height=0.6, color="#2E8B57", alpha=0.8, label='Train (3yr)')
    ax1.barh(y_pos+0.1, 8*30, left=3*365, height=0.15, color="#4169E1", alpha=0.8, label='Val Fold 1 (8mo)')
    ax1.barh(y_pos+0.25, 16*30, left=3*365, height=0.15, color="#1E90FF", alpha=0.8, label='Val Fold 2 (16mo)')
    ax1.barh(y_pos+0.4, 24*30, left=3*365, height=0.15, color="#87CEEB", alpha=0.8, label='Val Fold 3 (24mo)')
    ax1.barh(y_pos, 1*365, left=5*365, height=0.6, color="#DC143C", alpha=0.8, label='Test (1yr)')

    ax1.set_xlim(0, 6*365)
    ax1.set_ylim(0.5, 1.8)
    ax1.set_xlabel('Days')
    ax1.set_title('S&P 500 CV Scheme\n(6-year windows)', fontweight='bold')
    ax1.legend()
    ax1.grid(True, alpha=0.3)

    # Bitcoin scheme
    ax2.barh(y_pos, 2*365, left=0, height=0.6, color="#2E8B57", alpha=0.8, label='Train (2yr)')
    ax2.barh(y_pos+0.1, 4*30, left=2*365, height=0.15, color="#4169E1", alpha=0.8, label='Val Fold 1 (4mo)')
    ax2.barh(y_pos+0.25, 8*30, left=2*365, height=0.15, color="#1E90FF", alpha=0.8, label='Val Fold 2 (8mo)')
    ax2.barh(y_pos+0.4, 12*30, left=2*365, height=0.15, color="#87CEEB", alpha=0.8, label='Val Fold 3 (12mo)')

```

```

    ax2.barh(y_pos, 6*30, left=3*365, height=0.6, color="#DC143C", alpha=0.8, label='Test (6mo)')

    ax2.set_xlim(0, 3.5*365)
    ax2.set_ylim(0.5, 1.8)
    ax2.set_xlabel('Days')
    ax2.set_title('Bitcoin CV Scheme\n(~3.5-year windows)', fontweight='bold')
    ax2.legend()
    ax2.grid(True, alpha=0.3)

    plt.suptitle('Time Series Cross-Validation Scheme Comparison', fontsize=16, fontweight='bold')
    plt.tight_layout()
    return fig

# fig3 = create_cv_comparison_chart()
# plt.show()

# Create summary table
cv_scheme_summary = pd.DataFrame({
    'Asset': ['S&P 500', 'Bitcoin'],
    'Total_Windows': [len(sp500_cv_splits), len(bitcoin_cv_splits)],
    'Window_Length': ['6 years', '~3.5 years'],
    'Train_Period': ['3 years', '2 years'],
    'Validation_Folds': ['8/16/24 months', '4/8/12 months'],
    'Test_Period': ['1 year', '6 months'],
    'Window_Shift': ['1 year', '6 months'],
    'Test_Coverage': [
        f"{sp500_cv_splits[0]['test']['start'].strftime('%Y-%m-%d')} to {sp500_cv_splits[-1]['test']['end'].strftime('%Y-%m-%d')}",
        f"{bitcoin_cv_splits[0]['test']['start'].strftime('%Y-%m-%d')} to {bitcoin_cv_splits[-1]['test']['end'].strftime('%Y-%m-%d')}"
    ]
})

```

In []: # Utility Functions for Cross-Validation Data Access

```

def get_cv_data(cv_splits, window_id, fold=None, return_type='data'):
    split = next((s for s in cv_splits if s['window_id'] == window_id), None)
    if split is None:
        raise ValueError(f"Window ID {window_id} not found")

    if return_type == 'train':
        return split['train']['data']
    elif return_type == 'test':
        return split['test']['data']
    elif return_type == 'validation':
        if fold is None:

```

```
        raise ValueError("Fold number must be specified for validation data")
    if fold not in [1, 2, 3]:
        raise ValueError("Fold must be 1, 2, or 3")
    return split['validation'][fold-1]['data']
else:
    return split

# Example usage functions
def demonstrate_cv_usage():

    # Example 1: Get training data from first S&P 500 window
    train_data_sp500 = get_cv_data(sp500_cv_splits, window_id=1, return_type='train')
    print(f"S&P 500 Window 1 - Training data shape: {train_data_sp500.shape}")
    print(f"Training period: {train_data_sp500.index.min()} to {train_data_sp500.index.max()}")


    # Example 2: Get validation fold 2 data from first S&P 500 window
    val_data_sp500 = get_cv_data(sp500_cv_splits, window_id=1, fold=2, return_type='validation')
    print(f"\nS&P 500 Window 1 - Validation Fold 2 shape: {val_data_sp500.shape}")
    print(f"Validation period: {val_data_sp500.index.min()} to {val_data_sp500.index.max()}")


    # Example 3: Get test data from first Bitcoin window
    test_data_bitcoin = get_cv_data(bitcoin_cv_splits, window_id=1, return_type='test')
    print(f"\nBitcoin Window 1 - Test data shape: {test_data_bitcoin.shape}")
    print(f"Test period: {test_data_bitcoin.index.min()} to {test_data_bitcoin.index.max()}")


    return train_data_sp500, val_data_sp500, test_data_bitcoin

# Run demonstration
sample_train, sample_val, sample_test = demonstrate_cv_usage()

# Save CV splits for later use (optional)
cv_implementation_summary = {
    'sp500_cv_splits': sp500_cv_splits,
    'bitcoin_cv_splits': bitcoin_cv_splits,
    'sp500_summary': sp500_cv_summary,
    'bitcoin_summary': bitcoin_cv_summary,
    'combined_summary': cv_summary_combined,
    'scheme_comparison': cv_scheme_summary
}
```

In []:

ARIMA + SVM HYBRID

ARIMA

```
In [ ]: # ARIMA Model Implementation with AIC-based Selection
# Modern approach replacing traditional Box-Jenkins methodology

import itertools
from statsmodels.tsa.arima.model import ARIMA
from statsmodels.tsa.stattools import adfuller, kpss
from statsmodels.stats.diagnostic import acorr_ljungbox
import warnings
from sklearn.metrics import mean_squared_error, mean_absolute_error
from scipy import stats
import time

# Suppress convergence warnings for cleaner output
warnings.filterwarnings('ignore', category=UserWarning)
warnings.filterwarnings('ignore', category=RuntimeWarning)

print("== AIC-BASED ARIMA MODEL SELECTION IMPLEMENTATION ==\n")

def find_optimal_arima_order(data, max_p=5, max_d=2, max_q=5, seasonal=False,
                             information_criterion='aic', verbose=False):

    best_ic = np.inf
    best_params = None
    best_model = None
    results_log = []

    # Create parameter grid
    if seasonal:
        # For seasonal ARIMA (not implemented in this study)
        param_grid = itertools.product(range(max_p+1), range(max_d+1), range(max_q+1),
                                       range(2), range(2), range(2), [12])
    else:
        # Standard ARIMA grid search
        param_grid = itertools.product(range(max_p+1), range(max_d+1), range(max_q+1))

    total_combinations = (max_p+1) * (max_d+1) * (max_q+1)
```

```

if verbose:
    print(f"Testing {total_combinations} ARIMA parameter combinations...")
    start_time = time.time()

for i, params in enumerate(param_grid):
    p, d, q = params[:3]

    # Skip if model is too simple (all parameters zero)
    if p == 0 and d == 0 and q == 0:
        continue

    try:
        # Fit ARIMA model
        model = ARIMA(data, order=(p, d, q))
        fitted_model = model.fit()

        # Get information criterion value
        if information_criterion.lower() == 'aic':
            ic_value = fitted_model.aic
        elif information_criterion.lower() == 'bic':
            ic_value = fitted_model.bic
        elif information_criterion.lower() == 'hqic':
            ic_value = fitted_model.hqic
        else:
            ic_value = fitted_model.aic

        # Store results
        results_log.append({
            'order': (p, d, q),
            'aic': fitted_model.aic,
            'bic': fitted_model.bic,
            'hqic': fitted_model.hqic,
            'llf': fitted_model.llf,
            'converged': fitted_model.mle_retvals['converged'] if hasattr(fitted_model, 'mle_retvals') else True
        })

        # Update best model if current is better
        if ic_value < best_ic:
            best_ic = ic_value
            best_params = (p, d, q)
            best_model = fitted_model

    except Exception as e:
        # Log failed fits
        results_log.append({

```

```

        'order': (p, d, q),
        'aic': np.nan,
        'bic': np.nan,
        'hqic': np.nan,
        'llf': np.nan,
        'converged': False,
        'error': str(e)
    })

    if verbose and i % 10 == 0:
        print(f"Failed to fit ARIMA{params}: {str(e)[:50]}...")

    if verbose and (i + 1) % 20 == 0:
        elapsed = time.time() - start_time
        progress = (i + 1) / total_combinations * 100
        print(f"Progress: {progress:.1f}% ({i+1}/{total_combinations}) | "
              f"Best so far: ARIMA{best_params} ({information_criterion.upper()}={best_ic:.4f})")

    if verbose:
        total_time = time.time() - start_time
        print(f"\nGrid search completed in {total_time:.2f} seconds")
        print(f"Best model: ARIMA{best_params} with {information_criterion.upper()}={best_ic:.4f}")

# Create results summary
results_df = pd.DataFrame(results_log)
successful_fits = results_df[results_df['converged'] == True]

return {
    'best_order': best_params,
    'best_model': best_model,
    'best_ic_value': best_ic,
    'information_criterion': information_criterion,
    'results_df': results_df,
    'successful_fits': len(successful_fits),
    'total_attempts': len(results_log),
    'success_rate': len(successful_fits) / len(results_log) * 100
}

def evaluate_arima_model(model, train_data, test_data, model_order):
    # Generate forecasts
    n_forecast = len(test_data)
    forecast_result = model.get_forecast(steps=n_forecast)
    forecasts = forecast_result.predicted_mean
    forecast_ci = forecast_result.conf_int()

```

```

# Calculate performance metrics
mse = mean_squared_error(test_data, forecasts)
rmse = np.sqrt(mse)
mae = mean_absolute_error(test_data, forecasts)
mape = np.mean(np.abs((test_data - forecasts) / test_data)) * 100

# Calculate R2 score for fair comparison with SVM
ss_res = np.sum((test_data - forecasts) ** 2)
ss_tot = np.sum((test_data - np.mean(test_data)) ** 2)
r2 = 1 - (ss_res / ss_tot) if ss_tot > 0 else 0

# Direction accuracy (for returns)
direction_actual = np.sign(test_data.values[1:])
direction_forecast = np.sign(forecasts.values[1:])
direction_accuracy = np.mean(direction_actual == direction_forecast) * 100

# Residual diagnostics
residuals = model.resid

# Ljung-Box test for serial correlation in residuals
lb_test = acorr_ljungbox(residuals, lags=10, return_df=False)

# Normality test (Jarque-Bera)
jb_stat, jb_pvalue = stats.jarque_bera(residuals)

# Heteroskedasticity test (simple approach)
residuals_squared = residuals ** 2
arch_stat, arch_pvalue = acorr_ljungbox(residuals_squared, lags=5, return_df=False)

return {
    'model_order': model_order,
    'forecasts': forecasts,
    'forecast_ci': forecast_ci,
    'performance_metrics': {
        'mse': mse,
        'rmse': rmse,
        'mae': mae,
        'mape': mape,
        'r2': r2,
        'direction_accuracy': direction_accuracy
    },
    'diagnostic_tests': {
        'ljung_box_stat': lb_test['lb_stat'].iloc[-1],
        'ljung_box_pvalue': lb_test['lb_pvalue'].iloc[-1],
        'jarque_bera_stat': jb_stat,
    }
}

```

```

        'jarque_bera_pvalue': jb_pvalue,
        'arch_stat': arch_stat[-1] if isinstance(arch_stat, np.ndarray) else arch_stat,
        'arch_pvalue': arch_pvalue[-1] if isinstance(arch_pvalue, np.ndarray) else arch_pvalue
    },
    'residuals': residuals
}

```

HYBRID ARIMA-SVM

```

In [ ]: # Cross-Validation Integration for ARIMA Model Selection
# Implement the complete methodology with hyperparameter optimization

def run_arima_cross_validation(cv_splits, data_clean, asset_name, max_p=3, max_d=2, max_q=3,
                               information_criterion='aic', verbose=True):

    print(f"\n==== {asset_name.upper()} ARIMA CROSS-VALIDATION ===")
    print(f"Running AIC-based model selection across {len(cv_splits)} windows...")
    print(f"Parameter search space: p∈[0,{max_p}], d∈[0,{max_d}], q∈[0,{max_q}]")
    print("-" * 80)

    all_results = []
    model_selection_summary = []

    for window_idx, split in enumerate(cv_splits):
        window_id = split['window_id']

        if verbose:
            print(f"\n Processing Window {window_id}/{len(cv_splits)}...")
            print(f"  Train: {split['train']['start'].strftime('%Y-%m-%d')} to {split['train']['end'].strftime('%Y-%m-%d')}")
            print(f"  Test:  {split['test']['start'].strftime('%Y-%m-%d')} to {split['test']['end'].strftime('%Y-%m-%d')}")
            print("")

        # Extract data
        train_data = split['train']['data']['Log_Returns']
        test_data = split['test']['data']['Log_Returns']

        # STEP 1: Model Selection using Training Data
        if verbose:
            print(f"  Model selection using {information_criterion.upper()} criterion...")

        selection_result = find_optimal_arima_order(
            train_data,
            max_p=max_p,
            max_d=max_d,
            max_q=max_q,
            information_criterion=information_criterion)
    
```

```
    max_q=max_q,
    information_criterion=information_criterion,
    verbose=False # Keep individual window selection quiet
)

if selection_result['best_model'] is None:
    print(f"    Failed to find suitable model for Window {window_id}")
    continue

best_order = selection_result['best_order']

# STEP 2: Hyperparameter Validation using Validation Folds
if verbose:
    print(f"    Validating ARIMA{best_order} across 3 validation folds...")

validation_scores = []

for val_fold in split['validation']:
    fold_num = val_fold['fold']
    val_data = val_fold['data']['Log_Returns']

    try:
        # Fit model on training data and evaluate on validation fold
        val_model = ARIMA(train_data, order=best_order).fit()
        val_forecasts = val_model.get_forecast(steps=len(val_data)).predicted_mean
        val_rmse = np.sqrt(mean_squared_error(val_data, val_forecasts))
        validation_scores.append(val_rmse)

    except Exception as e:
        if verbose:
            print(f"        Validation fold {fold_num} failed: {str(e)[:50]}...")
        validation_scores.append(np.inf)

avg_validation_rmse = np.mean(validation_scores)

# STEP 3: Final Model Training and Out-of-Sample Evaluation
if verbose:
    print(f"    Final evaluation on test data...")

try:
    # Re-fit the model on training data
    final_model = ARIMA(train_data, order=best_order).fit()

    # Evaluate on test data
    evaluation = evaluate_arima_model(final_model, train_data, test_data, best_order)
```

```

# Store comprehensive results
window_result = {
    'window_id': window_id,
    'asset': asset_name,
    'train_period': f'{split["train"]['start'].strftime("%Y-%m-%d")}' to {split["train"]['end'].strftime("%Y-%m-%d")},
    'test_period': f'{split["test"]['start'].strftime("%Y-%m-%d")}' to {split["test"]['end'].strftime("%Y-%m-%d")}',
    'train_size': split['train']['size'],
    'test_size': split['test']['size'],
    'best_order': best_order,
    'model_selection': selection_result,
    'validation_scores': validation_scores,
    'avg_validation_rmse': avg_validation_rmse,
    'evaluation': evaluation,
    'final_model': final_model
}

all_results.append(window_result)

# Summary for quick reference
model_selection_summary.append({
    'Window': window_id,
    'Best_Order': f"ARIMA{best_order}",
    'AIC': selection_result['best_ic_value'],
    'Validation_RMSE': avg_validation_rmse,
    'Test_RMSE': evaluation['performance_metrics']['rmse'],
    'Test_MAE': evaluation['performance_metrics']['mae'],
    'Test_R2': evaluation['performance_metrics']['r2'],
    'Direction_Accuracy': evaluation['performance_metrics']['direction_accuracy'],
    'Ljung_Box_p': evaluation['diagnostic_tests']['ljung_box_pvalue']
})

if verbose:
    print(f"    ARIMA{best_order}: Test RMSE={evaluation['performance_metrics']['rmse']:.6f}, "
          f"Direction Acc={evaluation['performance_metrics']['direction_accuracy']:.1f}%")

except Exception as e:
    print(f"    Final evaluation failed for Window {window_id}: {str(e)}")
    continue

# Create summary DataFrame
summary_df = pd.DataFrame(model_selection_summary)

# Calculate overall performance statistics
if len(summary_df) > 0:

```

```

    performance_summary = {
        'total_windows': len(cv_splits),
        'successful_windows': len(summary_df),
        'success_rate': len(summary_df) / len(cv_splits) * 100,
        'avg_test_rmse': summary_df['Test_RMSE'].mean(),
        'std_test_rmse': summary_df['Test_RMSE'].std(),
        'avg_test_mae': summary_df['Test_MAE'].mean(),
        'avg_r2': summary_df['Test_R2'].mean(),
        'avg_direction_accuracy': summary_df['Direction_Accuracy'].mean(),
        'avg_validation_rmse': summary_df['Validation_RMSE'].mean(),
        'most_common_order': summary_df['Best_Order'].mode().iloc[0] if len(summary_df) > 0 else None
    }
else:
    performance_summary = None

print(f"\n{'='*80}")
print(f"{asset_name.upper()} ARIMA CROSS-VALIDATION COMPLETE")
print(f"\n{'='*80}")

if performance_summary:
    print(f" Successfully processed {performance_summary['successful_windows']} / {performance_summary['total_windows']} windows")
    print(f" Average Test RMSE: {performance_summary['avg_test_rmse']:.6f} ± {performance_summary['std_test_rmse']:.6f}")
    print(f" Average Direction Accuracy: {performance_summary['avg_direction_accuracy']:.2f}%")
    print(f" Most Common Model: {performance_summary['most_common_order']}")
else:
    print(" No successful model fits achieved")

return {
    'asset_name': asset_name,
    'all_results': all_results,
    'summary_df': summary_df,
    'performance_summary': performance_summary,
    'methodology': {
        'approach': 'AIC-based automated selection',
        'information_criterion': information_criterion,
        'parameter_space': f'p∈[0,{max_p}], d∈[0,{max_d}], q∈[0,{max_q}]',
        'cross_validation': '3-fold temporal validation',
        'evaluation_metric': 'Out-of-sample RMSE and direction accuracy'
    }
}

```

In []:

```

from sklearn.svm import SVR
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_squared_error, mean_absolute_error

```

```

def create_lagged_features_from_residuals(residuals, lookback=10):
    X, y = [], []
    for i in range(lookback, len(residuals)):
        X.append(residuals[i-lookback:i])
        y.append(residuals[i])
    return np.array(X), np.array(y)

def build_residual_svm_model(kernel='rbf', C=1.0, epsilon=0.1, gamma='scale'):
    model = SVR(
        kernel=kernel,
        C=C,
        epsilon=epsilon,
        gamma=gamma,
        cache_size=500
    )
    return model

def run_hybrid_arima_svm_cv(cv_splits, data_clean, asset_name,
                            max_p=3, max_d=1, max_q=3,
                            lookback=10, kernel='rbf', C=1.0,
                            epsilon=0.1, gamma='scale', verbose=True):

    print(f"\n{'='*100}")
    print(f"HYBRID ARIMA-SVM CROSS-VALIDATION: {asset_name.upper()}")
    print(f"{'='*100}")
    print(f"Architecture: ARIMA (linear) + SVM on residuals (non-linear)")
    print(f"ARIMA space: p∈[0,{max_p}], d∈[0,{max_d}], q∈[0,{max_q}]")
    print(f"SVM config: lookback={lookback}, kernel={kernel}, C={C}, epsilon={epsilon}")
    print(f"Total windows: {len(cv_splits)}")
    print("-" * 100)

    all_results = []
    scaler = StandardScaler()

    for window_idx, split in enumerate(cv_splits):
        window_id = split['window_id']

        if verbose:
            print(f"\n Window {window_id}/{len(cv_splits)}")
            print(f" Train: {split['train']['start'].strftime('%Y-%m-%d')} to {split['train']['end'].strftime('%Y-%m-%d')}")
            print(f" Test: {split['test']['start'].strftime('%Y-%m-%d')} to {split['test']['end'].strftime('%Y-%m-%d')}")
            print(f"{'='*100}")

```

```
try:
    # Extract data
    train_data = split['train']['data']['Log_Returns'].values
    test_data = split['test']['data']['Log_Returns'].values

    # STEP 1: FIT ARIMA MODEL (Linear Component)
    if verbose:
        print(f" [1/4] Fitting ARIMA model...")

    # Find optimal ARIMA order
    from statsmodels.tsa.arima.model import ARIMA
    best_aic = np.inf
    best_order = None
    best_model = None

    for p in range(max_p + 1):
        for d in range(max_d + 1):
            for q in range(max_q + 1):
                try:
                    model = ARIMA(train_data, order=(p, d, q))
                    fitted_model = model.fit()
                    if fitted_model.aic < best_aic:
                        best_aic = fitted_model.aic
                        best_order = (p, d, q)
                        best_model = fitted_model
                except:
                    continue

    if best_model is None:
        print(f" Failed to fit ARIMA for window {window_id}")
        continue

    if verbose:
        print(f" Best ARIMA{best_order}, AIC={best_aic:.2f}")

    # STEP 2: CALCULATE RESIDUALS (Non-linear Component to Model)
    if verbose:
        print(f" [2/4] Calculating ARIMA residuals...")

    # Get in-sample predictions and residuals
    arima_train_pred = best_model.fittedvalues
    train_residuals = train_data - arima_train_pred

    # Align arrays (ARIMA might drop initial values)
    if len(train_residuals) < len(train_data):
```

```

        n_dropped = len(train_data) - len(train_residuals)
        train_data_aligned = train_data[n_dropped:]
        train_residuals_aligned = train_residuals
    else:
        train_data_aligned = train_data
        train_residuals_aligned = train_residuals

    if verbose:
        print(f"      Residuals: mean={np.mean(train_residuals_aligned):.6f}, std={np.std(train_residuals_aligned):.6f}")

    # STEP 3: TRAIN SVM ON RESIDUALS

    if len(train_residuals_aligned) <= lookback:
        print(f"      Insufficient data for SVM (need > {lookback} points)")
        continue

    if verbose:
        print(f" [3/4] Training SVM on residuals...")

    # Create lagged features from residuals
    X_train, y_train = create_lagged_features_from_residuals(train_residuals_aligned, lookback)

    if len(X_train) == 0:
        print(f"      No features created from residuals")
        continue

    # Scale features
    X_train_scaled = scaler.fit_transform(X_train)

    # Build and train SVM
    svm_model = build_residual_svm_model(kernel=kernel, C=C, epsilon=epsilon, gamma=gamma)
    svm_model.fit(X_train_scaled, y_train)

    if verbose:
        print(f"      SVM trained with {len(X_train)} samples")

    # STEP 4: HYBRID FORECASTING ON TEST SET

    if verbose:
        print(f" [4/4] Generating hybrid forecasts...")

    # Make ARIMA forecasts for test period
    arima_test_pred = best_model.forecast(steps=len(test_data))

```

```

# For SVM residual prediction, we need historical residuals
# Use last 'lookback' residuals from training + generate iteratively for test

# Get last lookback residuals from training
last_residuals = train_residuals_aligned[-lookback:].flatten()
svm_residual_predictions = []

# Predict residuals for each test point
for t in range(len(test_data)):
    # Prepare input features (last lookback residuals)
    X_input = last_residuals[-lookback:].reshape(1, -1)
    X_input_scaled = scaler.transform(X_input)

    # Predict next residual
    residual_pred = svm_model.predict(X_input_scaled)[0]
    svm_residual_predictions.append(residual_pred)

    # Update sequence with actual residual for next iteration
    if t < len(test_data) - 1:
        actual_residual = test_data[t] - arima_test_pred[t]
        last_residuals = np.append(last_residuals[1:], actual_residual)

svm_residual_predictions = np.array(svm_residual_predictions)

# HYBRID PREDICTION = ARIMA + SVM_residuals
hybrid_predictions = arima_test_pred + svm_residual_predictions

# EVALUATE PERFORMANCE

# Metrics for ARIMA only
arima_rmse = np.sqrt(mean_squared_error(test_data, arima_test_pred))
arima_mae = mean_absolute_error(test_data, arima_test_pred)

# Metrics for HYBRID
hybrid_rmse = np.sqrt(mean_squared_error(test_data, hybrid_predictions))
hybrid_mae = mean_absolute_error(test_data, hybrid_predictions)

# Direction accuracy
arima_direction = np.mean(np.sign(test_data[1:]) == np.sign(arima_test_pred[1:])) * 100
hybrid_direction = np.mean(np.sign(test_data[1:]) == np.sign(hybrid_predictions[1:])) * 100

improvement = ((arima_rmse - hybrid_rmse) / arima_rmse) * 100

if verbose:
    print(f"    RESULTS:")

```

```

        print(f"      ARIMA only: RMSE={arima_rmse:.6f}, MAE={arima_mae:.6f}, Direction={arima_direction:.1f}%)")
        print(f"      HYBRID:     RMSE={hybrid_rmse:.6f}, MAE={hybrid_mae:.6f}, Direction={hybrid_direction:.1f}%)")
        print(f"      Improvement: {improvement:+.2f}% RMSE")

    # Store results
    window_results = {
        'window_id': window_id,
        'train_start': split['train']['start'],
        'train_end': split['train']['end'],
        'test_start': split['test']['start'],
        'test_end': split['test']['end'],
        'arima_order': best_order,
        'arima_aic': best_aic,
        # ARIMA metrics
        'arima_rmse': arima_rmse,
        'arima_mae': arima_mae,
        'arima_direction_accuracy': arima_direction,
        'arima_predictions': arima_test_pred,
        # SVM residual metrics
        'svm_residual_predictions': svm_residual_predictions,
        # HYBRID metrics
        'hybrid_rmse': hybrid_rmse,
        'hybrid_mae': hybrid_mae,
        'hybrid_direction_accuracy': hybrid_direction,
        'hybrid_predictions': hybrid_predictions,
        # Comparison
        'rmse_improvement_pct': improvement,
        'actuals': test_data
    }

    all_results.append(window_results)

except Exception as e:
    print(f"  Error in window {window_id}: {str(e)}")
    import traceback
    traceback.print_exc()
    continue

# SUMMARY
if len(all_results) > 0:
    avg_arima_rmse = np.mean([r['arima_rmse'] for r in all_results])
    avg_hybrid_rmse = np.mean([r['hybrid_rmse'] for r in all_results])
    avg_improvement = np.mean([r['rmse_improvement_pct'] for r in all_results])

    avg_arima_mae = np.mean([r['arima_mae'] for r in all_results])

```

```

avg_hybrid_mae = np.mean([r['hybrid_mae'] for r in all_results])

avg_arima_direction = np.mean([r['arima_direction_accuracy'] for r in all_results])
avg_hybrid_direction = np.mean([r['hybrid_direction_accuracy'] for r in all_results])

print(f"\n{='*100}'")
print(f" {asset_name.upper()} HYBRID MODEL COMPLETE")
print(f"{'='*100}'")
print(f"Windows processed: {len(all_results)}/{len(cv_splits)}")
print(f"\n AVERAGE PERFORMANCE:")
print(f"    ARIMA only: RMSE={avg_arima_rmse:.6f}, MAE={avg_arima_mae:.6f}, Direction={avg_arima_direction:.2f}%")
print(f"    HYBRID:      RMSE={avg_hybrid_rmse:.6f}, MAE={avg_hybrid_mae:.6f}, Direction={avg_hybrid_direction:.2f}%")
print(f"    Improvement: {avg_improvement:+.2f}% RMSE")
print(f"{'='*100}'")

results_dict = {
    'asset_name': asset_name,
    'model_type': 'HYBRID_ARIMA_SVM',
    'windows_processed': len(all_results),
    'total_windows': len(cv_splits),
    # ARIMA metrics
    'avg_arima_rmse': avg_arima_rmse,
    'avg_arima_direction': avg_arima_direction,
    'avg_arima_mae': avg_arima_mae,
    # HYBRID metrics
    'avg_hybrid_rmse': avg_hybrid_rmse,
    'avg_hybrid_direction': avg_hybrid_direction,
    'avg_hybrid_mae': avg_hybrid_mae,
    # Improvement
    'avg_improvement_pct': avg_improvement,
    'window_results': all_results,
    'hyperparameters': {
        'arima': f'p∈[0,{max_p}], d∈[0,{max_d}], q∈[0,{max_q}]',
        'svm_lookback': lookback,
        'svm_kernel': kernel,
        'svm_C': C,
        'svm_epsilon': epsilon,
        'svm_gamma': gamma
    }
}

return results_dict
else:

```

```
    print(f"\n No windows successfully processed for {asset_name}")
    return None
```

```
In [ ]: print("\n S&P 500 HYBRID ARIMA-SVM")

sp500_hybrid_results = run_hybrid_arima_svm_cv(
    cv_splits=sp500_cv_splits,
    data_clean=sp500_clean,
    asset_name='S&P 500',
    max_p=3,
    max_d=1,
    max_q=3,
    lookback=10,
    kernel='rbf',
    C=1.0,
    epsilon=0.1,
    gamma='scale',
    verbose=True
)

print("\n\n BITCOIN HYBRID ARIMA-SVM")

bitcoin_hybrid_results = run_hybrid_arima_svm_cv(
    cv_splits=bitcoin_cv_splits,
    data_clean=bitcoin_clean,
    asset_name='Bitcoin',
    max_p=3,
    max_d=1,
    max_q=3,
    lookback=10,
    kernel='rbf',
    C=1.0,
    epsilon=0.1,
    gamma='scale',
    verbose=True
)

print(" FINAL ")

summary_data = []

if sp500_hybrid_results:
    summary_data.append({
```

```

        'Asset': 'S&P 500',
        'ARIMA_RMSE': sp500_hybrid_results['avg_arima_rmse'],
        'Hybrid_RMSE': sp500_hybrid_results['avg_hybrid_rmse'],
        'Improvement_%': sp500_hybrid_results['avg_improvement_pct'],
        'ARIMA_MAE': sp500_hybrid_results['avg_arima_mae'],
        'Hybrid_MAE': sp500_hybrid_results['avg_hybrid_mae'],
        'ARIMA_Direction_%': sp500_hybrid_results['avg_arima_direction'],
        'Hybrid_Direction_%': sp500_hybrid_results['avg_hybrid_direction'],
        'Windows': f"{sp500_hybrid_results['windows_processed']}/{sp500_hybrid_results['total_windows']}"
    })

if bitcoin_hybrid_results:
    summary_data.append({
        'Asset': 'Bitcoin',
        'ARIMA_RMSE': bitcoin_hybrid_results['avg_arima_rmse'],
        'Hybrid_RMSE': bitcoin_hybrid_results['avg_hybrid_rmse'],
        'Improvement_%': bitcoin_hybrid_results['avg_improvement_pct'],
        'ARIMA_MAE': bitcoin_hybrid_results['avg_arima_mae'],
        'Hybrid_MAE': bitcoin_hybrid_results['avg_hybrid_mae'],
        'ARIMA_Direction_%': bitcoin_hybrid_results['avg_arima_direction'],
        'Hybrid_Direction_%': bitcoin_hybrid_results['avg_hybrid_direction'],
        'Windows': f"{bitcoin_hybrid_results['windows_processed']}/{bitcoin_hybrid_results['total_windows']}"
    })

if summary_data:
    summary_df = pd.DataFrame(summary_data)
    print("\n", summary_df.to_string(index=False))

    print("\n" + "=" * 100)
    print(" KEY INSIGHTS:")
    print("=" * 100)

    for data in summary_data:
        asset = data['Asset']
        improvement = data['Improvement_%']

        if improvement > 0:
            print(f"\u2713 {asset}: Hybrid model OUTPERFORMS ARIMA by {improvement:.2f}% RMSE")
        elif improvement < 0:
            print(f"\u2717 {asset}: Hybrid model underperforms ARIMA by {abs(improvement):.2f}% RMSE")
        else:
            print(f"\u2714 {asset}: Hybrid model equals ARIMA performance")

```

```
In [ ]: # Visualization Functions for Hybrid ARIMA-SVM Results
```

```
def plot_hybrid_decomposition(results, window_idx=0):
    if not results or len(results['window_results']) == 0:
        print("No results to plot")
        return

    window = results['window_results'][window_idx]
    asset_name = results['asset_name']

    # Extract data
    actuals = window['actuals']
    arima_pred = window['arima_predictions']
    svm_residual = window['svm_residual_predictions']
    hybrid_pred = window['hybrid_predictions']

    # Create figure with subplots
    fig, axes = plt.subplots(3, 2, figsize=(16, 12))
    fig.suptitle(f'{asset_name} - Hybrid ARIMA-SVM Decomposition (Window {window["window_id"]})',
                 fontsize=16, fontweight='bold')

    # Plot 1: Actual vs ARIMA
    ax1 = axes[0, 0]
    ax1.plot(actuals, label='Actual', color='black', linewidth=2, alpha=0.7)
    ax1.plot(arima_pred, label='ARIMA', color='blue', linewidth=1.5, linestyle='--', alpha=0.7)
    ax1.set_title('ARIMA Component (Linear)', fontsize=12, fontweight='bold')
    ax1.set_ylabel('Log Returns')
    ax1.legend()
    ax1.grid(True, alpha=0.3)
    ax1.text(0.02, 0.98, f'RMSE: {window["arima_rmse"]:.6f}',
             transform=ax1.transAxes, verticalalignment='top',
             bbox=dict(boxstyle='round', facecolor='wheat', alpha=0.5))

    # Plot 2: ARIMA Residuals
    ax2 = axes[0, 1]
    arima_residuals = actuals - arima_pred
    ax2.plot(arima_residuals, color='red', linewidth=1, alpha=0.7)
    ax2.axhline(y=0, color='black', linestyle='-', linewidth=0.5)
    ax2.fill_between(range(len(arima_residuals)), arima_residuals, 0, alpha=0.3, color='red')
    ax2.set_title('ARIMA Residuals (Non-linear to Model)', fontsize=12, fontweight='bold')
    ax2.set_ylabel('Residuals')
    ax2.grid(True, alpha=0.3)
    ax2.text(0.02, 0.98, f'Std: {np.std(arima_residuals):.6f}',
             transform=ax2.transAxes, verticalalignment='top',
```

```

bbox=dict(boxstyle='round', facecolor='wheat', alpha=0.5))

# Plot 3: SVM Residual Predictions
ax3 = axes[1, 0]
ax3.plot(arima_residuals, label='Actual Residuals', color='red', linewidth=1.5, alpha=0.7)
ax3.plot(svm_residual, label='SVM Predicted', color='green', linewidth=1.5, linestyle='--', alpha=0.7)
ax3.axhline(y=0, color='black', linestyle='-', linewidth=0.5)
ax3.set_title('SVM Residual Predictions', fontsize=12, fontweight='bold')
ax3.set_ylabel('Residuals')
ax3.legend()
ax3.grid(True, alpha=0.3)

# Plot 4: Actual vs Hybrid
ax4 = axes[1, 1]
ax4.plot(actuals, label='Actual', color='black', linewidth=2, alpha=0.7)
ax4.plot(hybrid_pred, label='Hybrid (ARIMA+SVM)', color='purple', linewidth=1.5, linestyle='--', alpha=0.7)
ax4.set_title('Final Hybrid Forecast', fontsize=12, fontweight='bold')
ax4.set_ylabel('Log Returns')
ax4.legend()
ax4.grid(True, alpha=0.3)
ax4.text(0.02, 0.98, f'RMSE: {window["hybrid_rmse"]:.6f}',
         transform=ax4.transAxes, verticalalignment='top',
         bbox=dict(boxstyle='round', facecolor='wheat', alpha=0.5))

# Plot 5: Component Comparison
ax5 = axes[2, 0]
sample_size = min(50, len(actuals))
x = range(sample_size)
ax5.plot(x, actuals[:sample_size], label='Actual', color='black', linewidth=2, alpha=0.7)
ax5.plot(x, arima_pred[:sample_size], label='ARIMA', color='blue', linewidth=1.5, linestyle='--', alpha=0.7)
ax5.plot(x, hybrid_pred[:sample_size], label='Hybrid', color='purple', linewidth=1.5, linestyle='--', alpha=0.7)
ax5.set_title(f'First {sample_size} Predictions Comparison', fontsize=12, fontweight='bold')
ax5.set_xlabel('Time Step')
ax5.set_ylabel('Log Returns')
ax5.legend()
ax5.grid(True, alpha=0.3)

# Plot 6: Performance Metrics
ax6 = axes[2, 1]
metrics = ['RMSE', 'MAE', 'Direction Acc (%)']
arima_metrics = [window['arima_rmse'], window['arima_mae'], window['arima_direction_accuracy']]
hybrid_metrics = [window['hybrid_rmse'], window['hybrid_mae'], window['hybrid_direction_accuracy']]

x_pos = np.arange(len(metrics))
width = 0.35

```

```

bars1 = ax6.bar(x_pos - width/2, arima_metrics, width, label='ARIMA', alpha=0.8)
bars2 = ax6.bar(x_pos + width/2, hybrid_metrics, width, label='Hybrid', alpha=0.8)

ax6.set_ylabel('Value')
ax6.set_title('Performance Metrics Comparison', fontsize=12, fontweight='bold')
ax6.set_xticks(x_pos)
ax6.set_xticklabels(metrics)
ax6.legend()
ax6.grid(True, alpha=0.3, axis='y')

# Add value labels
for bars in [bars1, bars2]:
    for bar in bars:
        height = bar.get_height()
        ax6.text(bar.get_x() + bar.get_width()/2., height,
                 f'{height:.4f}', ha='center', va='bottom', fontsize=8)

plt.tight_layout()
plt.show()

# Print detailed stats
print(f"\n{'='*80}")
print(f"DETAILED STATISTICS - Window {window['window_id']}")
print(f"{'='*80}")
print(f"Period: {window['test_start'].strftime('%Y-%m-%d')} to {window['test_end'].strftime('%Y-%m-%d')}")
print(f"ARIMA Order: {window['arima_order']}, AIC: {window['arima_aic']:.2f}")
print(f"\nARIMA Performance:")
print(f"    • RMSE: {window['arima_rmse']:.6f}")
print(f"    • MAE: {window['arima_mae']:.6f}")
print(f"    • Direction Accuracy: {window['arima_direction_accuracy']:.2f}%")
print(f"\nHybrid Performance:")
print(f"    • RMSE: {window['hybrid_rmse']:.6f}")
print(f"    • MAE: {window['hybrid_mae']:.6f}")
print(f"    • Direction Accuracy: {window['hybrid_direction_accuracy']:.2f}%")
print(f"\nImprovement: {window['rmse_improvement_pct']+2} % RMSE")
print(f"{'='*80}")

def plot_all_windows_comparison(results):
    if not results or len(results['window_results']) == 0:
        print("No results to plot")
        return

asset_name = results['asset_name']

```

```

windows = [r['window_id'] for r in results['window_results']]
arima_rmse = [r['arima_rmse'] for r in results['window_results']]
hybrid_rmse = [r['hybrid_rmse'] for r in results['window_results']]
improvement = [r['rmse_improvement_pct'] for r in results['window_results']]

fig, axes = plt.subplots(1, 2, figsize=(14, 5))
fig.suptitle(f'{asset_name} - Hybrid Model Performance Across All Windows',
             fontsize=14, fontweight='bold')

# Plot 1: RMSE Comparison
ax1 = axes[0]
ax1.plot(windows, arima_rmse, marker='o', label='ARIMA', linewidth=2, markersize=6)
ax1.plot(windows, hybrid_rmse, marker='s', label='Hybrid', linewidth=2, markersize=6)
ax1.set_xlabel('Window ID')
ax1.set_ylabel('RMSE')
ax1.set_title('RMSE by Window', fontsize=12, fontweight='bold')
ax1.legend()
ax1.grid(True, alpha=0.3)

# Plot 2: Improvement
ax2 = axes[1]
colors = ['green' if x > 0 else 'red' for x in improvement]
ax2.bar(windows, improvement, color=colors, alpha=0.7)
ax2.axhline(y=0, color='black', linestyle='-', linewidth=1)
ax2.set_xlabel('Window ID')
ax2.set_ylabel('RMSE Improvement (%)')
ax2.set_title('Hybrid Improvement over ARIMA', fontsize=12, fontweight='bold')
ax2.grid(True, alpha=0.3, axis='y')

plt.tight_layout()
plt.show()

# Summary statistics
print(f"\n{'='*80}")
print(f'{asset_name.upper()} - SUMMARY STATISTICS')
print(f"\n{'='*80}")
print(f"Total Windows: {len(windows)}")
print(f"Windows where Hybrid outperforms: {sum(1 for x in improvement if x > 0)}")
print(f"Windows where ARIMA outperforms: {sum(1 for x in improvement if x < 0)}")
print(f"Average Improvement: {np.mean(improvement):.2f}%")
print(f"Best Improvement: {max(improvement):.2f}% (Window {windows[improvement.index(max(improvement))]}"))
print(f"Worst Improvement: {min(improvement):.2f}% (Window {windows[improvement.index(min(improvement))]}"))
print(f"\n{'='*80}")

```

```
In [ ]: plot_all_windows_comparison(sp500_hybrid_results)
plot_all_windows_comparison(bitcoin_hybrid_results)
```

EURUSD Dataset Analysis

ARIMA + SVM Hybrid

Dataset: EUR/USD Exchange Rate (2009-08-11 to 2019-08-11)

```
In [ ]: # Download EURUSD data
print("Downloading EURUSD data...")
eurusd_data = yf.download("EURUSD=X", start="2009-08-11", end="2019-08-11", progress=False)

print(f"\nEURUSD Data Shape: {eurusd_data.shape}")
print(f"EURUSD Date Range: {eurusd_data.index.min()} to {eurusd_data.index.max()}")
print(f"Total EURUSD observations: {len(eurusd_data)}")

# Calculate log returns
eurusd_data['Log_RetURNS'] = np.log(eurusd_data['Close'] / eurusd_data['Close'].shift(1))

# Clean data
eurusd_clean = eurusd_data.dropna()

print(f"\nAfter cleaning: {len(eurusd_clean)} observations")
print(f"Mean daily return: {eurusd_clean['Log_RetURNS'].mean():.6f}")
print(f"Standard deviation: {eurusd_clean['Log_RetURNS'].std():.6f}")
```

```
In [ ]: from dateutil.relativedelta import relativedelta
from datetime import timedelta

def create_eurusd_cv_splits(data, start_date=None):

    if start_date is None:
        start_date = data.index.min()

    cv_splits = []
    window_start = start_date

    while True:
        # Define window boundaries
```

```
train_start = window_start
train_end = train_start + relativedelta(years=2) - timedelta(days=1)

# Validation periods (8, 16, 24 months)
val_start = train_end + timedelta(days=1)
val1_end = val_start + relativedelta(months=8) - timedelta(days=1)
val2_end = val_start + relativedelta(months=16) - timedelta(days=1)
val3_end = val_start + relativedelta(months=24) - timedelta(days=1)

# Test period (6 months)
test_start = val3_end + timedelta(days=1)
test_end = test_start + relativedelta(months=6) - timedelta(days=1)

# Check if we have enough data
if test_end > data.index.max():
    break

# Create splits for this window
train_data = data[(data.index >= train_start) & (data.index <= train_end)]

# Three validation folds
val1_data = data[(data.index >= val_start) & (data.index <= val1_end)]
val2_data = data[(data.index >= val_start) & (data.index <= val2_end)]
val3_data = data[(data.index >= val_start) & (data.index <= val3_end)]

test_data = data[(data.index >= test_start) & (data.index <= test_end)]

cv_splits.append({
    'window_id': len(cv_splits) + 1,
    'train': {
        'data': train_data,
        'start': train_start,
        'end': train_end,
        'size': len(train_data)
    },
    'validation': [
        {
            'fold': 1,
            'data': val1_data,
            'start': val_start,
            'end': val1_end,
            'size': len(val1_data),
            'months': 8
        },
        {

```

```

        'fold': 2,
        'data': val2_data,
        'start': val_start,
        'end': val2_end,
        'size': len(val2_data),
        'months': 16
    },
    {
        'fold': 3,
        'data': val3_data,
        'start': val_start,
        'end': val3_end,
        'size': len(val3_data),
        'months': 24
    }
],
'test': {
    'data': test_data,
    'start': test_start,
    'end': test_end,
    'size': len(test_data)
}
})
# Move window forward by 1 year
window_start += relativedelta(years=1)

return cv_splits
eurusd_cv_splits = create_eurusd_cv_splits(eurusd_clean)

```

In []: `print("\n\n New Dataset: EURUSD HYBRID ARIMA-SVM")`

```

eurusd_hybrid_results = run_hybrid_arima_svm_cv(
    cv_splits=eurusd_cv_splits,
    data_clean=eurusd_clean,
    asset_name='EURUSD',
    max_p=3,
    max_d=1,
    max_q=3,
    lookback=10,
    kernel='rbf',
    C=1.0,
    epsilon=0.1,

```

```
    gamma='scale',
    verbose=True
)
```

```
In [ ]: plot_all_windows_comparison(eurusd_hybrid_results)
```

Trading Performance Metrics

```
In [ ]: def volatility_predictions_to_returns_new(predictions, true_values, actual_returns, transaction_costs=0.0):

    # Ensure all arrays have matching length
    min_len = min(len(predictions), len(true_values), len(actual_returns))
    predictions = predictions[:min_len]
    true_values = true_values[:min_len]
    actual_returns = (actual_returns.iloc[:min_len]
                      if isinstance(actual_returns, pd.Series)
                      else actual_returns[:min_len])

    # Convert to numpy arrays for consistency
    actual_returns_array = (actual_returns.values
                            if isinstance(actual_returns, pd.Series)
                            else actual_returns)

    signal = np.where(predictions > transaction_costs, 1, -1)

    if transaction_costs > 0.0:
        signals = np.where(np.abs(actual_returns_array) > transaction_costs, signal, 0)

    strategy_returns = signals * actual_returns_array

    return pd.Series(strategy_returns)
```

```
In [ ]: def get_all_predictions(model_results, data_clean, model_type="S&P", window_indices=None):
    if model_type == "S&P":
        cost = 0.005
    elif model_type == "Bitcoin":
        cost = 0.01
    else:
        cost = 0.001
    all_strategy_returns = []
```

```

windows_to_use = model_results['window_results']
if window_indices is not None:
    windows_to_use = [w for w in windows_to_use if w['window_id'] in window_indices]

for window_result in windows_to_use:
    try:
        test_start = window_result['test_start']
        test_end = window_result['test_end']
        test_data = data_clean[test_start:test_end]

        predictions = window_result['hybrid_predictions']
        true_values = test_data['Log_Returns'].values[-len(predictions):]

        actual_returns = test_data['Log_Returns'].iloc[-len(predictions):]

        window_returns = volatility_predictions_to_returns_new(
            predictions, true_values, actual_returns.values, transaction_costs=cost
        )

        all_strategy_returns.append(window_returns)

    except Exception as e:
        print(f"Warning: Failed to process window {window_result.get('window_id', '?')}: {str(e)}")
        continue

if all_strategy_returns:
    return pd.concat(all_strategy_returns, ignore_index=True)
else:
    return pd.Series([])

```

```

In [ ]: def annualized_return(daily_returns):
    cumulative = (1 + daily_returns).prod()
    n = daily_returns.shape[0]
    return cumulative ** (TRADING_DAYS / n) - 1

def annualized_std(daily_returns):
    return daily_returns.std() * np.sqrt(TRADING_DAYS)

def max_drawdown(daily_returns):
    equity = (1 + daily_returns).cumprod()
    peak = equity.cummax()
    drawdown = (equity - peak) / peak

```

```
    return np.abs(drawdown.min()) # Paper uses absolute value

def information_ratio(strategy_returns, benchmark_returns):
    arc = annualized_return(strategy_returns)
    asd = annualized_std(strategy_returns)

    if asd == 0:
        return np.nan
    return arc / asd

def modified_information_ratio(strategy_returns, benchmark_returns):
    arc = annualized_return(strategy_returns)
    asd = annualized_std(strategy_returns)
    md = max_drawdown(strategy_returns)

    if asd == 0 or md == 0:
        return np.nan

    return (arc * np.sign(arc) * arc) / (asd * md)

def sortino_ratio(daily_returns, risk_free_rate=0):
    negative_returns = daily_returns[daily_returns < 0]

    if len(negative_returns) == 0:
        return np.nan

    downside_std = np.std(negative_returns, ddof=1)
    asd_downside = downside_std * np.sqrt(TRADING_DAYS)

    arc = annualized_return(daily_returns)

    if asd_downside == 0:
        return np.nan

    return arc / asd_downside

def compute_performance_indicators(strategy_returns, benchmark_returns):
    return {
        "ARC": annualized_return(strategy_returns),
        "ASD": annualized_std(strategy_returns),
        "MD": abs(max_drawdown(strategy_returns)),
```

```
        "IR": information_ratio(strategy_returns, benchmark_returns),
        "IR*": modified_information_ratio(strategy_returns, benchmark_returns),
        "SR": sortino_ratio(strategy_returns)
    }
```

S&P 500

```
In [ ]: TRADING_DAYS = 232
sp500_bnh_returns = sp500_clean['Log_Returns'].loc["2007-01-01":"2023-12-29"].values

sp500_hybrid_predictions = get_all_predictions(sp500_hybrid_results, sp500_clean)

sp500_hybrid_strategy_returns = sp500_hybrid_predictions

sp500_bnh_aligned = sp500_bnh_returns[-len(sp500_hybrid_strategy_returns):]

results_sp500 = []

# HYBRID
hybrid_metrics = compute_performance_indicators(
    pd.Series(sp500_hybrid_strategy_returns),
    pd.Series(sp500_bnh_aligned)
)
hybrid_metrics['Model'] = 'HYBRID'
hybrid_metrics['Num_Trades'] = int(np.sum(np.abs(np.diff(sp500_hybrid_strategy_returns > 0)) > 0))
results_sp500.append(hybrid_metrics)

table2_sp500 = pd.DataFrame(results_sp500)

print("TABLE: S&P 500 Long-Short Strategy Results")
print(table2_sp500[['Model', 'ARC', 'ASD', 'MD', 'IR', 'IR*', 'SR']].to_string(index=False))
table2_sp500.to_csv('table2_sp500.csv', index=False)
```

Bitcoin

```
In [ ]: TRADING_DAYS = 345
```

```

# Get benchmark returns (Buy-and-Hold)
bitcoin_bnh_returns = bitcoin_clean['Log_Returns'].values

bitcoin_hybrid_predictions = get_all_predictions(bitcoin_hybrid_results, bitcoin_clean, model_type="bitcoin")

bitcoin_hybrid_strategy_returns = bitcoin_hybrid_predictions

bitcoin_bnh_aligned = bitcoin_bnh_returns[-len(bitcoin_hybrid_strategy_returns):]

results_bitcoin = []

# HYBRID
hybrid_metrics = compute_performance_indicators(
    pd.Series(bitcoin_hybrid_strategy_returns),
    pd.Series(bitcoin_bnh_aligned)
)
hybrid_metrics['Model'] = 'HYBRID'
hybrid_metrics['Num_Trades'] = int(np.sum(np.abs(np.diff(bitcoin_hybrid_strategy_returns > 0)) > 0))
results_bitcoin.append(hybrid_metrics)

table2_bitcoin = pd.DataFrame(results_bitcoin)

print("TABLE: Bitcoin Long-Short Strategy Results")

print(table2_bitcoin[['Model', 'ARC', 'ASD', 'MD', 'IR', 'IR*', 'SR']].to_string(index=False))

table2_bitcoin.to_csv('table2_bitcoin.csv', index=False)

```

EURUSD

In []:

```

TRADING_DAYS = 252

# Get benchmark returns (Buy-and-Hold)
eurusd_bnh_returns = eurusd_clean['Log_Returns'].values

eurusd_hybrid_predictions = get_all_predictions(eurusd_hybrid_results, eurusd_clean, model_type="EURUSD")

eurusd_hybrid_strategy_returns = eurusd_hybrid_predictions

eurusd_bnh_aligned = eurusd_bnh_returns[-len(eurusd_hybrid_strategy_returns):]

```

```
results_eurusd = []

# HYBRID
hybrid_metrics = compute_performance_indicators(
    pd.Series(eurusd_hybrid_strategy_returns),
    pd.Series(eurusd_bnh_aligned)
)
hybrid_metrics['Model'] = 'HYBRID'
hybrid_metrics['Num_Trades'] = int(np.sum(np.abs(np.diff(eurusd_hybrid_strategy_returns > 0)) > 0))
results_eurusd.append(hybrid_metrics)

table2_eurusd = pd.DataFrame(results_eurusd)

print("TABLE: EURUSD Long-Short Strategy Results")

print(table2_eurusd[['Model', 'ARC', 'ASD', 'MD', 'IR', 'IR*', 'SR']].to_string(index=False))

table2_eurusd.to_csv('table2_eurusd.csv', index=False)
```