

```
In [ ]: import yfinance as yf
import pandas as pd
import math
from sklearn.metrics import mean_squared_error
import random
from datetime import datetime as dt
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
import numpy as np
import matplotlib.pyplot as plt
from statsmodels.tsa.stattools import adfuller
from statsmodels.tsa.arima.model import ARIMA
from sklearn.preprocessing import MinMaxScaler
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout, LSTM
import warnings
import tensorflow as tf
import random

warnings.filterwarnings('ignore')

# Set display options for better output
pd.set_option('display.max_columns', None)
pd.set_option('display.width', 1000)

random.seed(42)
np.random.seed(42)
tf.random.set_seed(42)
```

Prepare Dataset

```
In [ ]: # Data Download Configuration
# S&P 500: January 1, 2002 to December 31, 2023
# Bitcoin: January 1, 2015 to December 31, 2023

# Define date ranges
sp500_start = "2002-01-01"
sp500_end = "2023-12-31"
bitcoin_start = "2015-01-01"
bitcoin_end = "2023-12-31"

print("Downloading S&P 500 data...")
```

```
sp500_data = yf.download("^GSPC", start=sp500_start, end=sp500_end, progress=False)

print("Downloading Bitcoin data...")
bitcoin_data = yf.download("BTC-USD", start=bitcoin_start, end=bitcoin_end, progress=False)

# Display basic information about downloaded data
print(f"\nS&P 500 Data Shape: {sp500_data.shape}")
print(f"S&P 500 Date Range: {sp500_data.index.min()} to {sp500_data.index.max()}")
print(f"Total S&P 500 observations: {len(sp500_data)}")

print(f"\nBitcoin Data Shape: {bitcoin_data.shape}")
print(f"Bitcoin Date Range: {bitcoin_data.index.min()} to {bitcoin_data.index.max()}")
print(f"Total Bitcoin observations: {len(bitcoin_data)}")

sp500_data.columns = sp500_data.columns.set_levels(['Adj Close' if x == 'Close' else x for x in sp500_data.columns.levels[0]], level=0)
bitcoin_data.columns = bitcoin_data.columns.set_levels(['Adj Close' if x == 'Close' else x for x in bitcoin_data.columns.levels[0]], level=0)

print("\nVerification - S&P 500 has Adj Close:", ('Adj Close', '^GSPC') in sp500_data.columns)
print("Verification - Bitcoin has Adj Close:", ('Adj Close', 'BTC-USD') in bitcoin_data.columns)
```

```
In [ ]: sp500_data['Log_Returns'] = np.log(sp500_data['Adj Close']) / sp500_data['Adj Close'].shift(1)

bitcoin_data['Log_Returns'] = np.log(bitcoin_data['Adj Close']) / bitcoin_data['Adj Close'].shift(1)

sp500_clean = sp500_data.dropna()
bitcoin_clean = bitcoin_data.dropna()
```

```
In [ ]: statistics_data = {
    'Observations': [len(sp500_clean), len(bitcoin_clean)],
    'Mean_Daily_Return': [sp500_clean['Log_Returns'].mean(), bitcoin_clean['Log_Returns'].mean()],
    'Standard_Deviation': [sp500_clean['Log_Returns'].std(), bitcoin_clean['Log_Returns'].std()],
    'Minimum_Return': [sp500_clean['Log_Returns'].min(), bitcoin_clean['Log_Returns'].min()],
    'Maximum_Return': [sp500_clean['Log_Returns'].max(), bitcoin_clean['Log_Returns'].max()],
    'Skewness': [sp500_clean['Log_Returns'].skew(), bitcoin_clean['Log_Returns'].skew()],
    'Kurtosis': [sp500_clean['Log_Returns'].kurtosis(), bitcoin_clean['Log_Returns'].kurtosis()],
    'Infinite_Values': [np.isinf(sp500_clean['Log_Returns']).sum(), np.isinf(bitcoin_clean['Log_Returns']).sum()],
    'NaN_Values': [sp500_clean['Log_Returns'].isnull().sum(), bitcoin_clean['Log_Returns'].isnull().sum()]
}

# Create DataFrame with asset names as index
returns_statistics = pd.DataFrame(statistics_data, index=['S&P_500', 'Bitcoin'])
```

```
returns_statistics.T
```

```
In [ ]: # Data Visualization and Final Verification

print("== FINAL DATA VERIFICATION ==\n")

# Create visualizations
fig, ((ax1, ax2), (ax3, ax4)) = plt.subplots(2, 2, figsize=(15, 10))

# S&P 500 Price Series
ax1.plot(sp500_clean.index, sp500_clean['Adj Close'])
ax1.set_title('S&P 500 Adjusted Close Price (2002-2023)')
ax1.set_ylabel('Price ($)')
ax1.grid(True, alpha=0.3)

# Bitcoin Price Series
ax2.plot(bitcoin_clean.index, bitcoin_clean['Adj Close'])
ax2.set_title('Bitcoin Price (2015-2023)')
ax2.set_ylabel('Price ($)')
ax2.grid(True, alpha=0.3)

# S&P 500 Log Returns
ax3.plot(sp500_clean.index, sp500_clean['Log_Returns'])
ax3.set_title('S&P 500 Logarithmic Returns')
ax3.set_ylabel('Log Returns')
ax3.set_xlabel('Date')
ax3.axhline(y=0, color='red', linestyle='--', alpha=0.5)
ax3.grid(True, alpha=0.3)

# Bitcoin Log Returns
ax4.plot(bitcoin_clean.index, bitcoin_clean['Log_Returns'])
ax4.set_title('Bitcoin Logarithmic Returns')
ax4.set_ylabel('Log Returns')
ax4.set_xlabel('Date')
ax4.axhline(y=0, color='red', linestyle='--', alpha=0.5)
ax4.grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

# Display sample data
print("\n== SAMPLE DATA PREVIEW ==\n")
print("S&P 500 Data (First 5 rows):")
```

```
print(sp500_clean[['Open', 'High', 'Low', 'Adj Close', 'Volume', 'Log_Returns']].head())

print("\nBitcoin Data (First 5 rows):")
print(bitcoin_clean[['Open', 'High', 'Low', 'Adj Close', 'Volume', 'Log_Returns']].head())
```

Stationary Test

```
In [ ]: from statsmodels.tsa.stattools import adfuller
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf

sp500_adf_test = adfuller(sp500_clean['Log_Returns'])
# Output the results
print('ADF Statistic: %f' % sp500_adf_test[0])
print('p-value: %f' % sp500_adf_test[1])

plot_acf(sp500_clean['Log_Returns'], lags=40)
plot_pacf(sp500_clean['Log_Returns'], lags=40)
plt.show()
```

```
In [ ]: bitcoin_adf_test = adfuller(bitcoin_clean['Log_Returns'])
# Output the results
print('ADF Statistic: %f' % bitcoin_adf_test[0])
print('p-value: %f' % bitcoin_adf_test[1])

plot_acf(bitcoin_clean['Log_Returns'], lags=40)
plot_pacf(bitcoin_clean['Log_Returns'], lags=40)
plt.show()
```

Cross Validation

```
In [ ]: # Time Series Cross-Validation Implementation
# Novel 3-fold cross-validation scheme with rolling windows

import matplotlib.dates as mdates
from datetime import datetime, timedelta
from dateutil.relativedelta import relativedelta

# print("== TIME SERIES CROSS-VALIDATION IMPLEMENTATION ==\n")

def create_sp500_cv_splits(data, start_date=None):
```

```

if start_date is None:
    start_date = data.index.min()

cv_splits = []
window_start = start_date

while True:
    # Define window boundaries
    train_start = window_start
    train_end = train_start + relativedelta(years=3) - timedelta(days=1)

    # Validation periods (8, 16, 24 months)
    val_start = train_end + timedelta(days=1)
    val1_end = val_start + relativedelta(months=8) - timedelta(days=1) # 8 months
    val2_end = val_start + relativedelta(months=16) - timedelta(days=1) # 16 months
    val3_end = val_start + relativedelta(months=24) - timedelta(days=1) # 24 months (2 years)

    # Test period (1 year)
    test_start = val3_end + timedelta(days=1)
    test_end = test_start + relativedelta(years=1) - timedelta(days=1)

    # Check if we have enough data
    if test_end.year > 2024:
        break

    # Create splits for this window
    train_data = data[(data.index >= train_start) & (data.index <= train_end)]

    # Three validation folds
    val1_data = data[(data.index >= val_start) & (data.index <= val1_end)]
    val2_data = data[(data.index >= val_start) & (data.index <= val2_end)]
    val3_data = data[(data.index >= val_start) & (data.index <= val3_end)]

    test_data = data[(data.index >= test_start) & (data.index <= test_end)]

    cv_splits.append({
        'window_id': len(cv_splits) + 1,
        'train': {
            'data': train_data,
            'start': train_start,
            'end': train_end,
            'size': len(train_data)
        },
        'validation': [

```

```

        {
            'fold': 1,
            'data': val1_data,
            'start': val_start,
            'end': val1_end,
            'size': len(val1_data),
            'months': 8
        },
        {
            'fold': 2,
            'data': val2_data,
            'start': val_start,
            'end': val2_end,
            'size': len(val2_data),
            'months': 16
        },
        {
            'fold': 3,
            'data': val3_data,
            'start': val_start,
            'end': val3_end,
            'size': len(val3_data),
            'months': 24
        }
    ],
    'test': {
        'data': test_data,
        'start': test_start,
        'end': test_end,
        'size': len(test_data)
    }
}

# Move window forward by 1 year
window_start += relativedelta(years=1)

return cv_splits

```

def create_bitcoin_cv_splits(data, start_date=**None**):

if start_date **is** **None**:

Start from a date that allows for proper window construction

start_date = datetime(2015, 1, 1)

cv_splits = []

```
window_start = start_date

# Define the testing period constraint
test_period_start = datetime(2018, 1, 1)
test_period_end = datetime(2023, 12, 31)

while True:
    # Define window boundaries
    train_start = window_start
    train_end = train_start + relativedelta(years=2) - timedelta(days=1)

    # Validation periods (4, 8, 12 months)
    val_start = train_end + timedelta(days=1)
    val1_end = val_start + relativedelta(months=4) - timedelta(days=1) # 4 months
    val2_end = val_start + relativedelta(months=8) - timedelta(days=1) # 8 months
    val3_end = val_start + relativedelta(months=12) - timedelta(days=1) # 12 months

    # Test period (6 months)
    test_start = val3_end + timedelta(days=1)
    test_end = test_start + relativedelta(months=6) - timedelta(days=1)

    # Check constraints
    if test_end.year > 2023:
        break

    # Only include windows where test period is within 2018–2023
    if test_start < test_period_start:
        window_start += relativedelta(months=6)
        continue

    # Create splits for this window
    train_data = data[(data.index >= train_start) & (data.index <= train_end)]

    # Three validation folds
    val1_data = data[(data.index >= val_start) & (data.index <= val1_end)]
    val2_data = data[(data.index >= val_start) & (data.index <= val2_end)]
    val3_data = data[(data.index >= val_start) & (data.index <= val3_end)]

    test_data = data[(data.index >= test_start) & (data.index <= test_end)]

    cv_splits.append({
        'window_id': len(cv_splits) + 1,
        'train': {
            'data': train_data,
            'start': train_start,
```

```
        'end': train_end,
        'size': len(train_data)
    },
    'validation': [
        {
            'fold': 1,
            'data': val1_data,
            'start': val_start,
            'end': val1_end,
            'size': len(val1_data),
            'months': 4
        },
        {
            'fold': 2,
            'data': val2_data,
            'start': val_start,
            'end': val2_end,
            'size': len(val2_data),
            'months': 8
        },
        {
            'fold': 3,
            'data': val3_data,
            'start': val_start,
            'end': val3_end,
            'size': len(val3_data),
            'months': 12
        }
    ],
    'test': {
        'data': test_data,
        'start': test_start,
        'end': test_end,
        'size': len(test_data)
    }
}

# Move window forward by 6 months
window_start += relativedelta(months=6)

return cv_splits
```

In []: # Apply Cross-Validation Schemes to Data

```
print("==> GENERATING CROSS-VALIDATION SPLITS ==>\n")

# Generate S&P 500 cross-validation splits
sp500_cv_splits = create_sp500_cv_splits(sp500_clean)

# Generate Bitcoin cross-validation splits
bitcoin_cv_splits = create_bitcoin_cv_splits(bitcoin_clean)

# Create summary DataFrames
def create_cv_summary(cv_splits, asset_name):
    summary_data = []

    for split in cv_splits:
        # Add training data info
        summary_data.append({
            'Asset': asset_name,
            'Window_ID': split['window_id'],
            'Split_Type': 'Train',
            'Fold': 'N/A',
            'Start_Date': split['train']['start'].strftime('%Y-%m-%d'),
            'End_Date': split['train']['end'].strftime('%Y-%m-%d'),
            'Size': split['train']['size'],
            'Duration_Months': 'N/A'
        })

        # Add validation data info
        for val_fold in split['validation']:
            summary_data.append({
                'Asset': asset_name,
                'Window_ID': split['window_id'],
                'Split_Type': 'Validation',
                'Fold': val_fold['fold'],
                'Start_Date': val_fold['start'].strftime('%Y-%m-%d'),
                'End_Date': val_fold['end'].strftime('%Y-%m-%d'),
                'Size': val_fold['size'],
                'Duration_Months': val_fold['months']
            })

        # Add test data info
        summary_data.append({
            'Asset': asset_name,
            'Window_ID': split['window_id'],
            'Split_Type': 'Test',
```

```

        'Fold': 'N/A',
        'Start_Date': split['test']['start'].strftime('%Y-%m-%d'),
        'End_Date': split['test']['end'].strftime('%Y-%m-%d'),
        'Size': split['test']['size'],
        'Duration_Months': 'N/A'
    })

    return pd.DataFrame(summary_data)

# Create summary DataFrames
sp500_cv_summary = create_cv_summary(sp500_cv_splits, 'S&P_500')
bitcoin_cv_summary = create_cv_summary(bitcoin_cv_splits, 'Bitcoin')

# Combined summary
cv_summary_combined = pd.concat([sp500_cv_summary, bitcoin_cv_summary], ignore_index=True)

# Display first few windows for each asset
print("\n==== S&P 500 CV WINDOWS (First set) ===")
for i, split in enumerate(sp500_cv_splits[:1]):
    print(f"\nWindow {split['window_id']}:")
    print(f" Train: {split['train']['start'].strftime('%Y-%m-%d')} to {split['train']['end'].strftime('%Y-%m-%d')} ({split['train']['months']} months)")
    print(f" Validation Folds:")
    for val_fold in split['validation']:
        print(f"   Fold {val_fold['fold']} ({val_fold['months']} months): {val_fold['start'].strftime('%Y-%m-%d')} to {val_fold['end'].strftime('%Y-%m-%d')}")
    print(f" Test: {split['test']['start'].strftime('%Y-%m-%d')} to {split['test']['end'].strftime('%Y-%m-%d')} ({split['test']['months']} months)")

print("\n==== BITCOIN CV WINDOWS (First set) ===")
for i, split in enumerate(bitcoin_cv_splits[:1]):
    print(f"\nWindow {split['window_id']}:")
    print(f" Train: {split['train']['start'].strftime('%Y-%m-%d')} to {split['train']['end'].strftime('%Y-%m-%d')} ({split['train']['months']} months)")
    print(f" Validation Folds:")
    for val_fold in split['validation']:
        print(f"   Fold {val_fold['fold']} ({val_fold['months']} months): {val_fold['start'].strftime('%Y-%m-%d')} to {val_fold['end'].strftime('%Y-%m-%d')}")
    print(f" Test: {split['test']['start'].strftime('%Y-%m-%d')} to {split['test']['end'].strftime('%Y-%m-%d')} ({split['test']['months']} months)")

# Display summary statistics
print("\n==== CV SPLIT STATISTICS ===")
split_stats = cv_summary_combined.groupby(['Asset', 'Split_Type']).agg({
    'Size': ['mean', 'std', 'min', 'max'],
    'Window_ID': 'count'
}).round(0)
print(split_stats)

```

```
In [ ]: # Visualize Cross-Validation Scheme
```

```
def plot_cv_timeline(cv_splits, asset_name, max_windows=8):
    fig, ax = plt.subplots(figsize=(16, max(6, len(cv_splits[:max_windows]) * 1.5)))

    # Colors for different split types
    colors = {
        'train': '#2E8B57',
        'val_fold1': '#4169E1',
        'val_fold2': '#1E90FF',
        'val_fold3': '#87CEEB',
        'test': '#DC143C'
    }

    y_positions = []

    for i, split in enumerate(cv_splits[:max_windows]):
        y_pos = len(cv_splits[:max_windows]) - i - 1
        y_positions.append(y_pos)

        # Plot training period
        ax.barh(y_pos, (split['train']['end'] - split['train']['start']).days,
                left=split['train']['start'], height=0.6,
                color=colors['train'], alpha=0.8, label='Train' if i == 0 else "")

        # Plot validation periods
        val_colors = ['val_fold1', 'val_fold2', 'val_fold3']
        for j, val_fold in enumerate(split['validation']):
            ax.barh(y_pos + 0.1 + j*0.15, (val_fold['end'] - val_fold['start']).days,
                    left=val_fold['start'], height=0.12,
                    color=colors[val_colors[j]], alpha=0.8,
                    label=f'Val Fold {j+1} ({val_fold["months"]}mo)' if i == 0 else "")

        # Plot test period
        ax.barh(y_pos, (split['test']['end'] - split['test']['start']).days,
                left=split['test']['start'], height=0.6,
                color=colors['test'], alpha=0.8, label='Test' if i == 0 else "")

        # Add window labels
        ax.text(split['train']['start'], y_pos, f'W{split["window_id"]}',
                verticalalignment='center', fontsize=9, fontweight='bold')

    # Formatting
    ax.set_ylim(-0.5, len(cv_splits[:max_windows]) - 0.5)
```

```

    ax.set_ylabel('CV Windows (Newest to Oldest)', fontsize=12)
    ax.set_xlabel('Time Period', fontsize=12)
    ax.set_title(f'{asset_name} Cross-Validation Timeline\n{len(cv_splits)} Total Windows, Showing First {min(max_windows, len(cv_splits))} windows', fontsize=14, fontweight='bold')

    # Format x-axis
    ax.xaxis.set_major_formatter(mdates.DateFormatter('%Y-%m'))
    ax.xaxis.set_major_locator(mdates.YearLocator())
    plt.xticks(rotation=45)

    # Add legend
    ax.legend(loc='upper right', bbox_to_anchor=(1, 1), frameon=True, fancybox=True, shadow=True)

    # Add grid
    ax.grid(True, alpha=0.3, axis='x')

    plt.tight_layout()
    return fig, ax

print("== CROSS-VALIDATION SCHEME VISUALIZATION ==\n")

def create_cv_comparison_chart():
    fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(20, 8))

    # S&P 500 scheme visualization
    y_pos = 1

    # S&P 500 scheme
    ax1.barh(y_pos, 3*365, left=0, height=0.6, color="#2E8B57", alpha=0.8, label='Train (3yr)')
    ax1.barh(y_pos+0.1, 8*30, left=3*365, height=0.15, color="#4169E1", alpha=0.8, label='Val Fold 1 (8mo)')
    ax1.barh(y_pos+0.25, 16*30, left=3*365, height=0.15, color="#1E90FF", alpha=0.8, label='Val Fold 2 (16mo)')
    ax1.barh(y_pos+0.4, 24*30, left=3*365, height=0.15, color="#87CEEB", alpha=0.8, label='Val Fold 3 (24mo)')
    ax1.barh(y_pos, 1*365, left=5*365, height=0.6, color="#DC143C", alpha=0.8, label='Test (1yr)')

    ax1.set_xlim(0, 6*365)
    ax1.set_ylim(0.5, 1.8)
    ax1.set_xlabel('Days')
    ax1.set_title('S&P 500 CV Scheme\n(6-year windows)', fontweight='bold')
    ax1.legend()
    ax1.grid(True, alpha=0.3)

    # Bitcoin scheme
    ax2.barh(y_pos, 2*365, left=0, height=0.6, color="#2E8B57", alpha=0.8, label='Train (2yr)')
    ax2.barh(y_pos+0.1, 4*30, left=2*365, height=0.15, color="#4169E1", alpha=0.8, label='Val Fold 1 (4mo)')
    ax2.barh(y_pos+0.25, 8*30, left=2*365, height=0.15, color="#1E90FF", alpha=0.8, label='Val Fold 2 (8mo)')

```

```

ax2.barh(y_pos+0.4, 12*30, left=2*365, height=0.15, color='#87CEEB', alpha=0.8, label='Val Fold 3 (12mo)')
ax2.barh(y_pos, 6*30, left=3*365, height=0.6, color='#DC143C', alpha=0.8, label='Test (6mo)')

ax2.set_xlim(0, 3.5*365)
ax2.set_ylim(0.5, 1.8)
ax2.set_xlabel('Days')
ax2.set_title('Bitcoin CV Scheme\n(~3.5-year windows)', fontweight='bold')
ax2.legend()
ax2.grid(True, alpha=0.3)

plt.suptitle('Time Series Cross-Validation Scheme Comparison', fontsize=16, fontweight='bold')
plt.tight_layout()
return fig

# fig3 = create_cv_comparison_chart()
# plt.show()

# Create summary table
cv_scheme_summary = pd.DataFrame({
    'Asset': ['S&P 500', 'Bitcoin'],
    'Total_Windows': [len(sp500_cv_splits), len(bitcoin_cv_splits)],
    'Window_Length': ['6 years', '~3.5 years'],
    'Train_Period': ['3 years', '2 years'],
    'Validation_Folds': ['8/16/24 months', '4/8/12 months'],
    'Test_Period': ['1 year', '6 months'],
    'Window_Shift': ['1 year', '6 months'],
    'Test_Coverage': [
        f'{sp500_cv_splits[0]["test"]["start"].strftime("%Y-%m-%d")}' + ' to ' + {sp500_cv_splits[-1]["test"]["end"].strftime("%Y-%m-%d")},
        f'{bitcoin_cv_splits[0]["test"]["start"].strftime("%Y-%m-%d")}' + ' to ' + {bitcoin_cv_splits[-1]["test"]["end"].strftime("%Y-%m-%d")}
    ]
})

```

In []: # Utility Functions for Cross-Validation Data Access

```

def get_cv_data(cv_splits, window_id, fold=None, return_type='data'):

    split = next((s for s in cv_splits if s['window_id'] == window_id), None)
    if split is None:
        raise ValueError(f"Window ID {window_id} not found")

    if return_type == 'train':
        return split['train']['data']
    elif return_type == 'test':
        return split['test']['data']

```

```

    elif return_type == 'validation':
        if fold is None:
            raise ValueError("Fold number must be specified for validation data")
        if fold not in [1, 2, 3]:
            raise ValueError("Fold must be 1, 2, or 3")
        return split['validation'][fold-1]['data']
    else:
        return split

# Example usage functions
def demonstrate_cv_usage():
    print("== CROSS-VALIDATION DATA ACCESS EXAMPLES ==\n")

    # Example 1: Get training data from first S&P 500 window
    train_data_sp500 = get_cv_data(sp500_cv_splits, window_id=1, return_type='train')
    print(f"\nS&P 500 Window 1 - Training data shape: {train_data_sp500.shape}")
    print(f"Training period: {train_data_sp500.index.min()} to {train_data_sp500.index.max()}\n")

    # Example 2: Get validation fold 2 data from first S&P 500 window
    val_data_sp500 = get_cv_data(sp500_cv_splits, window_id=1, fold=2, return_type='validation')
    print(f"\nS&P 500 Window 1 - Validation Fold 2 shape: {val_data_sp500.shape}")
    print(f"Validation period: {val_data_sp500.index.min()} to {val_data_sp500.index.max()}\n")

    # Example 3: Get test data from first Bitcoin window
    test_data_bitcoin = get_cv_data(bitcoin_cv_splits, window_id=1, return_type='test')
    print(f"\nBitcoin Window 1 - Test data shape: {test_data_bitcoin.shape}")
    print(f"Test period: {test_data_bitcoin.index.min()} to {test_data_bitcoin.index.max()}\n")

    return train_data_sp500, val_data_sp500, test_data_bitcoin

# Run demonstration
sample_train, sample_val, sample_test = demonstrate_cv_usage()

# Save CV splits for later use (optional)
cv_implementation_summary = {
    'sp500_cv_splits': sp500_cv_splits,
    'bitcoin_cv_splits': bitcoin_cv_splits,
    'sp500_summary': sp500_cv_summary,
    'bitcoin_summary': bitcoin_cv_summary,
    'combined_summary': cv_summary_combined,
    'scheme_comparison': cv_scheme_summary
}

```

ARIMA-LSTM HYBRID METHOD 2 (Non-additive)

ARIMA

```
In [ ]: # ARIMA Model Implementation with AIC-based Selection
# Modern approach replacing traditional Box-Jenkins methodology

import itertools
from statsmodels.tsa.arima.model import ARIMA
from statsmodels.tsa.stattools import adfuller, kpss
from statsmodels.stats.diagnostic import acorr_ljungbox
import warnings
from sklearn.metrics import mean_squared_error, mean_absolute_error
from scipy import stats
import time

# Suppress convergence warnings for cleaner output
warnings.filterwarnings('ignore', category=UserWarning)
warnings.filterwarnings('ignore', category=RuntimeWarning)

def find_optimal_arima_order(data, max_p=5, max_d=2, max_q=5, seasonal=False,
                             information_criterion='aic', verbose=False):

    best_ic = np.inf
    best_params = None
    best_model = None
    results_log = []

    # Create parameter grid
    if seasonal:
        # For seasonal ARIMA (not implemented in this study)
        param_grid = itertools.product(range(max_p+1), range(max_d+1), range(max_q+1),
                                       range(2), range(2), range(2), [12])
    else:
        # Standard ARIMA grid search
        param_grid = itertools.product(range(max_p+1), range(max_d+1), range(max_q+1))

    total_combinations = (max_p+1) * (max_d+1) * (max_q+1)

    if verbose:
```

```

print(f"Testing {total_combinations} ARIMA parameter combinations...")
start_time = time.time()

for i, params in enumerate(param_grid):
    p, d, q = params[:3]

    # Skip if model is too simple (all parameters zero)
    if p == 0 and d == 0 and q == 0:
        continue

    try:
        # Fit ARIMA model
        model = ARIMA(data, order=(p, d, q))
        fitted_model = model.fit()

        # Get information criterion value
        if information_criterion.lower() == 'aic':
            ic_value = fitted_model.aic
        elif information_criterion.lower() == 'bic':
            ic_value = fitted_model.bic
        elif information_criterion.lower() == 'hqic':
            ic_value = fitted_model.hqic
        else:
            ic_value = fitted_model.aic

        # Store results
        results_log.append({
            'order': (p, d, q),
            'aic': fitted_model.aic,
            'bic': fitted_model.bic,
            'hqic': fitted_model.hqic,
            'llf': fitted_model.llf,
            'converged': fitted_model.mle_retvals['converged'] if hasattr(fitted_model, 'mle_retvals') else True
        })

        # Update best model if current is better
        if ic_value < best_ic:
            best_ic = ic_value
            best_params = (p, d, q)
            best_model = fitted_model

    except Exception as e:
        # Log failed fits
        results_log.append({
            'order': (p, d, q),

```

```

        'aic': np.nan,
        'bic': np.nan,
        'hqic': np.nan,
        'llf': np.nan,
        'converged': False,
        'error': str(e)
    })

    if verbose and i % 10 == 0:
        print(f"Failed to fit ARIMA{params}: {str(e)[:50]}...")

    if verbose and (i + 1) % 20 == 0:
        elapsed = time.time() - start_time
        progress = (i + 1) / total_combinations * 100
        print(f"Progress: {progress:.1f}% ({i+1}/{total_combinations}) | "
              f"Best so far: ARIMA{best_params} ({information_criterion.upper()}={best_ic:.4f})")

    if verbose:
        total_time = time.time() - start_time
        print(f"\nGrid search completed in {total_time:.2f} seconds")
        print(f"Best model: ARIMA{best_params} with {information_criterion.upper()}={best_ic:.4f}")

# Create results summary
results_df = pd.DataFrame(results_log)
successful_fits = results_df[results_df['converged'] == True]

return {
    'best_order': best_params,
    'best_model': best_model,
    'best_ic_value': best_ic,
    'information_criterion': information_criterion,
    'results_df': results_df,
    'successful_fits': len(successful_fits),
    'total_attempts': len(results_log),
    'success_rate': len(successful_fits) / len(results_log) * 100
}

def evaluate_arima_model(model, train_data, test_data, model_order):
    # Generate forecasts
    n_forecast = len(test_data)
    forecast_result = model.get_forecast(steps=n_forecast)
    forecasts = forecast_result.predicted_mean
    forecast_ci = forecast_result.conf_int()

    # Calculate performance metrics

```

```

mse = mean_squared_error(test_data, forecasts)
rmse = np.sqrt(mse)
mae = mean_absolute_error(test_data, forecasts)
mape = np.mean(np.abs((test_data - forecasts) / test_data)) * 100

# Calculate R2 score for fair comparison with LSTM and SVM
ss_res = np.sum((test_data - forecasts) ** 2)
ss_tot = np.sum((test_data - np.mean(test_data)) ** 2)
r2 = 1 - (ss_res / ss_tot) if ss_tot > 0 else 0

# Direction accuracy (for returns)
direction_actual = np.sign(test_data.values[1:])
direction_forecast = np.sign(forecasts.values[1:])
direction_accuracy = np.mean(direction_actual == direction_forecast) * 100

# Residual diagnostics
residuals = model.resid

# Ljung-Box test for serial correlation in residuals
lb_test = acorr_ljungbox(residuals, lags=10, return_df=False)

# Normality test (Jarque-Bera)
jb_stat, jb_pvalue = stats.jarque_bera(residuals)

# Heteroskedasticity test (simple approach)
residuals_squared = residuals ** 2
arch_stat, arch_pvalue = acorr_ljungbox(residuals_squared, lags=5, return_df=False)

return {
    'model_order': model_order,
    'forecasts': forecasts,
    'forecast_ci': forecast_ci,
    'performance_metrics': {
        'mse': mse,
        'rmse': rmse,
        'mae': mae,
        'mape': mape,
        'r2': r2,
        'direction_accuracy': direction_accuracy
    },
    'diagnostic_tests': {
        'ljung_box_stat': lb_test['lb_stat'].iloc[-1],
        'ljung_box_pvalue': lb_test['lb_pvalue'].iloc[-1],
        'jarque_bera_stat': jb_stat,
        'jarque_bera_pvalue': jb_pvalue,
    }
}

```

```

        'arch_stat': arch_stat[-1] if isinstance(arch_stat, np.ndarray) else arch_stat,
        'arch_pvalue': arch_pvalue[-1] if isinstance(arch_pvalue, np.ndarray) else arch_pvalue
    },
    'residuals': residuals
}

```

HYBRID ARIMA-LSTM MODEL - METHOD 2 (Non-additive)

```

In [ ]: # Cross-Validation Integration for ARIMA Model Selection
# Implement the complete methodology with hyperparameter optimization

def run_arima_cross_validation(cv_splits, data_clean, asset_name, max_p=3, max_d=2, max_q=3,
                               information_criterion='aic', verbose=True):
    print(f"\n{asset_name.upper()} ARIMA CROSS-VALIDATION ===")
    print(f"Running AIC-based model selection across {len(cv_splits)} windows...")
    print(f"Parameter search space: p∈[0,{max_p}], d∈[0,{max_d}], q∈[0,{max_q}]")
    print("-" * 80)

    all_results = []
    model_selection_summary = []

    for window_idx, split in enumerate(cv_splits):
        window_id = split['window_id']

        if verbose:
            print(f"\nProcessing Window {window_id}/{len(cv_splits)}...")
            print(f"  Train: {split['train']['start'].strftime('%Y-%m-%d')} to {split['train']['end'].strftime('%Y-%m-%d')}")
            print(f"  Test:  {split['test']['start'].strftime('%Y-%m-%d')} to {split['test']['end'].strftime('%Y-%m-%d')}")
            print(f"{'-' * 80}\n")

        # Extract data
        train_data = split['train']['data']['Log_Returns']
        test_data = split['test']['data']['Log_Returns']

        # STEP 1: Model Selection using Training Data
        if verbose:
            print(f"  Model selection using {information_criterion.upper()} criterion...")

        selection_result = find_optimal_arima_order(
            train_data,
            max_p=max_p,
            max_d=max_d,
            max_q=max_q,
            information_criterion=information_criterion,

```

```
    verbose=False # Keep individual window selection quiet
)

if selection_result['best_model'] is None:
    print(f"    Failed to find suitable model for Window {window_id}")
    continue

best_order = selection_result['best_order']

# STEP 2: Hyperparameter Validation using Validation Folds
if verbose:
    print(f"    📈 Validating ARIMA{best_order} across 3 validation folds...")

validation_scores = []

for val_fold in split['validation']:
    fold_num = val_fold['fold']
    val_data = val_fold['data']['Log_Returns']

    try:
        # Fit model on training data and evaluate on validation fold
        val_model = ARIMA(train_data, order=best_order).fit()
        val_forecasts = val_model.get_forecast(steps=len(val_data)).predicted_mean
        val_rmse = np.sqrt(mean_squared_error(val_data, val_forecasts))
        validation_scores.append(val_rmse)

    except Exception as e:
        if verbose:
            print(f"    Validation fold {fold_num} failed: {str(e)[:50]}...")
        validation_scores.append(np.inf)

avg_validation_rmse = np.mean(validation_scores)

# STEP 3: Final Model Training and Out-of-Sample Evaluation
if verbose:
    print(f"    Final evaluation on test data...")

try:
    # Re-fit the model on training data
    final_model = ARIMA(train_data, order=best_order).fit()

    # Evaluate on test data
    evaluation = evaluate_arima_model(final_model, train_data, test_data, best_order)

    # Store comprehensive results
```

```

window_result = {
    'window_id': window_id,
    'asset': asset_name,
    'train_period': f'{split["train"]['start'].strftime("%Y-%m-%d")} to {split["train"]['end'].strftime("%Y-%m-%d")}',
    'test_period': f'{split["test"]['start'].strftime("%Y-%m-%d")} to {split["test"]['end'].strftime("%Y-%m-%d")}',
    'train_size': split['train']['size'],
    'test_size': split['test']['size'],
    'best_order': best_order,
    'model_selection': selection_result,
    'validation_scores': validation_scores,
    'avg_validation_rmse': avg_validation_rmse,
    'evaluation': evaluation,
    'final_model': final_model
}

all_results.append(window_result)

# Summary for quick reference
model_selection_summary.append({
    'Window': window_id,
    'Best_Order': f"ARIMA{best_order}",
    'AIC': selection_result['best_ic_value'],
    'Validation_RMSE': avg_validation_rmse,
    'Test_RMSE': evaluation['performance_metrics']['rmse'],
    'Test_MAE': evaluation['performance_metrics']['mae'],
    'Test_R2': evaluation['performance_metrics']['r2'],
    'Direction_Accuracy': evaluation['performance_metrics']['direction_accuracy'],
    'Ljung_Box_p': evaluation['diagnostic_tests']['ljung_box_pvalue']
})

if verbose:
    print(f"    ARIMA{best_order}: Test RMSE={evaluation['performance_metrics']['rmse']:.6f}, "
          f"Direction Acc={evaluation['performance_metrics']['direction_accuracy']:.1f}%")

except Exception as e:
    print(f"    Final evaluation failed for Window {window_id}: {str(e)}")
    continue

# Create summary DataFrame
summary_df = pd.DataFrame(model_selection_summary)

# Calculate overall performance statistics
if len(summary_df) > 0:
    performance_summary = {
        'total_windows': len(cv_splits),

```

```

'successful_windows': len(summary_df),
'success_rate': len(summary_df) / len(cv_splits) * 100,
'avg_test_rmse': summary_df['Test_RMSE'].mean(),
'std_test_rmse': summary_df['Test_RMSE'].std(),
'avg_test_mae': summary_df['Test_MAE'].mean(),
'avg_r2': summary_df['Test_R2'].mean(),
'avg_direction_accuracy': summary_df['Direction_Accuracy'].mean(),
'avg_validation_rmse': summary_df['Validation_RMSE'].mean(),
'most_common_order': summary_df['Best_Order'].mode().iloc[0] if len(summary_df) > 0 else None
}
else:
    performance_summary = None

print(f"\n{'='*80}")
print(f"{asset_name.upper()} ARIMA CROSS-VALIDATION COMPLETE")
print(f"\n{'='*80}")

if performance_summary:
    print(f" Successfully processed {performance_summary['successful_windows']}/{performance_summary['total_windows']} windows")
    print(f" Average Test RMSE: {performance_summary['avg_test_rmse']:.6f} ± {performance_summary['std_test_rmse']:.6f}")
    print(f" Average Direction Accuracy: {performance_summary['avg_direction_accuracy']:.2f}%")
    print(f" Most Common Model: {performance_summary['most_common_order']}")
else:
    print("No successful model fits achieved")

return {
    'asset_name': asset_name,
    'all_results': all_results,
    'summary_df': summary_df,
    'performance_summary': performance_summary,
    'methodology': {
        'approach': 'AIC-based automated selection',
        'information_criterion': information_criterion,
        'parameter_space': f'p∈[0,{max_p}], d∈[0,{max_d}], q∈[0,{max_q}]',
        'cross_validation': '3-fold temporal validation',
        'evaluation_metric': 'Out-of-sample RMSE and direction accuracy'
    }
}
}

```

```
In [ ]: import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense, Dropout
from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau
from sklearn.preprocessing import MinMaxScaler
```

```

from sklearn.metrics import mean_squared_error, mean_absolute_error

def create_sequences_with_arima_pred(original_data, arima_predictions, lookback=60):
    X, y = [], []

    for i in range(lookback, len(original_data)):
        # Lagged observations: yt-lookback to yt-1
        lagged_obs = original_data[i-lookback:i]

        # ARIMA predictions: includes prediction at time t
        # We use ARIMA preds from i-lookback to i (includes current timestep)
        arima_features = arima_predictions[i-lookback:i]

        # Stack features: each timestep has [observation, arima_pred]
        features = np.column_stack([lagged_obs, arima_features])

        X.append(features)
        y.append(original_data[i])

    return np.array(X), np.array(y)

def build_hybrid_lstm_model(lookback=60, n_features=2, units=50, dropout=0.2, learning_rate=0.001):
    model = Sequential([
        LSTM(units=units, return_sequences=True, input_shape=(lookback, n_features)),
        Dropout(dropout),
        LSTM(units=units, return_sequences=False),
        Dropout(dropout),
        Dense(units=1)
    ])

    optimizer = tf.keras.optimizers.Adam(learning_rate=learning_rate)
    model.compile(optimizer=optimizer, loss='mean_squared_error', metrics=['mae'])

    return model

def run_hybrid_arima_lstm_cv(cv_splits, data_clean, asset_name,
                             max_p=3, max_d=1, max_q=3,
                             lookback=60, lstm_units=50, dropout=0.2,
                             epochs=50, batch_size=32, verbose=True):

    print(f"\n{'='*100}")
    print(f"HYBRID ARIMA-LSTM CROSS-VALIDATION: {asset_name.upper()}")
    print(f"{'='*100}")

```



```

        continue

    if best_model is None:
        print(f"    Failed to fit ARIMA for window {window_id}")
        continue

    if verbose:
        print(f"        Best ARIMA{best_order}, AIC={best_aic:.2f}")

# =====
# STEP 2: GET ARIMA FITTED VALUES (In-sample predictions)
# =====
    if verbose:
        print(f"    [2/4] Extracting ARIMA fitted values...")

# Get in-sample predictions (fitted values)
arima_train_pred = best_model.fittedvalues

# Align arrays (ARIMA might drop initial values)
if len(arima_train_pred) < len(train_data):
    n_dropped = len(train_data) - len(arima_train_pred)
    train_data_aligned = train_data[n_dropped:]
    arima_train_pred_aligned = arima_train_pred
else:
    train_data_aligned = train_data
    arima_train_pred_aligned = arima_train_pred

if verbose:
    print(f"    ARIMA predictions: mean={np.mean(arima_train_pred_aligned):.6f}, std={np.std(arima_train_pred_aligned):.6f}")

# =====
# STEP 3: TRAIN LSTM WITH LAGGED OBS + ARIMA PREDICTIONS
# =====
    if len(train_data_aligned) <= lookback:
        print(f"    Insufficient data for LSTM (need > {lookback} points)")
        continue

    if verbose:
        print(f"    [3/4] Training LSTM with lagged observations + ARIMA predictions...")

# Scale both original data and ARIMA predictions
train_data_reshaped = train_data_aligned.reshape(-1, 1)
arima_pred_reshaped = arima_train_pred_aligned.reshape(-1, 1)

# Use separate scalers for data and ARIMA predictions

```

```

data_scaler = MinMaxScaler(feature_range=(0, 1))
arima_scaler = MinMaxScaler(feature_range=(0, 1))

scaled_train_data = data_scaler.fit_transform(train_data_reshaped)
scaled_arima_pred = arima_scaler.fit_transform(arima_pred_reshaped)

# Create sequences with both lagged observations and ARIMA predictions
X_train, y_train = create_sequences_with_arima_pred(
    scaled_train_data.flatten(),
    scaled_arima_pred.flatten(),
    lookback
)

if len(X_train) == 0:
    print(f"  No sequences created")
    continue

# X_train already has shape (n_samples, lookback, 2) from create_sequences_with_arima_pred

# Build and train LSTM
lstm_model = build_hybrid_lstm_model(lookback=lookback, n_features=2, units=lstm_units, dropout=dropout)

early_stop = EarlyStopping(monitor='loss', patience=5, restore_best_weights=True, verbose=0)
reduce_lr = ReduceLROnPlateau(monitor='loss', factor=0.5, patience=3, min_lr=1e-6, verbose=0)

history = lstm_model.fit(
    X_train, y_train,
    epochs=epochs,
    batch_size=batch_size,
    callbacks=[early_stop, reduce_lr],
    verbose=0,
    validation_split=0.1
)

if verbose:
    print(f"      LSTM trained ({len(history.history['loss'])} epochs)")

# =====
# STEP 4: HYBRID FORECASTING ON TEST SET (Method 2: Non-additive)
# =====

if verbose:
    print(f"  [4/4] Generating hybrid forecasts...")

# Make ARIMA forecasts for test period
arima_test_pred = best_model.forecast(steps=len(test_data))

```

```

# For LSTM prediction, we need:
# 1. Last 'lookback' observations from training
# 2. ARIMA prediction at each timestep

# Get last lookback observations from training (scaled)
last_observations = scaled_train_data[-lookback:].flatten()
hybrid_predictions = []

for t in range(len(test_data)):
    arima_pred_t = arima_scaler.transform([[arima_test_pred[t]]])[0, 0]

    obs_sequence = last_observations[-lookback:]

    if t == 0:
        # First prediction: use training ARIMA predictions
        arima_sequence = scaled_arima_pred[-lookback:].flatten()
    else:
        # Subsequent predictions: use recent ARIMA predictions
        arima_sequence = np.append(arima_sequence[1:], arima_pred_t)

    # Stack features: shape (lookback, 2)
    features = np.column_stack([obs_sequence, arima_sequence])
    X_input = features.reshape(1, lookback, 2)

    # Predict final value (LSTM directly outputs the forecast)
    pred_scaled = lstm_model.predict(X_input, verbose=0)
    pred = data_scaler.inverse_transform(pred_scaled)[0, 0]
    hybrid_predictions.append(pred)

    # Update observation sequence with actual value for next iteration
    if t < len(test_data) - 1:
        actual_scaled = data_scaler.transform([[test_data[t]]])[0, 0]
        last_observations = np.append(last_observations[1:], actual_scaled)

hybrid_predictions = np.array(hybrid_predictions)

# Metrics for ARIMA only
arima_rmse = np.sqrt(mean_squared_error(test_data, arima_test_pred))
arima_mae = mean_absolute_error(test_data, arima_test_pred)

# Metrics for HYBRID
hybrid_rmse = np.sqrt(mean_squared_error(test_data, hybrid_predictions))
hybrid_mae = mean_absolute_error(test_data, hybrid_predictions)

```

```

# Direction accuracy
arima_direction = np.mean(np.sign(test_data[1:]) == np.sign(arima_test_pred[1:])) * 100
hybrid_direction = np.mean(np.sign(test_data[1:]) == np.sign(hybrid_predictions[1:])) * 100

improvement = ((arima_rmse - hybrid_rmse) / arima_rmse) * 100

if verbose:
    print(f"    RESULTS:")
    print(f"        ARIMA only: RMSE={arima_rmse:.6f}, MAE={arima_mae:.6f}, Direction={arima_direction:.1f}%")
    print(f"        HYBRID: RMSE={hybrid_rmse:.6f}, MAE={hybrid_mae:.6f}, Direction={hybrid_direction:.1f}%")
    print(f"        Improvement: {improvement:+.2f}% RMSE")

# Store results
window_results = {
    'window_id': window_id,
    'train_start': split['train']['start'],
    'train_end': split['train']['end'],
    'test_start': split['test']['start'],
    'test_end': split['test']['end'],
    'arima_order': best_order,
    'arima_aic': best_aic,
    # ARIMA metrics
    'arima_rmse': arima_rmse,
    'arima_mae': arima_mae,
    'arima_direction_accuracy': arima_direction,
    'arima_predictions': arima_test_pred,
    # Note: No separate LSTM residual predictions in Method 2 (non-additive)
    # HYBRID metrics
    'hybrid_rmse': hybrid_rmse,
    'hybrid_mae': hybrid_mae,
    'hybrid_direction_accuracy': hybrid_direction,
    'hybrid_predictions': hybrid_predictions,
    # Comparison
    'rmse_improvement_pct': improvement,
    'actuals': test_data
}
all_results.append(window_results)

except Exception as e:
    print(f"Error in window {window_id}: {str(e)}")
    import traceback
    traceback.print_exc()
    continue

```

```

if len(all_results) > 0:
    avg_arima_rmse = np.mean([r['arima_rmse'] for r in all_results])
    avg_hybrid_rmse = np.mean([r['hybrid_rmse'] for r in all_results])
    avg_improvement = np.mean([r['rmse_improvement_pct'] for r in all_results])

    avg_arima_mae = np.mean([r['arima_mae'] for r in all_results])
    avg_hybrid_mae = np.mean([r['hybrid_mae'] for r in all_results])

    avg_arima_direction = np.mean([r['arima_direction_accuracy'] for r in all_results])
    avg_hybrid_direction = np.mean([r['hybrid_direction_accuracy'] for r in all_results])

print(f"\n{'='*100}")
print(f"{asset_name.upper()} HYBRID MODEL COMPLETE")
print(f"\n{'='*100}")
print(f"Windows processed: {len(all_results)}/{len(cv_splits)}")
print(f"\nAVERAGE PERFORMANCE:")
print(f"    ARIMA only: RMSE={avg_arima_rmse:.6f}, MAE={avg_arima_mae:.6f}, Direction={avg_arima_direction:.2f}%")
print(f"    HYBRID:      RMSE={avg_hybrid_rmse:.6f}, MAE={avg_hybrid_mae:.6f}, Direction={avg_hybrid_direction:.2f}%")
print(f"    Improvement: {avg_improvement:+.2f}% RMSE")
print(f"\n{'='*100}")

results_dict = {
    'asset_name': asset_name,
    'model_type': 'HYBRID_ARIMA_LSTM',
    'windows_processed': len(all_results),
    'total_windows': len(cv_splits),
    # ARIMA metrics
    'avg_arima_rmse': avg_arima_rmse,
    'avg_arima_direction': avg_arima_direction,
    'avg_arima_mae': avg_arima_mae,
    # HYBRID metrics
    'avg_hybrid_rmse': avg_hybrid_rmse,
    'avg_hybrid_direction': avg_hybrid_direction,
    'avg_hybrid_mae': avg_hybrid_mae,
    # Improvement
    'avg_improvement_pct': avg_improvement,
    'window_results': all_results,
    'hyperparameters': {
        'arima': f'p∈[0,{max_p}], d∈[0,{max_d}], q∈[0,{max_q}]',
        'lstm_lookback': lookback,
        'lstm_units': lstm_units,
        'dropout': dropout,
        'epochs': epochs,
        'batch_size': batch_size
}

```

```
        }

    return results_dict
else:
    print(f"\n No windows successfully processed for {asset_name}")
    return None
```

```
In [ ]: print("\nPHASE 1: S&P 500 HYBRID ARIMA-LSTM")

sp500_hybrid_results = run_hybrid_arima_lstm_cv(
    cv_splits=sp500_cv_splits,
    data_clean=sp500_clean,
    asset_name='S&P 500',
    max_p=3,
    max_d=1,
    max_q=3,
    lookback=60,
    lstm_units=50,
    dropout=0.2,
    epochs=50,
    batch_size=32,
    verbose=True
)

print("\n\nPHASE 2: BITCOIN HYBRID ARIMA-LSTM")

bitcoin_hybrid_results = run_hybrid_arima_lstm_cv(
    cv_splits=bitcoin_cv_splits,
    data_clean=bitcoin_clean,
    asset_name='Bitcoin',
    max_p=3,
    max_d=1,
    max_q=3,
    lookback=60,
    lstm_units=50,
    dropout=0.2,
    epochs=50,
    batch_size=32,
    verbose=True
)

print("\n" + "=" * 100)
print("FINAL SUMMARY: HYBRID MODEL PERFORMANCE ACROSS BOTH ASSETS")
```

```

print("=" * 100)

summary_data = []

if sp500_hybrid_results:
    summary_data.append({
        'Asset': 'S&P 500',
        'ARIMA_RMSE': sp500_hybrid_results['avg_arima_rmse'],
        'Hybrid_RMSE': sp500_hybrid_results['avg_hybrid_rmse'],
        'Improvement_%': sp500_hybrid_results['avg_improvement_pct'],
        'ARIMA_MAE': sp500_hybrid_results['avg_arima_mae'],
        'Hybrid_MAE': sp500_hybrid_results['avg_hybrid_mae'],
        'ARIMA_Direction_%': sp500_hybrid_results['avg_arima_direction'],
        'Hybrid_Direction_%': sp500_hybrid_results['avg_hybrid_direction'],
        'Windows': f"{sp500_hybrid_results['windows_processed']}/{sp500_hybrid_results['total_windows']}"
    })

if bitcoin_hybrid_results:
    summary_data.append({
        'Asset': 'Bitcoin',
        'ARIMA_RMSE': bitcoin_hybrid_results['avg_arima_rmse'],
        'Hybrid_RMSE': bitcoin_hybrid_results['avg_hybrid_rmse'],
        'Improvement_%': bitcoin_hybrid_results['avg_improvement_pct'],
        'ARIMA_MAE': bitcoin_hybrid_results['avg_arima_mae'],
        'Hybrid_MAE': bitcoin_hybrid_results['avg_hybrid_mae'],
        'ARIMA_Direction_%': bitcoin_hybrid_results['avg_arima_direction'],
        'Hybrid_Direction_%': bitcoin_hybrid_results['avg_hybrid_direction'],
        'Windows': f"{bitcoin_hybrid_results['windows_processed']}/{bitcoin_hybrid_results['total_windows']}"
    })

if summary_data:
    summary_df = pd.DataFrame(summary_data)
    print("\n", summary_df.to_string(index=False))

    print("\n" + "=" * 100)
    print("KEY INSIGHTS:")
    print("=" * 100)

    for data in summary_data:
        asset = data['Asset']
        improvement = data['Improvement_%']

        if improvement > 0:
            print(f"{asset}: Hybrid model OUTPERFORMS ARIMA by {improvement:.2f}% RMSE")
        elif improvement < 0:

```

```
        print(f'{asset}: Hybrid model underperforms ARIMA by {abs(improvement):.2f}% RMSE')
    else:
        print(f'= {asset}: Hybrid model equals ARIMA performance')
```

```
In [ ]: def plot_hybrid_decomposition(results, window_idx=0):
    if not results or len(results['window_results']) == 0:
        print("No results to plot")
        return

    window = results['window_results'][window_idx]
    asset_name = results['asset_name']

    # Extract data
    actuals = window['actuals']
    arima_pred = window['arima_predictions']
    hybrid_pred = window['hybrid_predictions']

    # Create figure with subplots
    fig, axes = plt.subplots(3, 2, figsize=(16, 12))
    fig.suptitle(f'{asset_name} - Hybrid Method 2 (Non-additive) Analysis (Window {window["window_id"]})',
                 fontsize=16, fontweight='bold')

    # Plot 1: Actual vs ARIMA
    ax1 = axes[0, 0]
    ax1.plot(actuals, label='Actual', color='black', linewidth=2, alpha=0.7)
    ax1.plot(arima_pred, label='ARIMA', color='blue', linewidth=1.5, linestyle='--', alpha=0.7)
    ax1.set_title('ARIMA Component (Linear)', fontsize=12, fontweight='bold')
    ax1.set_ylabel('Log Returns')
    ax1.legend()
    ax1.grid(True, alpha=0.3)
    ax1.text(0.02, 0.98, f'RMSE: {window["arima_rmse"]:.6f}',
             transform=ax1.transAxes, verticalalignment='top',
             bbox=dict(boxstyle='round', facecolor='wheat', alpha=0.5))

    # Plot 2: Hybrid vs ARIMA Predictions
    ax2 = axes[0, 1]
    ax2.plot(arima_pred, label='ARIMA', color='blue', linewidth=1.5, alpha=0.7)
    ax2.plot(hybrid_pred, label='Hybrid (Method 2)', color='purple', linewidth=1.5, linestyle='--', alpha=0.7)
    ax2.set_title('ARIMA vs Hybrid Predictions', fontsize=12, fontweight='bold')
    ax2.set_ylabel('Log Returns')
    ax2.legend()
    ax2.grid(True, alpha=0.3)
    ax2.text(0.02, 0.98, f'ARIMA RMSE: {window["arima_rmse"]:.6f}\nHybrid RMSE: {window["hybrid_rmse"]:.6f}',
             transform=ax2.transAxes, verticalalignment='top',
```

```

bbox=dict(boxstyle='round', facecolor='wheat', alpha=0.5))

# Plot 3: ARIMA Prediction Errors
ax3 = axes[1, 0]
arima_residuals = actuals - arima_pred
ax3.plot(arima_residuals, label='ARIMA Errors', color='red', linewidth=1.5, alpha=0.7)
ax3.axhline(y=0, color='black', linestyle='-', linewidth=0.5)
ax3.fill_between(range(len(arima_residuals)), arima_residuals, 0, alpha=0.3, color='red')
ax3.set_title('ARIMA Prediction Errors', fontsize=12, fontweight='bold')
ax3.set_ylabel('Errors')
ax3.legend()
ax3.grid(True, alpha=0.3)

# Plot 4: Actual vs Hybrid
ax4 = axes[1, 1]
ax4.plot(actuals, label='Actual', color='black', linewidth=2, alpha=0.7)
ax4.plot(hybrid_pred, label='Hybrid Method 2', color='purple', linewidth=1.5, linestyle='--', alpha=0.7)
ax4.set_title('Final Hybrid Forecast (Method 2)', fontsize=12, fontweight='bold')
ax4.set_ylabel('Log Returns')
ax4.legend()
ax4.grid(True, alpha=0.3)
ax4.text(0.02, 0.98, f'RMSE: {window["hybrid_rmse"]:.6f}',
         transform=ax4.transAxes, verticalalignment='top',
         bbox=dict(boxstyle='round', facecolor='wheat', alpha=0.5))

# Plot 5: Component Comparison
ax5 = axes[2, 0]
sample_size = min(50, len(actuals))
x = range(sample_size)
ax5.plot(x, actuals[:sample_size], label='Actual', color='black', linewidth=2, alpha=0.7)
ax5.plot(x, arima_pred[:sample_size], label='ARIMA', color='blue', linewidth=1.5, linestyle='--', alpha=0.7)
ax5.plot(x, hybrid_pred[:sample_size], label='Hybrid', color='purple', linewidth=1.5, linestyle='--', alpha=0.7)
ax5.set_title(f'First {sample_size} Predictions Comparison', fontsize=12, fontweight='bold')
ax5.set_xlabel('Time Step')
ax5.set_ylabel('Log Returns')
ax5.legend()
ax5.grid(True, alpha=0.3)

# Plot 6: Performance Metrics
ax6 = axes[2, 1]
metrics = ['RMSE', 'MAE', 'Direction Acc (%)']
arima_metrics = [window['arima_rmse'], window['arima_mae'], window['arima_direction_accuracy']]
hybrid_metrics = [window['hybrid_rmse'], window['hybrid_mae'], window['hybrid_direction_accuracy']]

x_pos = np.arange(len(metrics))

```

```

width = 0.35

bars1 = ax6.bar(x_pos - width/2, arima_metrics, width, label='ARIMA', alpha=0.8)
bars2 = ax6.bar(x_pos + width/2, hybrid_metrics, width, label='Hybrid', alpha=0.8)

ax6.set_ylabel('Value')
ax6.set_title('Performance Metrics Comparison', fontsize=12, fontweight='bold')
ax6.set_xticks(x_pos)
ax6.set_xticklabels(metrics)
ax6.legend()
ax6.grid(True, alpha=0.3, axis='y')

# Add value labels
for bars in [bars1, bars2]:
    for bar in bars:
        height = bar.get_height()
        ax6.text(bar.get_x() + bar.get_width()/2., height,
                 f'{height:.4f}', ha='center', va='bottom', fontsize=8)

plt.tight_layout()
plt.show()

# Print detailed stats
print(f"\n{'*'*80}")
print(f"DETAILED STATISTICS - Window {window['window_id']}")
print(f"\n{'*'*80}")
print(f"Period: {window['test_start'].strftime('%Y-%m-%d')} to {window['test_end'].strftime('%Y-%m-%d')}")
print(f"ARIMA Order: {window['arima_order']}, AIC: {window['arima_aic']:.2f}")
print(f"\nARIMA Performance:")
print(f"    • RMSE: {window['arima_rmse']:.6f}")
print(f"    • MAE: {window['arima_mae']:.6f}")
print(f"    • Direction Accuracy: {window['arima_direction_accuracy']:.2f}%")
print(f"\nHybrid Performance:")
print(f"    • RMSE: {window['hybrid_rmse']:.6f}")
print(f"    • MAE: {window['hybrid_mae']:.6f}")
print(f"    • Direction Accuracy: {window['hybrid_direction_accuracy']:.2f}%")
print(f"\nImprovement: {window['rmse_improvement_pct']+2f}% RMSE")
print(f"\n{'*'*80}")

def plot_all_windows_comparison(results):
    if not results or len(results['window_results']) == 0:
        print("No results to plot")
    return

```

```

asset_name = results['asset_name']
windows = [r['window_id'] for r in results['window_results']]
arima_rmse = [r['arima_rmse'] for r in results['window_results']]
hybrid_rmse = [r['hybrid_rmse'] for r in results['window_results']]
improvement = [r['rmse_improvement_pct'] for r in results['window_results']]

fig, axes = plt.subplots(1, 2, figsize=(14, 5))
fig.suptitle(f'{asset_name} - Hybrid Model Performance Across All Windows',
             fontsize=14, fontweight='bold')

# Plot 1: RMSE Comparison
ax1 = axes[0]
ax1.plot(windows, arima_rmse, marker='o', label='ARIMA', linewidth=2, markersize=6)
ax1.plot(windows, hybrid_rmse, marker='s', label='Hybrid', linewidth=2, markersize=6)
ax1.set_xlabel('Window ID')
ax1.set_ylabel('RMSE')
ax1.set_title('RMSE by Window', fontsize=12, fontweight='bold')
ax1.legend()
ax1.grid(True, alpha=0.3)

# Plot 2: Improvement
ax2 = axes[1]
colors = ['green' if x > 0 else 'red' for x in improvement]
ax2.bar(windows, improvement, color=colors, alpha=0.7)
ax2.axhline(y=0, color='black', linestyle='-', linewidth=1)
ax2.set_xlabel('Window ID')
ax2.set_ylabel('RMSE Improvement (%)')
ax2.set_title('Hybrid Improvement over ARIMA', fontsize=12, fontweight='bold')
ax2.grid(True, alpha=0.3, axis='y')

plt.tight_layout()
plt.show()

```

In []: plot_all_windows_comparison(sp500_hybrid_results)
plot_all_windows_comparison(bitcoin_hybrid_results)

EURUSD Dataset Analysis

ARIMA-LSTM Hybrid Method 2 (Non-additive)

Dataset: EUR/USD Exchange Rate (2009-08-11 to 2019-08-11)

```
In [ ]: eurusd_data = yf.download("EURUSD=X", start="2009-08-11", end="2019-08-11", progress=False)

print(f"\nEURUSD Data Shape: {eurusd_data.shape}")
print(f"EURUSD Date Range: {eurusd_data.index.min()} to {eurusd_data.index.max()}")
print(f"Total EURUSD observations: {len(eurusd_data)}")

# Calculate log returns
eurusd_data['Log_Returns'] = np.log(eurusd_data['Close'] / eurusd_data['Close'].shift(1))

# Clean data
eurusd_clean = eurusd_data.dropna()

print(f"\nAfter cleaning: {len(eurusd_clean)} observations")
print(f"Mean daily return: {eurusd_clean['Log_Returns'].mean():.6f}")
print(f"Standard deviation: {eurusd_clean['Log_Returns'].std():.6f}")
```

```
In [ ]: from dateutil.relativedelta import relativedelta
from datetime import timedelta

def create_eurusd_cv_splits(data, start_date=None):

    if start_date is None:
        start_date = data.index.min()

    cv_splits = []
    window_start = start_date

    while True:
        # Define window boundaries
        train_start = window_start
        train_end = train_start + relativedelta(years=2) - timedelta(days=1)

        # Validation periods (8, 16, 24 months)
        val_start = train_end + timedelta(days=1)
```

```
val1_end = val_start + relativedelta(months=8) - timedelta(days=1)
val2_end = val_start + relativedelta(months=16) - timedelta(days=1)
val3_end = val_start + relativedelta(months=24) - timedelta(days=1)

# Test period (6 months)
test_start = val3_end + timedelta(days=1)
test_end = test_start + relativedelta(months=6) - timedelta(days=1)

# Check if we have enough data
if test_end > data.index.max():
    break

# Create splits for this window
train_data = data[(data.index >= train_start) & (data.index <= train_end)]

# Three validation folds
val1_data = data[(data.index >= val_start) & (data.index <= val1_end)]
val2_data = data[(data.index >= val_start) & (data.index <= val2_end)]
val3_data = data[(data.index >= val_start) & (data.index <= val3_end)]

test_data = data[(data.index >= test_start) & (data.index <= test_end)]

cv_splits.append({
    'window_id': len(cv_splits) + 1,
    'train': {
        'data': train_data,
        'start': train_start,
        'end': train_end,
        'size': len(train_data)
    },
    'validation': [
        {
            'fold': 1,
            'data': val1_data,
            'start': val_start,
            'end': val1_end,
            'size': len(val1_data),
            'months': 8
        },
        {
            'fold': 2,
            'data': val2_data,
            'start': val_start,
            'end': val2_end,
            'size': len(val2_data),
            'months': 8
        }
    ]
})
```

```

        'months': 16
    },
    {
        'fold': 3,
        'data': val3_data,
        'start': val_start,
        'end': val3_end,
        'size': len(val3_data),
        'months': 24
    }
],
'test': {
    'data': test_data,
    'start': test_start,
    'end': test_end,
    'size': len(test_data)
}
})
)

# Move window forward by 1 year
window_start += relativedelta(years=1)

return cv_splits

# Create EURUSD CV splits
eurusd_cv_splits = create_eurusd_cv_splits(eurusd_clean)

print(f"\n==== EURUSD Cross-Validation Setup ===")
print(f"Total CV windows: {len(eurusd_cv_splits)}")
print(f"\nFirst window details:")
print(f"  Train: {eurusd_cv_splits[0]['train']['start'].date()} to {eurusd_cv_splits[0]['train']['end'].date()} ({eurusd_cv_splits[0]['train']['size']} months)")
print(f"  Val 1: {eurusd_cv_splits[0]['validation'][0]['start'].date()} to {eurusd_cv_splits[0]['validation'][0]['end'].date()} ({eurusd_cv_splits[0]['validation'][0]['size']} months)")
print(f"  Test: {eurusd_cv_splits[0]['test']['start'].date()} to {eurusd_cv_splits[0]['test']['end'].date()} ({eurusd_cv_splits[0]['test']['size']} months)")

```

```
In [ ]: eurusd_hybrid_results = run_hybrid_arima_lstm_cv(
    cv_splits=eurusd_cv_splits,
    data_clean=eurusd_clean,
    asset_name='EURUSD',
    max_p=3,
    max_d=1,
    max_q=3,
    lookback=60,
    lstm_units=50,
    dropout=0.2,
```

```
    epochs=50,  
    batch_size=32,  
    verbose=True  
)
```

```
In [ ]: plot_all_windows_comparison(eurusd_hybrid_results)
```

Trading Performance Metrics

```
In [ ]: def volatility_predictions_to_returns_new(predictions, true_values, actual_returns, transaction_costs=0.0):  
    # Ensure all arrays have matching length  
    min_len = min(len(predictions), len(true_values), len(actual_returns))  
    predictions = predictions[:min_len]  
    true_values = true_values[:min_len]  
    actual_returns = (actual_returns.iloc[:min_len]  
                      if isinstance(actual_returns, pd.Series)  
                      else actual_returns[:min_len])  
  
    # Convert to numpy arrays for consistency  
    actual_returns_array = (actual_returns.values  
                           if isinstance(actual_returns, pd.Series)  
                           else actual_returns)  
  
    signal = np.where(predictions > transaction_costs, 1, -1)  
  
    if transaction_costs > 0.0:  
        signals = np.where(np.abs(actual_returns_array) > transaction_costs, signal, 0)  
  
    strategy_returns = signals * actual_returns_array  
  
    return pd.Series(strategy_returns)
```

```
In [ ]: def get_all_predictions(model_results, data_clean, model_type="S&P", window_indices=None):  
  
    if model_type == "S&P":  
        cost = 0.005  
    elif model_type == "Bitcoin":  
        cost = 0.01  
    else:  
        cost = 0.001  
    all_strategy_returns = []
```

```

windows_to_use = model_results['window_results']
if window_indices is not None:
    windows_to_use = [w for w in windows_to_use if w['window_id'] in window_indices]

for window_result in windows_to_use:
    try:
        test_start = window_result['test_start']
        test_end = window_result['test_end']
        test_data = data_clean[test_start:test_end]

        predictions = window_result['hybrid_predictions']
        true_values = test_data['Log_Returns'].values[-len(predictions):]

        actual_returns = test_data['Log_Returns'].iloc[-len(predictions):]

        window_returns = volatility_predictions_to_returns_new(
            predictions, true_values, actual_returns.values, transaction_costs=cost
        )

        all_strategy_returns.append(window_returns)

    except Exception as e:
        print(f"Warning: Failed to process window {window_result.get('window_id', '?')}: {str(e)}")
        continue

# Concatenate all returns
if all_strategy_returns:
    return pd.concat(all_strategy_returns, ignore_index=True)
else:
    return pd.Series([])

```

In []:

```

def annualized_return(daily_returns):
    cumulative = (1 + daily_returns).prod()
    n = daily_returns.shape[0]
    return cumulative ** (TRADING_DAYS / n) - 1

def annualized_std(daily_returns):
    return daily_returns.std() * np.sqrt(TRADING_DAYS)

def max_drawdown(daily_returns):
    equity = (1 + daily_returns).cumprod()

```

```

peak = equity.cummax()
drawdown = (equity - peak) / peak
return np.abs(drawdown.min()) # Paper uses absolute value

def information_ratio(strategy_returns, benchmark_returns):
    arc = annualized_return(strategy_returns)
    asd = annualized_std(strategy_returns)

    if asd == 0:
        return np.nan
    return arc / asd

def modified_information_ratio(strategy_returns, benchmark_returns):
    arc = annualized_return(strategy_returns)
    asd = annualized_std(strategy_returns)
    md = max_drawdown(strategy_returns)

    if asd == 0 or md == 0:
        return np.nan

    return (arc * np.sign(arc) * arc) / (asd * md)

def sortino_ratio(daily_returns, risk_free_rate=0):
    negative_returns = daily_returns[daily_returns < 0]

    if len(negative_returns) == 0:
        return np.nan

    downside_std = np.std(negative_returns, ddof=1)
    asd_downside = downside_std * np.sqrt(TRADING_DAYS)

    arc = annualized_return(daily_returns)

    if asd_downside == 0:
        return np.nan

    return arc / asd_downside

def compute_performance_indicators(strategy_returns, benchmark_returns):
    return {
        "ARC": annualized_return(strategy_returns),

```

```
        "ASD": annualized_std(strategy_returns),
        "MD": abs(max_drawdown(strategy_returns)),
        "IR": information_ratio(strategy_returns, benchmark_returns),
        "IR*": modified_information_ratio(strategy_returns, benchmark_returns),
        "SR": sortino_ratio(strategy_returns)
    }
```

S&P 500

```
In [ ]: TRADING_DAYS = 232

# Get benchmark returns (Buy-and-Hold)
sp500_bnh_returns = sp500_clean['Log_Returns'].loc["2007-01-01":"2023-12-29"].values

sp500_hybrid_predictions = get_all_predictions(sp500_hybrid_results, sp500_clean)

sp500_hybrid_strategy_returns = sp500_hybrid_predictions

sp500_bnh_aligned = sp500_bnh_returns[-len(sp500_hybrid_strategy_returns):]

results_sp500 = []

# HYBRID
hybrid_metrics = compute_performance_indicators(
    pd.Series(sp500_hybrid_strategy_returns),
    pd.Series(sp500_bnh_aligned)
)
hybrid_metrics['Model'] = 'HYBRID'
hybrid_metrics['Num_Trades'] = int(np.sum(np.abs(np.diff(sp500_hybrid_strategy_returns > 0)) > 0))
results_sp500.append(hybrid_metrics)

table2_sp500 = pd.DataFrame(results_sp500)

print("TABLE: S&P 500 Long-Short Strategy Results")

print(table2_sp500[['Model', 'ARC', 'ASD', 'MD', 'IR', 'IR*', 'SR']].to_string(index=False))

table2_sp500.to_csv('table2_sp500.csv', index=False)
```

Bitcoin

```
In [ ]: TRADING_DAYS = 345

bitcoin_bnh_returns = bitcoin_clean['Log_Returns'].values

bitcoin_hybrid_predictions = get_all_predictions(bitcoin_hybrid_results, bitcoin_clean, model_type="bitcoin")

bitcoin_hybrid_strategy_returns = bitcoin_hybrid_predictions

bitcoin_bnh_aligned = bitcoin_bnh_returns[-len(bitcoin_hybrid_strategy_returns):]

results_bitcoin = []

# HYBRID
hybrid_metrics = compute_performance_indicators(
    pd.Series(bitcoin_hybrid_strategy_returns),
    pd.Series(bitcoin_bnh_aligned)
)
hybrid_metrics['Model'] = 'HYBRID'
hybrid_metrics['Num_Trades'] = int(np.sum(np.abs(np.diff(bitcoin_hybrid_strategy_returns > 0)) > 0))
results_bitcoin.append(hybrid_metrics)

table2_bitcoin = pd.DataFrame(results_bitcoin)

print("TABLE: Bitcoin Long-Short Strategy Results")
print(table2_bitcoin[['Model', 'ARC', 'ASD', 'MD', 'IR', 'IR*', 'SR']].to_string(index=False))

table2_bitcoin.to_csv('table2_bitcoin.csv', index=False)
```

EURUSD

```
In [ ]: TRADING_DAYS = 252

# Get benchmark returns (Buy-and-Hold)
eurusd_bnh_returns = eurusd_clean['Log_Returns'].values

eurusd_hybrid_predictions = get_all_predictions(eurusd_hybrid_results, eurusd_clean, model_type="EURUSD")

eurusd_hybrid_strategy_returns = eurusd_hybrid_predictions
```

```
eurusd_bnh_aligned = eurusd_bnh_returns[-len(eurusd_hybrid_strategy_returns):]

results_eurusd = []

# HYBRID
hybrid_metrics = compute_performance_indicators(
    pd.Series(eurusd_hybrid_strategy_returns),
    pd.Series(eurusd_bnh_aligned)
)
hybrid_metrics['Model'] = 'HYBRID'
hybrid_metrics['Num_Trades'] = int(np.sum(np.abs(np.diff(eurusd_hybrid_strategy_returns > 0)) > 0))
results_eurusd.append(hybrid_metrics)

table2_eurusd = pd.DataFrame(results_eurusd)

print(table2_eurusd[['Model', 'ARC', 'ASD', 'MD', 'IR', 'IR*', 'SR']].to_string(index=False))

print("TABLE: EURUSD Long-Short Strategy Results")
table2_eurusd.to_csv('table2_eurusd.csv', index=False)
print("\n✓ Results saved to 'table2_eurusd.csv'")
```